

The Hardware/Software Interface

CSE351 Spring 2013

Integers

Roadmap

C:

```
car *c = malloc(sizeof(car));
c->miles = 100;
c->gals = 17;
float mpg = get_mpg(c);
free(c);
```

Java:

```
Car c = new Car();
c.setMiles(100);
c.setGals(17);
float mpg =
    c.getMPG();
```

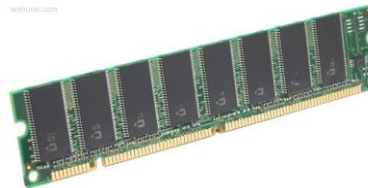
Assembly
language:

```
get_mpg:
    pushq    %rbp
    movq     %rsp, %rbp
    ...
    popq     %rbp
    ret
```

Machine
code:

```
0111010000011000
100011010000010000000010
1000100111000010
110000011111101000011111
```

Computer
system:



Data & addressing
Integers & floats
 Machine code & C
 x86 assembly
 programming
 Procedures &
 stacks
 Arrays & structs
 Memory & caches
 Processes
 Virtual memory
 Memory allocation
 Java vs. C

OS:




Today's Topics

- Representation of integers: unsigned and signed
- Casting
- Arithmetic and shifting
- Sign extension

But before we get to integers....

- How about encoding a standard deck of playing cards?
- 52 cards in 4 suits
 - How do we encode suits, face cards?
- What operations do we want to make easy to implement?
 - Which is the higher value card?
 - Are they the same suit?



Two possible representations

- 52 cards – 52 bits with bit corresponding to card set to 1



low-order 52 bits of 64-bit word

- “One-hot” encoding
- Drawbacks:
 - Hard to compare values and suits
 - Large number of bits required

Two possible representations

- 52 cards – 52 bits with bit corresponding to card set to 1



low-order 52 bits of 64-bit word

- “One-hot” encoding
- Drawbacks:
 - Hard to compare values and suits
 - Large number of bits required

- 4 bits for suit, 13 bits for card value – 17 bits with two set to 1



- “Two-hot” (?) encoding
- Easier to compare suits and values
 - Still an excessive number of bits

Two better representations

■ Binary encoding of all 52 cards – only 6 bits needed



low-order 6 bits of a byte

- Fits in one byte
- Smaller than one-hot or two-hot encoding, but how can we make value and suit comparisons easier?

■ Binary encoding of suit (2 bits) and value (4 bits) separately



suit

value

- Also fits in one byte, and easy to do comparisons


Some basic operations

■ Checking if two cards have the same suit:

```
#define SUIT_MASK 0x30
char array[5];          // represents a 5 card hand
char card1, card2;      // two cards to compare
card1 = array[0];
card2 = array[1];
...
if sameSuitP(card1, card2) {
...

```

SUIT_MASK = 0x30;



```
bool sameSuitP(char card1, char card2) {
    return (! (card1 & SUIT_MASK) ^ (card2 & SUIT_MASK));
    //return (card1 & SUIT_MASK) == (card2 & SUIT_MASK);
}
```


Some basic operations

■ Comparing the values of two cards:

```
#define SUIT_MASK 0x30
#define VALUE_MASK 0x0F

char array[5];          // represents a 5 card hand
char card1, card2;      // two cards to compare
card1 = array[0];
card2 = array[1];
...
if greaterValue(card1, card2) {
...
bool greaterValue(char card1, char card2) {
    return ((unsigned int)(card1 & VALUE_MASK) >
            (unsigned int)(card2 & VALUE_MASK));
}
```

VALUE_MASK = 0x0F;

0	0	0	0	1	1	1	1
---	---	---	---	---	---	---	---

suit

value

Encoding Integers

- **The hardware (and C) supports two flavors of integers:**
 - unsigned – only the non-negatives
 - signed – both negatives and non-negatives
- **There are only 2^W distinct bit patterns of W bits, so...**
 - Can't represent all the integers
 - Unsigned values are $0 \dots 2^W - 1$
 - Signed values are $-2^{W-1} \dots 2^{W-1} - 1$
- **Reminder: terminology for binary representations:**

“Most-significant” or
“high-order” bit(s)

“Least-significant” or
“low-order” bit(s)



0110010110101001

Unsigned Integers

- Unsigned values are just what you expect

- $b_7b_6b_5b_4b_3b_2b_1b_0 = b_72^7 + b_62^6 + b_52^5 + \dots + b_12^1 + b_02^0$
 - Useful formula: $1+2+4+8+\dots+2^{N-1} = 2^N - 1$

- You add/subtract them using the normal “carry/borrow” rules, just in binary

00111111
+00000001
01000000

63
+ 8
71

- How would you make *signed* integers?

Signed Integers


- **Let's do the natural thing for the positives**
 - They correspond to the unsigned integers of the same value
 - Example (8 bits): $0x00 = 0$, $0x01 = 1$, ..., $0x7F = 127$
- **But, we need to let about half of them be negative**
 - Use the high order bit to indicate *negative*: call it the “sign bit”
 - Call this a “sign-and-magnitude” representation
 - Examples (8 bits):
 - $0x00 = 00000000_2$ is non-negative, because the sign bit is 0
 - $0x7F = 01111111_2$ is non-negative
 - $0x85 = 10000101_2$ is negative
 - $0x80 = 10000000_2$ is negative...

Sign-and-Magnitude Negatives

■ How should we represent -1 in binary?

- Sign-and-magnitude: 10000001_2

Use the MSB for + or -, and the other bits to give magnitude




Sign-and-Magnitude Negatives

■ How should we represent -1 in binary?

- Sign-and-magnitude: 10000001_2

Use the MSB for + or -, and the other bits to give magnitude

(Unfortunate side effect: there are two representations of 0!)



Sign-and-Magnitude Negatives

■ How should we represent -1 in binary?


- Sign-and-magnitude: 10000001_2
Use the MSB for + or -, and the other bits to give magnitude
(Unfortunate side effect: there are two representations of 0!)
- Another problem: math is cumbersome

- Example:

$$4 - 3 \neq 4 + (-3)$$

↓


0100
+ 1011
1111



How do we solve these problems?

Two's Complement Negatives

- How should we represent -1 in binary?




Two's Complement Negatives

■ How should we represent -1 in binary?

- Rather than a sign bit, let MSB have same value, but *negative* weight
 - W-bit word: Bits 0, 1, ..., W-2 add $2^0, 2^1, \dots, 2^{W-2}$ to value of integer when set, but bit W-1 adds -2^{W-1} when set
 - e.g. unsigned 1010_2 : $1*2^3 + 0*2^2 + 1*2^1 + 0*2^0 = 10_{10}$
 2's comp. 1010_2 : $-1*2^3 + 0*2^2 + 1*2^1 + 0*2^0 = -6_{10}$

- So -1 represented as 1111_2 ; all negative integers still have MSB = 1
- Advantages of two's complement: only one zero, simple arithmetic
- To get negative representation of any integer, take bitwise complement and then add one!

$$\sim x + 1 = -x$$



Two's Complement Arithmetic

- The same addition procedure works for both unsigned and two's complement integers
 - Simplifies hardware: only one adder needed
 - Algorithm: simple addition, discard the highest carry bit
 - Called “modular” addition: result is sum *modulo* 2^w
- Examples:

4 0100	4 0100	− 4 1100
+ 3 + 0011	− 3 + 1101	+ 3 + 0011
= 7 = 0111	= 1 1 0001	− 1 1111
	drop carry = 0001	

Two's Complement

■ Why does it work?

- Put another way: given the bit representation of a positive integer, we want the negative bit representation to always sum to 0 (ignoring the carry-out bit) when added to the positive representation
- This turns out to be the *bitwise complement plus one*
 - What should the 8-bit representation of -1 be?

$$\begin{array}{r}
 00000001 \\
 + \text{????????} \\
 \hline
 00000000
 \end{array}
 \quad \text{(we want whichever bit string gives the right result)}$$

$$\begin{array}{r}
 00000010 \\
 + \text{????????} \\
 \hline
 00000000
 \end{array}
 \qquad
 \begin{array}{r}
 00000011 \\
 + \text{????????} \\
 \hline
 00000000
 \end{array}$$

Two's Complement

■ Why does it work?

- Put another way: given the bit representation of a positive integer, we want the negative bit representation to always sum to 0 (ignoring the carry-out bit) when added to the positive representation
- This turns out to be the *bitwise complement plus one*
 - What should the 8-bit representation of -1 be?

$$\begin{array}{r}
 00000001 \\
 +11111111 \\
 \hline
 100000000
 \end{array}
 \quad \text{(we want whichever bit string gives the right result)}$$

$$\begin{array}{r}
 00000010 \\
 +???????? \\
 \hline
 00000000
 \end{array}
 \qquad
 \begin{array}{r}
 00000011 \\
 +???????? \\
 \hline
 00000000
 \end{array}$$

Two's Complement

■ Why does it work?

- Put another way: given the bit representation of a positive integer, we want the negative bit representation to always sum to 0 (ignoring the carry-out bit) when added to the positive representation
- This turns out to be the *bitwise complement plus one*
 - What should the 8-bit representation of -1 be?

$$\begin{array}{r}
 00000001 \\
 +11111111 \\
 \hline
 100000000
 \end{array}
 \quad \text{(we want whichever bit string gives the right result)}$$

$$\begin{array}{r}
 00000010 \\
 +11111110 \\
 \hline
 100000000
 \end{array}
 \qquad
 \begin{array}{r}
 00000011 \\
 +11111101 \\
 \hline
 100000000
 \end{array}$$

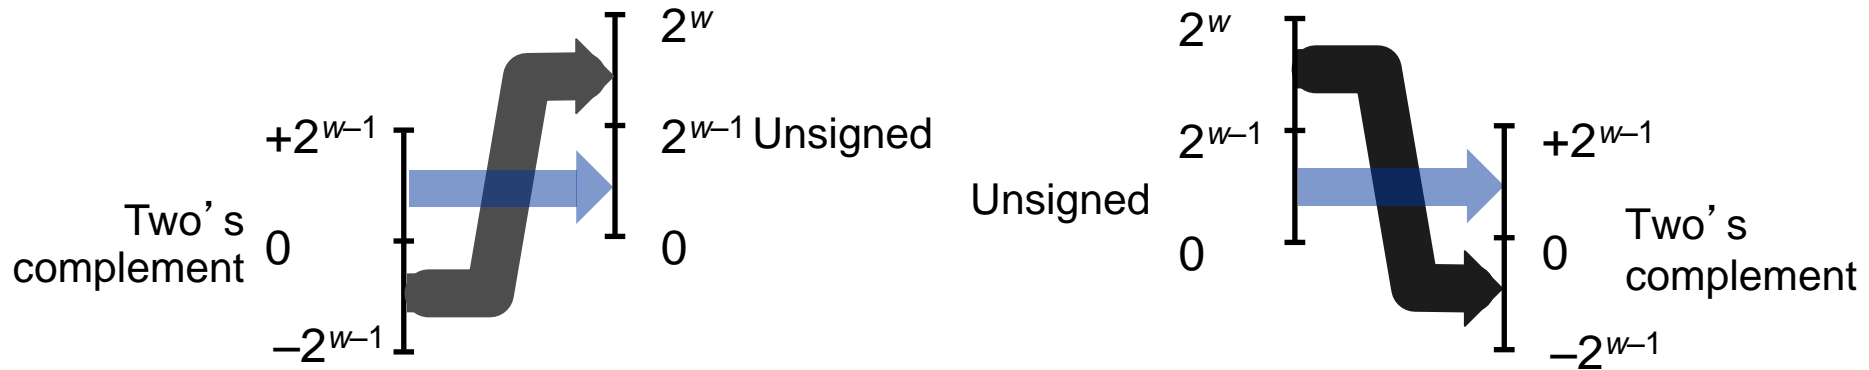
Unsigned & Signed Numeric Values

X	Unsigned	Signed
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

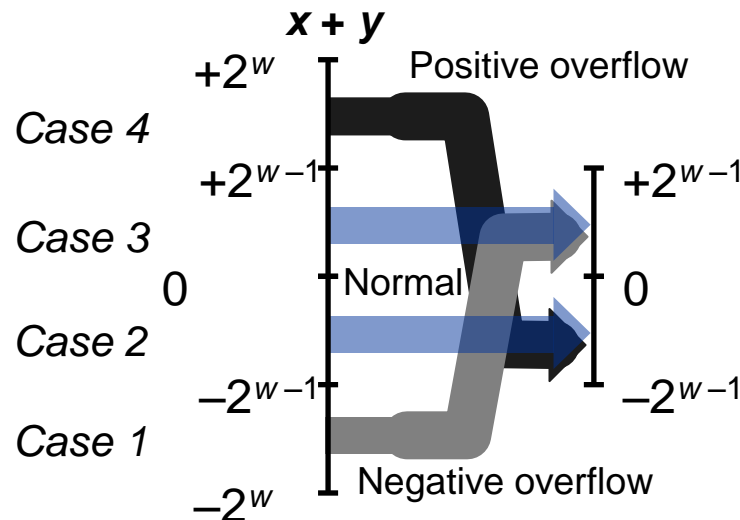
- Both signed and unsigned integers have limits
 - If you compute a number that is too big, you wrap: $6 + 4 = ?$ $15U + 2U = ?$
 - If you compute a number that is too small, you wrap: $-7 - 3 = ?$ $0U - 2U = ?$
 - Answers are only correct mod 2^b
- The CPU may be capable of “throwing an exception” for overflow on signed values
 - It won't for unsigned
- But C and Java just cruise along silently when overflow occurs...

Visualizations

- Same W bits interpreted as signed vs. unsigned:



- Two's complement (signed) addition: x and y are W bits wide



Values To Remember

■ Unsigned Values

- UMin = 0
 - 000...0
- UMax = $2^w - 1$
 - 111...1

■ Two's Complement Values

- TMin = -2^{w-1}
 - 100...0
- TMax = $2^{w-1} - 1$
 - 011...1
- Negative 1
 - 111...1 0xFFFFFFFF (32 bits)

Values for $W = 16$

	Decimal	Hex	Binary
UMax	65535	FF FF	11111111 11111111
TMax	32767	7F FF	01111111 11111111
TMin	-32768	80 00	10000000 00000000
-1	-1	FF FF	11111111 11111111
0	0	00 00	00000000 00000000

Values for Different Word Sizes

	W			
	8	16	32	64
UMax	255	65,535	4,294,967,295	18,446,744,073,709,551,615
TMax	127	32,767	2,147,483,647	9,223,372,036,854,775,807
TMin	-128	-32,768	-2,147,483,648	-9,223,372,036,854,775,808

■ Observations

- $|TMin| = TMax + 1$
 - Asymmetric range
- $UMax = 2 * TMax + 1$

■ C Programming

- `#include <limits.h>`
- Declares constants, e.g.:
 - `ULONG_MAX`
 - `LONG_MAX`
 - `LONG_MIN`
- Values are platform specific
- See: `/usr/include/limits.h` on Linux

Signed vs. Unsigned in C

■ Constants

- By default are considered to be signed integers
- Use “U” suffix to force unsigned:
 - `0U, 4294967259U`

Signed vs. Unsigned in C

■ Casting

- `int tx, ty;`
- `unsigned ux, uy;`
- Explicit casting between signed & unsigned:
 - `tx = (int) ux;`
 - `uy = (unsigned) ty;`
- Implicit casting also occurs via assignments and function calls:
 - `tx = ux;`
 - `uy = ty;`
 - The gcc flag *-Wsign-conversion* produces warnings for implicit casts, but *-Wall* does not!
- How does casting between signed and unsigned work – what values are going to be produced?
 - *Bits are unchanged*, just interpreted differently!

Casting Surprises

■ Expression Evaluation

- If you mix unsigned and signed in a single expression, then ***signed values implicitly cast to unsigned***
- Including comparison operations $<$, $>$, $==$, $<=$, $>=$
- Examples for $W = 32$: **TMIN = -2,147,483,648** **TMAX = 2,147,483,647**

■ Constant ₁	Constant ₂	Relation	Evaluation
0	0U	==	unsigned
-1	0	<	signed
-1	0U	>	unsigned
2147483647	-2147483648	>	signed
2147483647U	-2147483648	<	unsigned
-1	-2	>	signed
(unsigned)-1	-2	>	unsigned
2147483647	2147483648U	<	unsigned
2147483647	(int) 2147483648U	>	signed

Shift Operations

- **Left shift:** $x \ll y$
 - Shift bit-vector x left by y positions
 - Throw away extra bits on left
 - Fill with 0s on right
 - Equivalent to multiplying by 2^y (if no bits lost)
- **Right shift:** $x \gg y$
 - Shift bit-vector x right by y positions
 - Throw away extra bits on right
 - Logical shift (for unsigned values)
 - Fill with 0s on left
 - Arithmetic shift (for signed values)
 - Replicate most significant bit on left
 - Maintains sign of x
 - Equivalent to dividing by 2^y
 - Correct rounding (towards 0) requires some care with signed numbers

Argument x	01100010
$\ll 3$	00010000
Logical $\gg 2$	00011000
Arithmetic $\gg 2$	00011000

Argument x	10100010
$\ll 3$	00010000
Logical $\gg 2$	00101000
Arithmetic $\gg 2$	11101000

**Undefined behavior when
 $y < 0$ or $y \geq \text{word_size}$**

Using Shifts and Masks

■ Extract the 2nd most significant byte of an integer:

- First shift, then mask: $(x \gg 16) \& 0xFF$

x	01100001 01100010 01100011 01100100
$x \gg 16$	00000000 00000000 01100001 01100010
$(x \gg 16) \& 0xFF$	<div>00000000 00000000 00000000 11111111</div> <hr/> <div>00000000 00000000 00000000 01100010</div>

■ Extract the sign bit of a signed integer:

- $(x \gg 31) \& 1$ - need the “& 1” to clear out all other bits except LSB

■ Conditionals as Boolean expressions (**assuming x is 0 or 1**)

- if (x) a=y else a=z; which is the same as $a = x ? y : z;$
- Can be re-written (assuming arithmetic right shift) as:

$$a = ((x \ll 31) \gg 31) \& y + ((!x) \ll 31) \gg 31 \& z;$$


Sign Extension

■ Task:

- Given w -bit signed integer x
- Convert it to $w+k$ -bit integer *with same value*

■ Rule:

- Make k copies of sign bit:
- $X' = \underbrace{x_{w-1}, \dots, x_{w-1}}_{k \text{ copies of MSB}}, x_{w-1}, x_{w-2}, \dots, x_0$



Sign Extension Example

- Converting from smaller to larger integer data type
- C automatically performs sign extension

```
short int x = 12345;
int      ix = (int) x;
short int y = -12345;
int      iy = (int) y;
```

	Decimal	Hex	Binary
x	12345	30 39	00110000 01101101
ix	12345	00 00 30 39	00000000 00000000 00110000 01101101
y	-12345	CF C7	11001111 11000111
iy	-12345	FF FF CF C7	11111111 11111111 11001111 11000111