

The Hardware/Software Interface

CSE351 Spring 2013

Basics of Machine Programming

Roadmap

C:

```
car *c = malloc(sizeof(car));
c->miles = 100;
c->gals = 17;
float mpg = get_mpg(c);
free(c);
```

Java:

```
Car c = new Car();
c.setMiles(100);
c.setGals(17);
float mpg =
    c.getMPG();
```

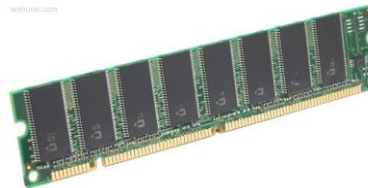
Assembly
language:

```
get_mpg:
    pushq    %rbp
    movq     %rsp, %rbp
    ...
    popq     %rbp
    ret
```

Machine
code:

```
0111010000011000
100011010000010000000010
1000100111000010
110000011111101000011111
```

Computer
system:



Data & addressing
Integers & floats
Machine code & C
x86 assembly
programming
Procedures &
stacks
Arrays & structs
Memory & caches
Processes
Virtual memory
Memory allocation
Java vs. C

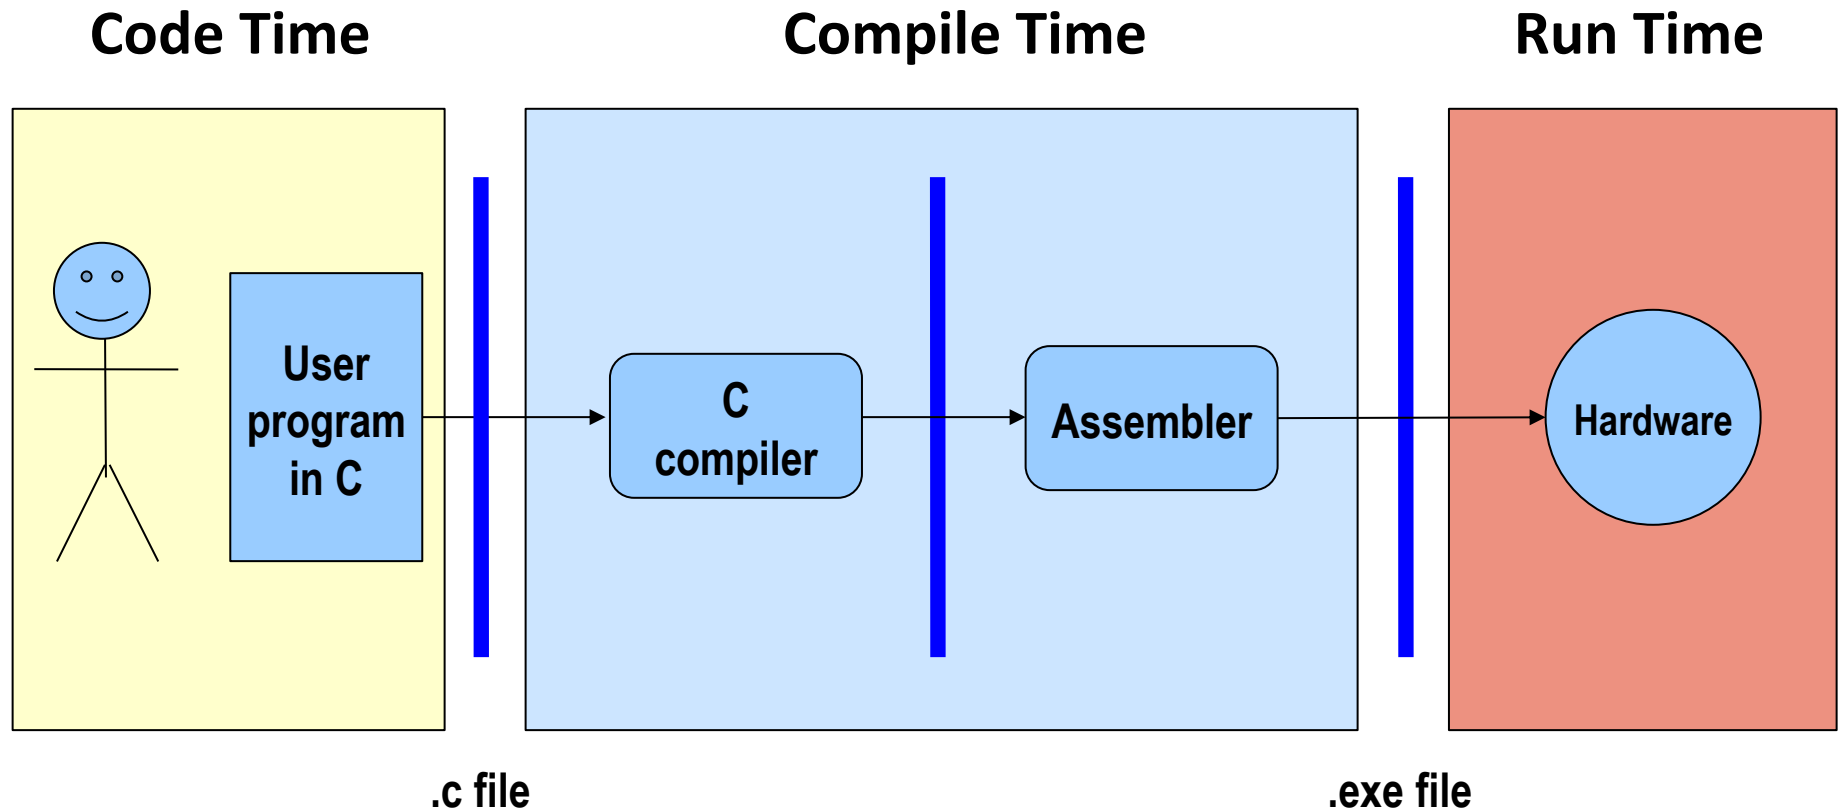
OS:



Today's Topics

- What is an ISA (Instruction Set Architecture)?
- A brief history of Intel processors and architectures
- C, assembly, machine code
- x86 basics: registers

Translation



What makes programs run fast?

Translation Impacts Performance

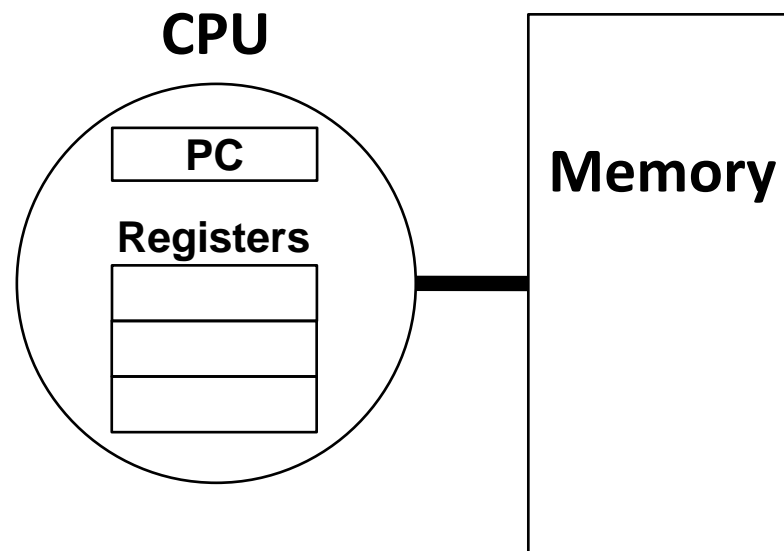
- **The time required to execute a program depends on:**
 - *The program* (as written in C, for instance)
 - *The compiler*: what set of assembler instructions it translates the C program into
 - *The instruction set architecture* (ISA): what set of instructions it makes available to the compiler
 - *The hardware implementation*: how much time it takes to execute an instruction

What should the HW/SW interface contain?

Instruction Set Architectures

■ The ISA defines:

- The system's state (e.g. registers, memory, program counter)
- The instructions the CPU can execute
- The effect that each of these instructions will have on the system state



General ISA Design Decisions

■ Instructions

- What instructions are available? What do they do?
- How are they encoded?

■ Registers

- How many registers are there?
- How wide are they?

■ Memory

- How do you specify a memory location?

x86

- **Processors that implement the x86 ISA completely dominate the server, desktop and laptop markets**
- **Evolutionary design**
 - Backwards compatible up until 8086, introduced in 1978
 - Added more features as time goes on
- **Complex instruction set computer (CISC)**
 - Many different instructions with many different formats
 - But, only small subset encountered with Linux programs
 - (as opposed to Reduced Instruction Set Computers (RISC), which use simpler instructions)

Intel x86 Evolution: Milestones

<i>Name</i>	<i>Date</i>	<i>Transistors</i>	<i>MHz</i>
■ 8086	1978	29K	5-10
<ul style="list-style-type: none">■ First 16-bit processor. Basis for IBM PC & DOS■ 1MB address space			
■ 386	1985	275K	16-33
<ul style="list-style-type: none">■ First 32 bit processor, referred to as IA32■ Added “flat addressing”■ Capable of running Unix■ 32-bit Linux/gcc targets i386 by default			
■ Pentium 4F	2005	230M	2800-3800
<ul style="list-style-type: none">■ First 64-bit Intel x86 processor, referred to as x86-64			

Intel x86 Processors

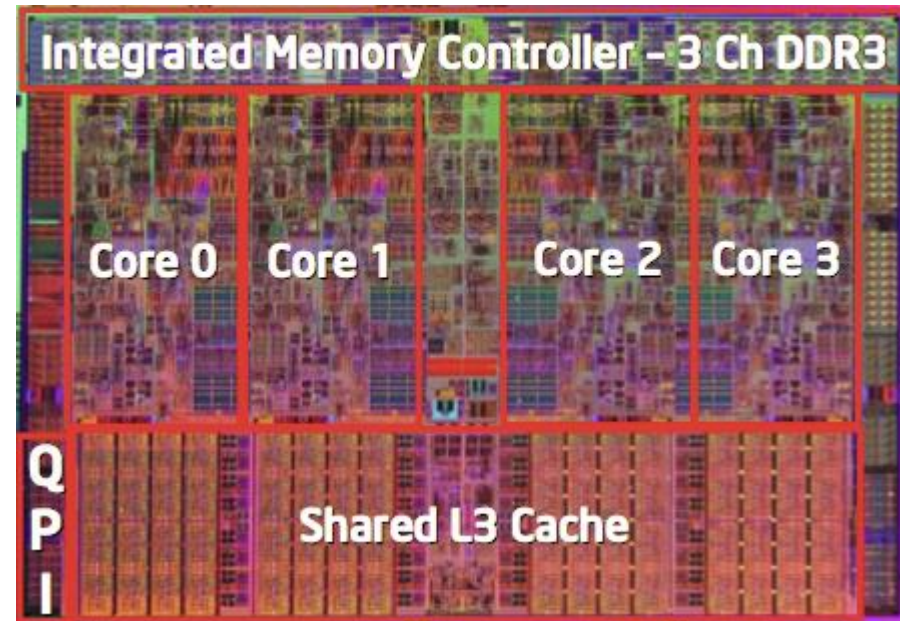
■ Machine Evolution

■ 486	1989
■ Pentium	1993
■ Pentium/MMX	1997
■ PentiumPro	1995
■ Pentium III	1999
■ Pentium 4	2001
■ Core 2 Duo	2006
■ Core i7	2008

■ Added Features

- Instructions to support multimedia operations
 - Parallel operations on 1, 2, and 4-byte data
- Instructions to enable more efficient conditional operations
- More cores!

Intel Core i7



More information

■ References for Intel processor specifications:

- Intel's "automated relational knowledgebase":
 - <http://ark.intel.com/>
- Wikipedia:
 - http://en.wikipedia.org/wiki/List_of_Intel_microprocessors

x86 Clones: Advanced Micro Devices (AMD)

■ Historically

- AMD has followed just behind Intel
- A little bit slower, a lot cheaper

■ Then

- Recruited top circuit designers from Digital Equipment and other downward trending companies
- Built Opteron: tough competitor to Pentium 4
- Developed x86-64, their own extension of x86 to 64 bits

Intel's Transition to 64-Bit

- **Intel attempted radical shift from IA32 to IA64 (2001)**
 - Totally different architecture (Itanium) and ISA than x86
 - Executes IA32 code only as legacy
 - Performance disappointing
- **AMD stepped in with *evolutionary* solution (2003)**
 - x86-64 (also called “AMD64”)
- **Intel felt obligated to focus on IA64**
 - Hard to admit mistake or that AMD is better
- **Intel announces “EM64T” extension to IA32 (2004)**
 - Extended Memory 64-bit Technology
 - Almost identical to AMD64!
- **Today: all but low-end x86 processors support x86-64**
 - But, lots of code out there is still just IA32

Our Coverage in 351

- **IA32**

- The traditional x86

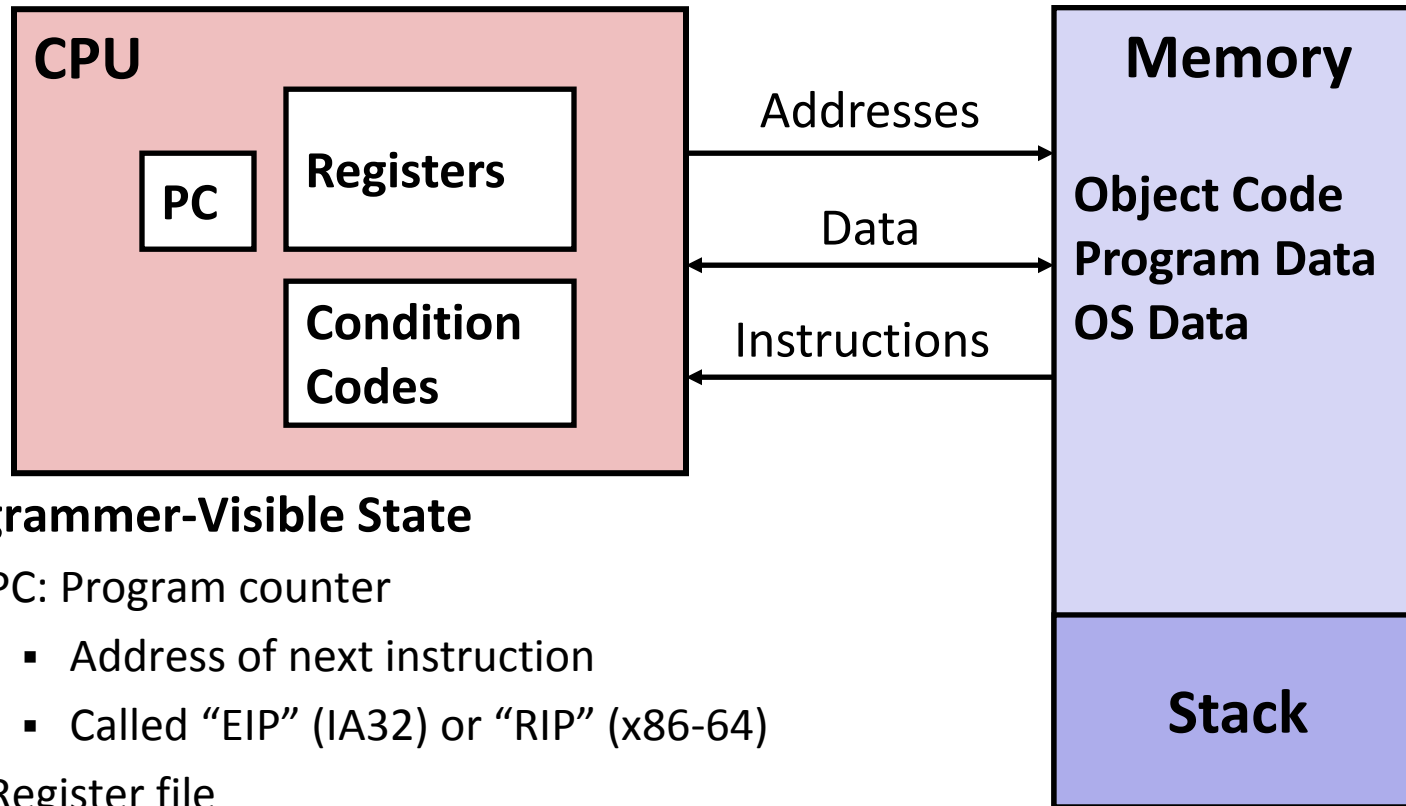
- **x86-64**

- The emerging standard – all lab assignments use x86-64!

Definitions

- **Architecture:** (also instruction set architecture or ISA)
The parts of a processor design that one needs to understand to write assembly code
 - “What is directly visible to software”
- **Microarchitecture: Implementation of the architecture**
 - Includes: CPU frequency, cache sizes, other implementation details
- Is cache size “architecture”?
- How about core frequency?
- And number of registers?

Assembly Programmer's View



■ Programmer-Visible State

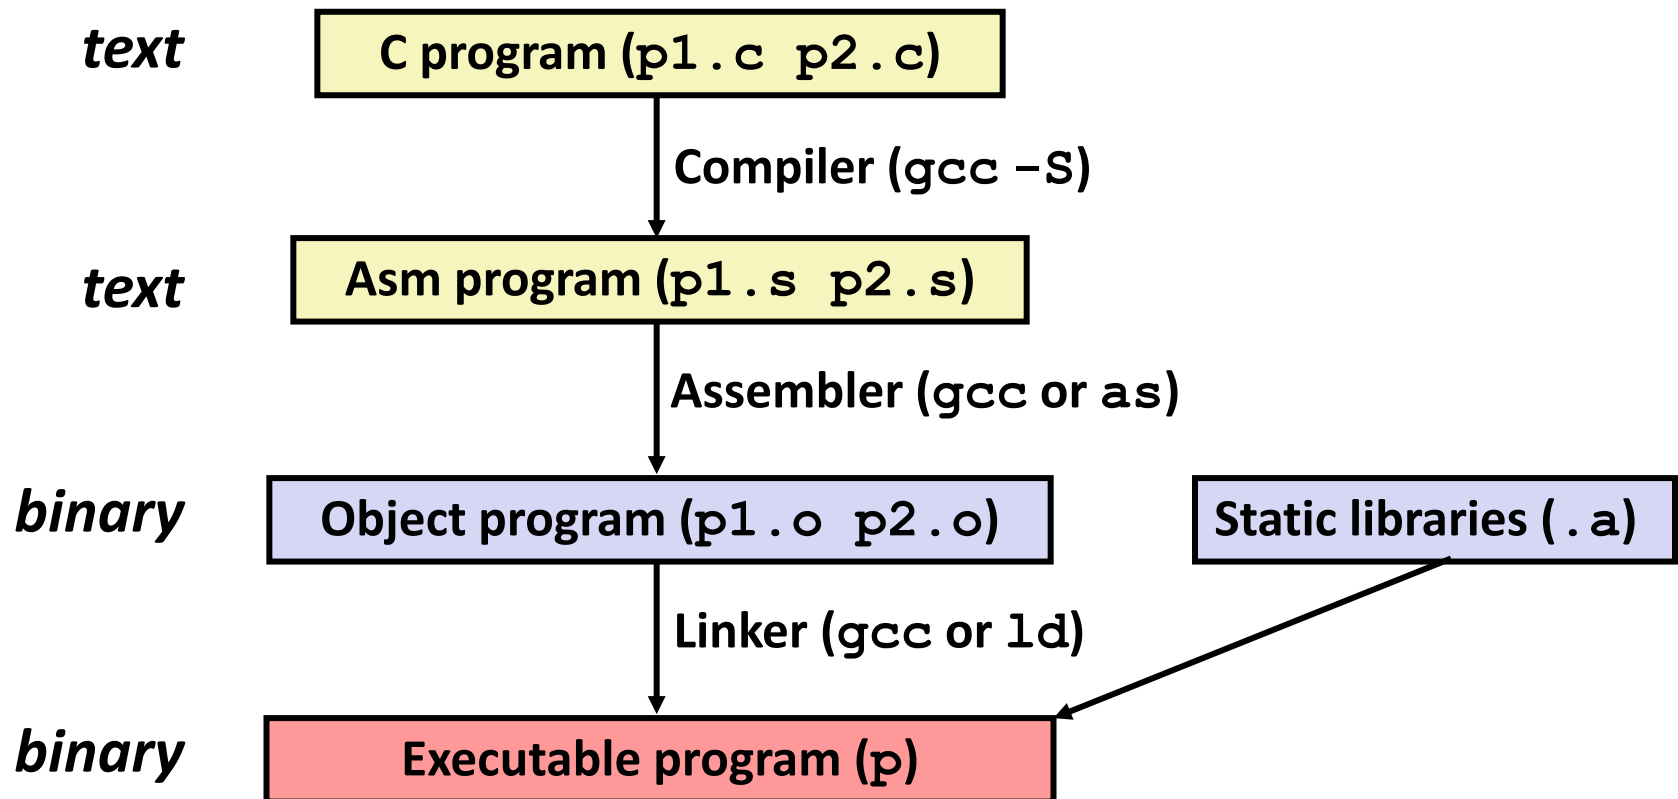
- PC: Program counter
 - Address of next instruction
 - Called “EIP” (IA32) or “RIP” (x86-64)
- Register file
 - Heavily used program data
- Condition codes
 - Store status information about most recent arithmetic operation
 - Used for conditional branching

■ Memory

- Byte addressable array
- Code, user data, (some) OS data
- Includes stack used to support procedures (we'll come back to that)

Turning C into Object Code

- Code in files `p1.c` `p2.c`
- Compile with command: `gcc -O1 p1.c p2.c -o p`
 - Use basic optimizations (`-O1`)
 - Put resulting binary in file `p`



Compiling Into Assembly

C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

Generated IA32 Assembly

```
sum:
    pushl %ebp
    movl %esp,%ebp
    movl 12(%ebp),%eax
    addl 8(%ebp),%eax
    movl %ebp,%esp
    popl %ebp
    ret
```

Obtain with command

```
gcc -O1 -S code.c
```

Produces file code.s

Three Basic Kinds of Instructions

- **Perform arithmetic function on register or memory data**
- **Transfer data between memory and register**
 - Load data from memory into register
 - Store register data into memory
- **Transfer control**
 - Unconditional jumps to/from procedures
 - Conditional branches

Assembly Characteristics: Data Types

- **“Integer” data of 1, 2, 4 (IA32), or 8 (just in x86-64) bytes**
 - Data values
 - Addresses (untyped pointers)
- **Floating point data of 4, 8, or 10 bytes**
- **What about “aggregate” types such as arrays or structs?**
 - No aggregate types, just contiguously allocated bytes in memory

Object Code

Code for `sum`

0x401040 <sum>:

0x55

0x89

0xe5

0x8b

0x45

0x0c

0x03

0x45

0x08

0x89

0xec

0x5d

0xc3

- Total of 13 bytes
- Each instruction 1, 2, or 3 bytes
- Starts at address 0x401040
- Not at all obvious where each instruction starts and ends

■ Assembler

- Translates `.s` into `.o`
- Binary encoding of each instruction
- Nearly-complete image of executable code
- Missing links between code in different files

■ Linker

- Resolves references between object files and (re)locates their data
- Combines with static run-time libraries
 - E.g., code for `malloc`, `printf`
- Some libraries are *dynamically linked*
 - Linking occurs when program begins execution

Machine Instruction Example

```
int t = x+y;
```

```
addl 8(%ebp), %eax
```

Similar to expression:

```
x += y
```

More precisely:

```
int eax;
```

```
int *ebp;
```

```
eax += ebp[2]
```

```
0x401046:    03 45 08
```

■ **C Code:** add two signed integers

■ **Assembly**

- Add two 4-byte integers
 - “Long” words in GCC speak
 - Same instruction whether signed or unsigned

■ Operands:

x: Register **%eax**

y: Memory **M[%ebp+8]**

t: Register **%eax**

–Return function value in **%eax**

■ **Object Code**

- 3-byte instruction
- Stored at address **0x401046**

Disassembling Object Code

Disassembled

```

00401040 <_sum>:
    0:      55          push    %ebp
    1:      89 e5       mov     %esp, %ebp
    3:      8b 45 0c    mov     0xc(%ebp), %eax
    6:      03 45 08    add     0x8(%ebp), %eax
    9:      89 ec       mov     %ebp, %esp
    b:      5d        pop     %ebp
    c:      c3        ret

```

■ Disassembler

`objdump -d p`

- Useful tool for examining object code (`man 1 objdump`)
- Analyzes bit pattern of series of instructions (delineates instructions)
- Produces near-exact rendition of assembly code
- Can be run on either `p` (complete executable) or `p1.o / p2.o` file

Alternate Disassembly

Object

```
0x401040:
  0x55
  0x89
  0xe5
  0x8b
  0x45
  0x0c
  0x03
  0x45
  0x08
  0x89
  0xec
  0x5d
  0xc3
```

Disassembled

```
0x401040 <sum>:      push    %ebp
0x401041 <sum+1>:     mov     %esp, %ebp
0x401043 <sum+3>:     mov     0xc(%ebp), %eax
0x401046 <sum+6>:     add     0x8(%ebp), %eax
0x401049 <sum+9>:     mov     %ebp, %esp
0x40104b <sum+11>:    pop     %ebp
0x40104c <sum+12>:    ret
```

■ Within gdb debugger

```
gdb p
```

```
disassemble sum
```

```
(disassemble function)
```

```
x/13b sum
```

```
(examine the 13 bytes starting at sum)
```

What Can be Disassembled?

```
% objdump -d WINWORD.EXE
```

```
WINWORD.EXE:      file format pei-i386
```

```
No symbols in "WINWORD.EXE".
```

```
Disassembly of section .text:
```

```
30001000 <.text>:
```

```
30001000:  55                push    %ebp
30001001:  8b ec            mov     %esp, %ebp
30001003:  6a ff            push    $0xffffffff
30001005:  68 90 10 00 30    push    $0x30001090
3000100a:  68 91 dc 4c 30    push    $0x304cdc91
```

- Anything that can be interpreted as executable code
- Disassembler examines bytes and reconstructs assembly source

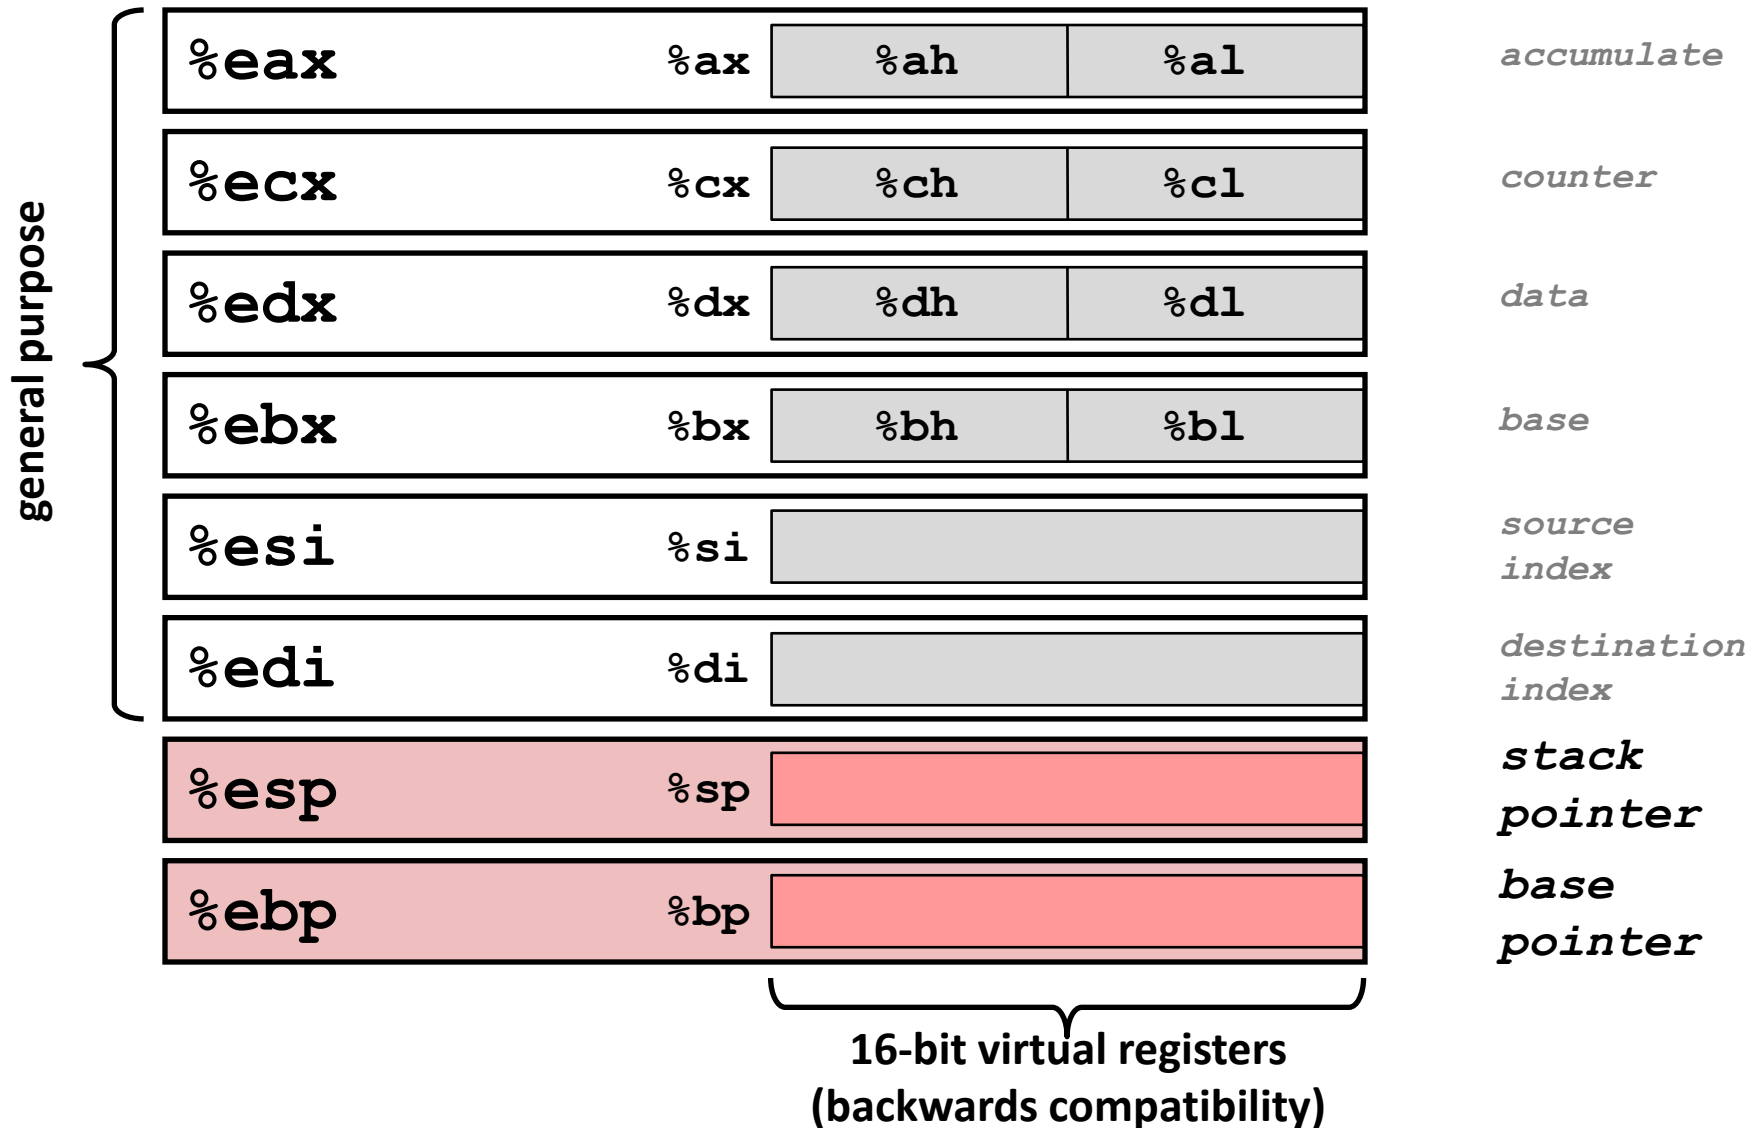
What Is A Register?

- A location in the CPU that stores a small amount of data, which can be accessed very quickly (once every clock cycle)
- Registers are at the heart of assembly programming
 - They are a precious commodity in all architectures, but *especially* x86

Integer Registers (IA32)



Integer Registers (IA32)



x86-64 Integer Registers

64-bits wide

%rax	%eax
%rbx	%ebx
%rcx	%ecx
%rdx	%edx
%rsi	%esi
%rdi	%edi
%rsp	%esp
%rbp	%ebp

%r8	%r8d
%r9	%r9d
%r10	%r10d
%r11	%r11d
%r12	%r12d
%r13	%r13d
%r14	%r14d
%r15	%r15d

- Extend existing registers, and add 8 new ones; *all* accessible as 8, 16, 32, 64 bits.

Summary: Machine Programming

■ What is an ISA (Instruction Set Architecture)?

- Defines the system's state and instructions that are available to the software

■ History of Intel processors and architectures

- Evolutionary design leads to many quirks and artifacts

■ C, assembly, machine code

- Compiler must transform statements, expressions, procedures into low-level instruction sequences

■ x86 registers

- Very limited number
- Not all general-purpose