

# **The Hardware/Software Interface**

CSE351 Spring 2013

## **Procedures and Stacks I**

# Roadmap

C:

```
car *c = malloc(sizeof(car));
c->miles = 100;
c->gals = 17;
float mpg = get_mpg(c);
free(c);
```

Java:

```
Car c = new Car();
c.setMiles(100);
c.setGals(17);
float mpg =
    c.getMPG();
```

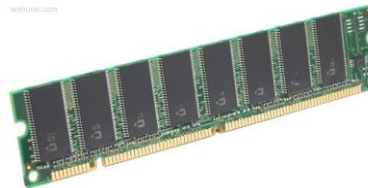
Assembly  
language:

```
get_mpg:
    pushq    %rbp
    movq     %rsp, %rbp
    ...
    popq     %rbp
    ret
```

Machine  
code:

```
0111010000011000
100011010000010000000010
1000100111000010
110000011111101000011111
```

Computer  
system:



Data & addressing  
Integers & floats  
Machine code & C  
x86 assembly  
programming  
**Procedures &  
stacks**  
Arrays & structs  
Memory & caches  
Processes  
Virtual memory  
Memory allocation  
Java vs. C

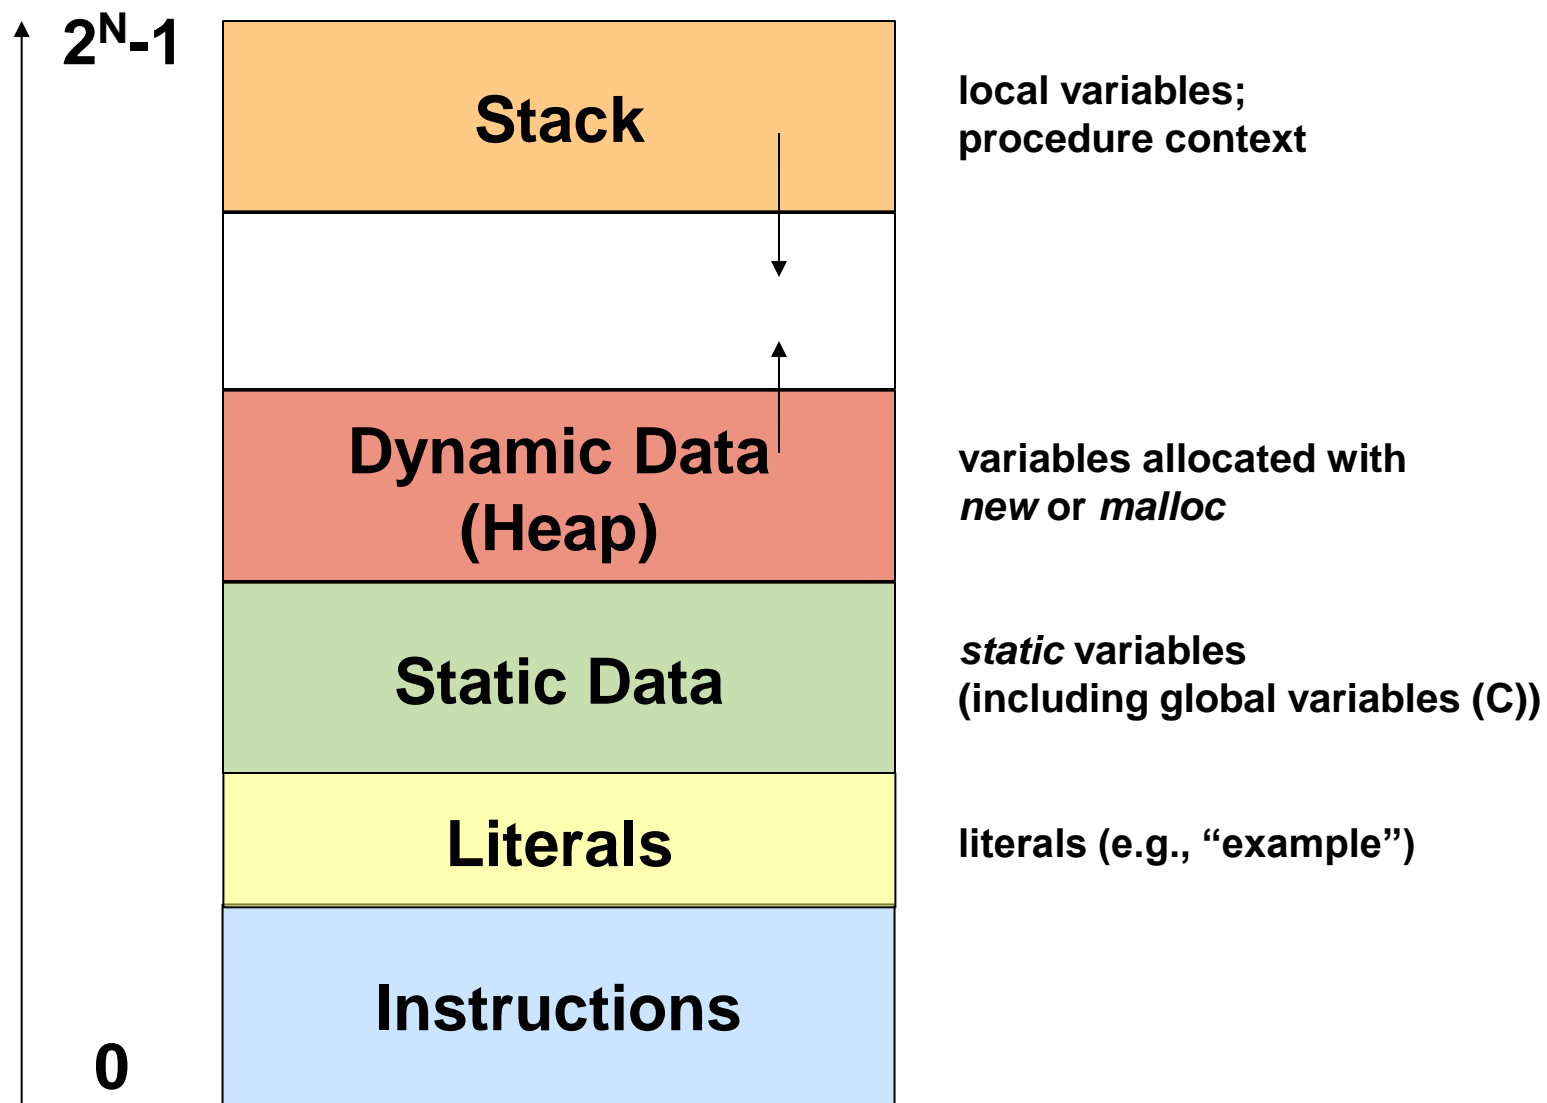
OS:



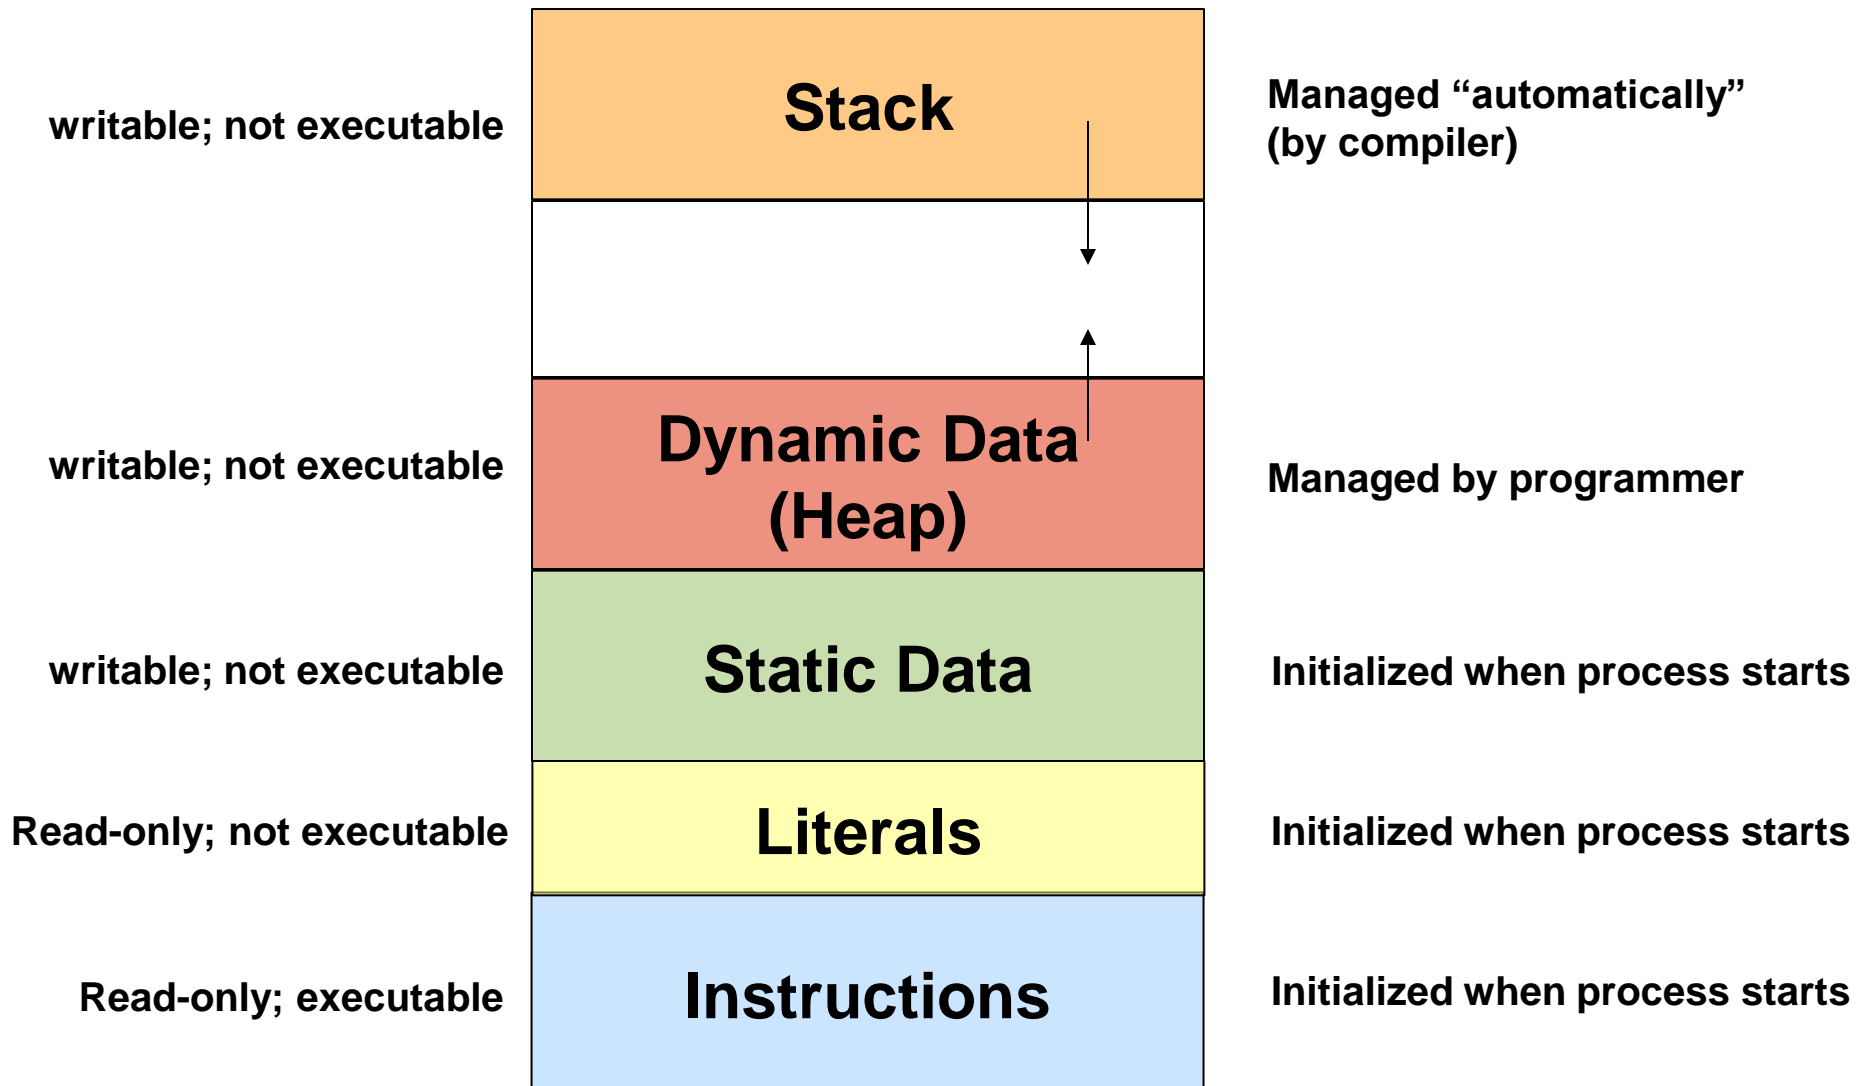
# Procedures and Call Stacks

- How do I pass arguments to a procedure?
  - How do I get a return value from a procedure?
  - Where do I put local variables?
  - When a function returns, how does it know where to return to?
- 
- To answer these questions, we need a *call stack* ...

# Memory Layout

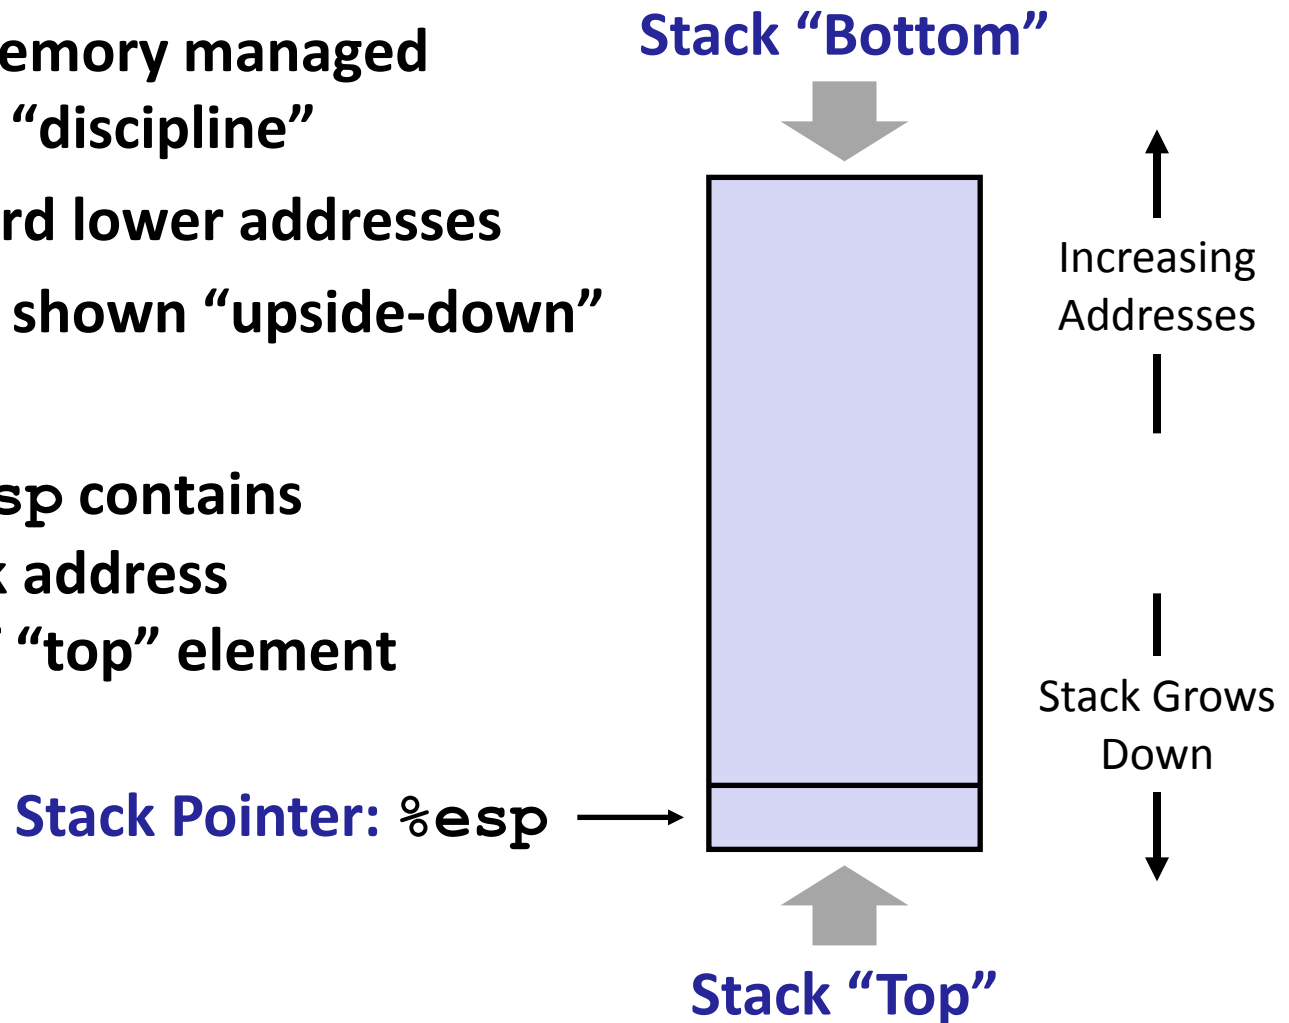


# Memory Layout



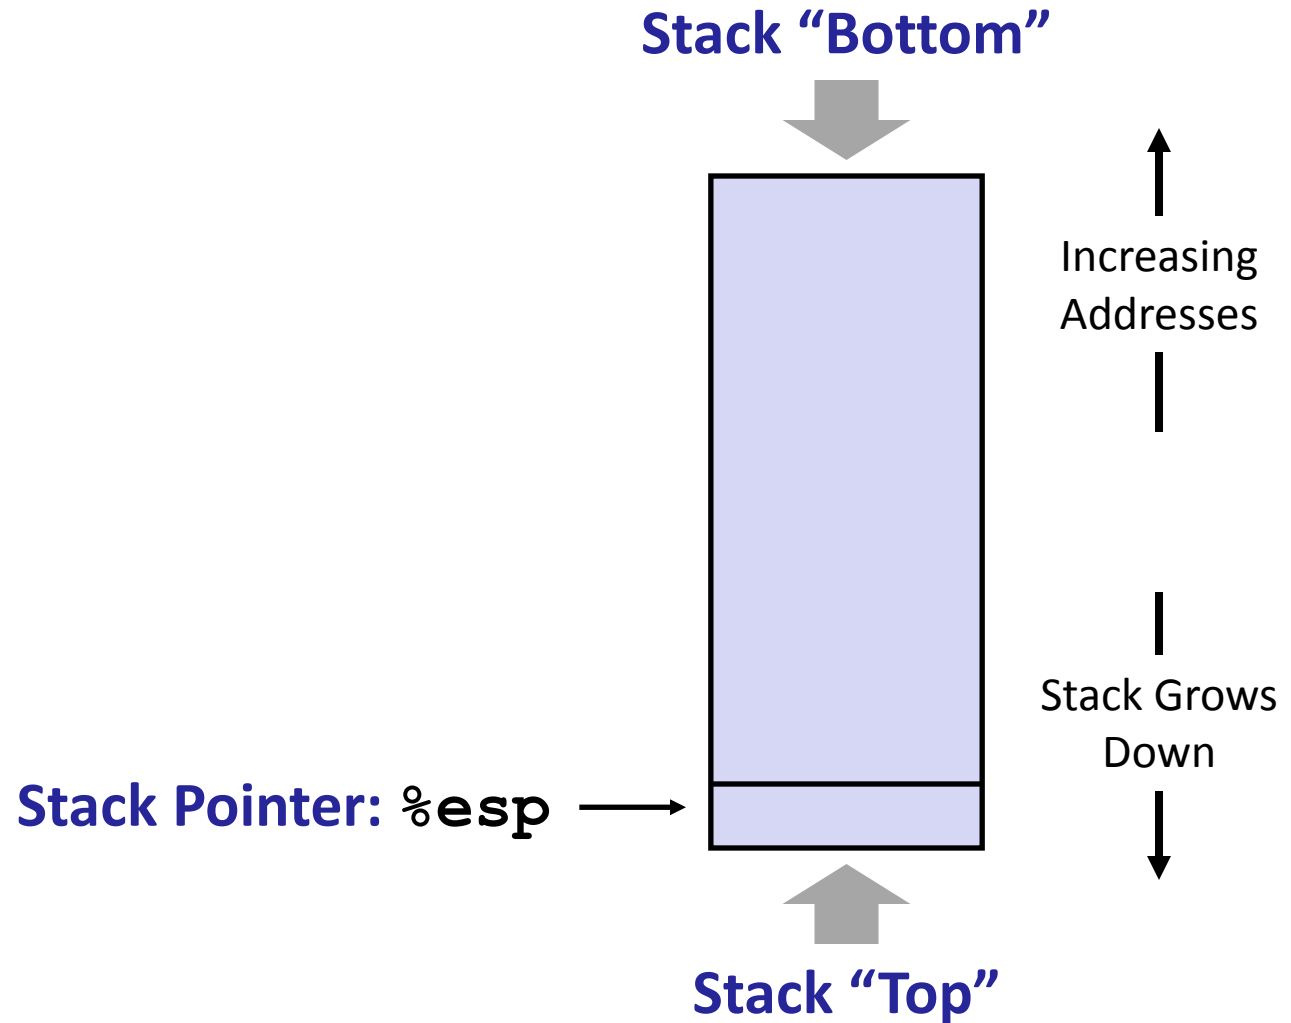
# IA32 Call Stack

- Region of memory managed with a stack “discipline”
- Grows toward lower addresses
- Customarily shown “upside-down”
- Register `%esp` contains lowest stack address = address of “top” element



# IA32 Call Stack: Push

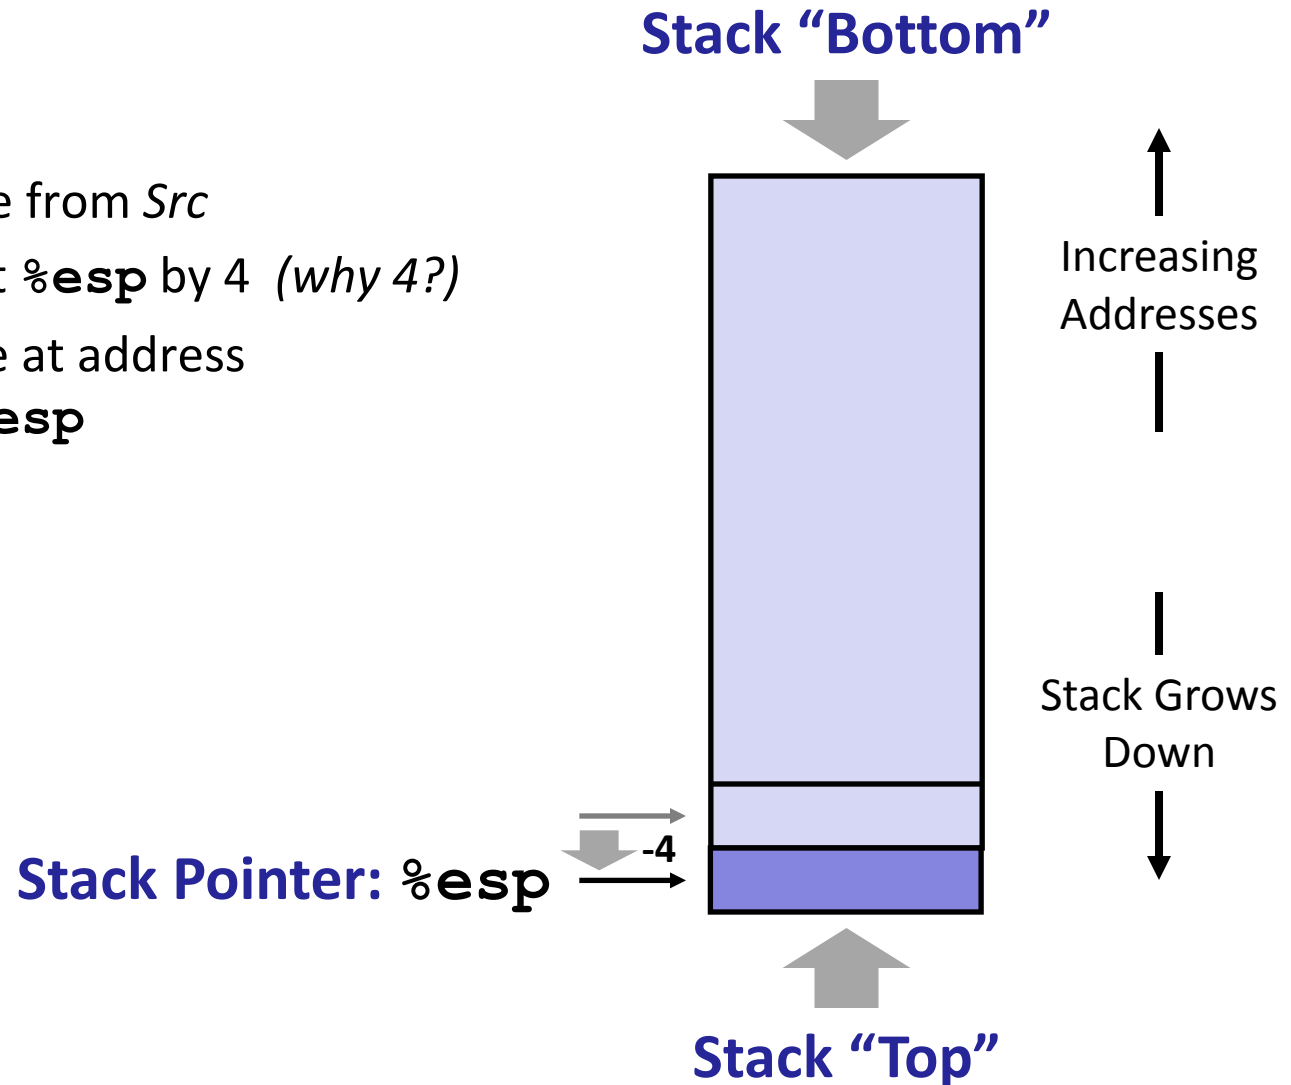
■ `pushl Src`



# IA32 Call Stack: Push

## ■ `pushl Src`

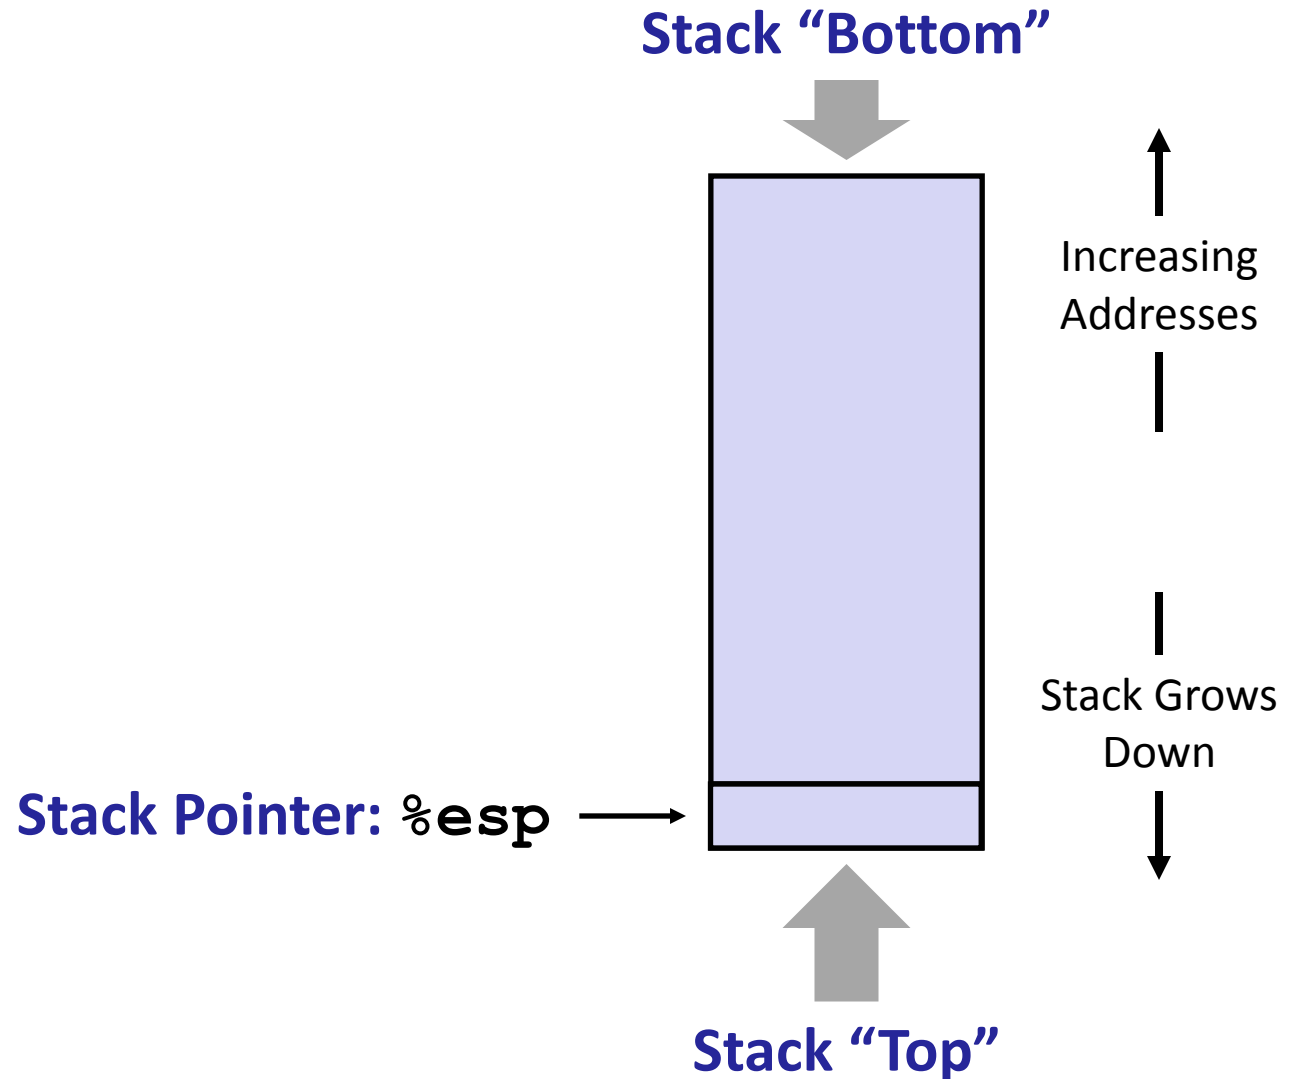
- Fetch value from *Src*
- Decrement `%esp` by 4 (*why 4?*)
- Store value at address given by `%esp`





# IA32 Call Stack: Pop

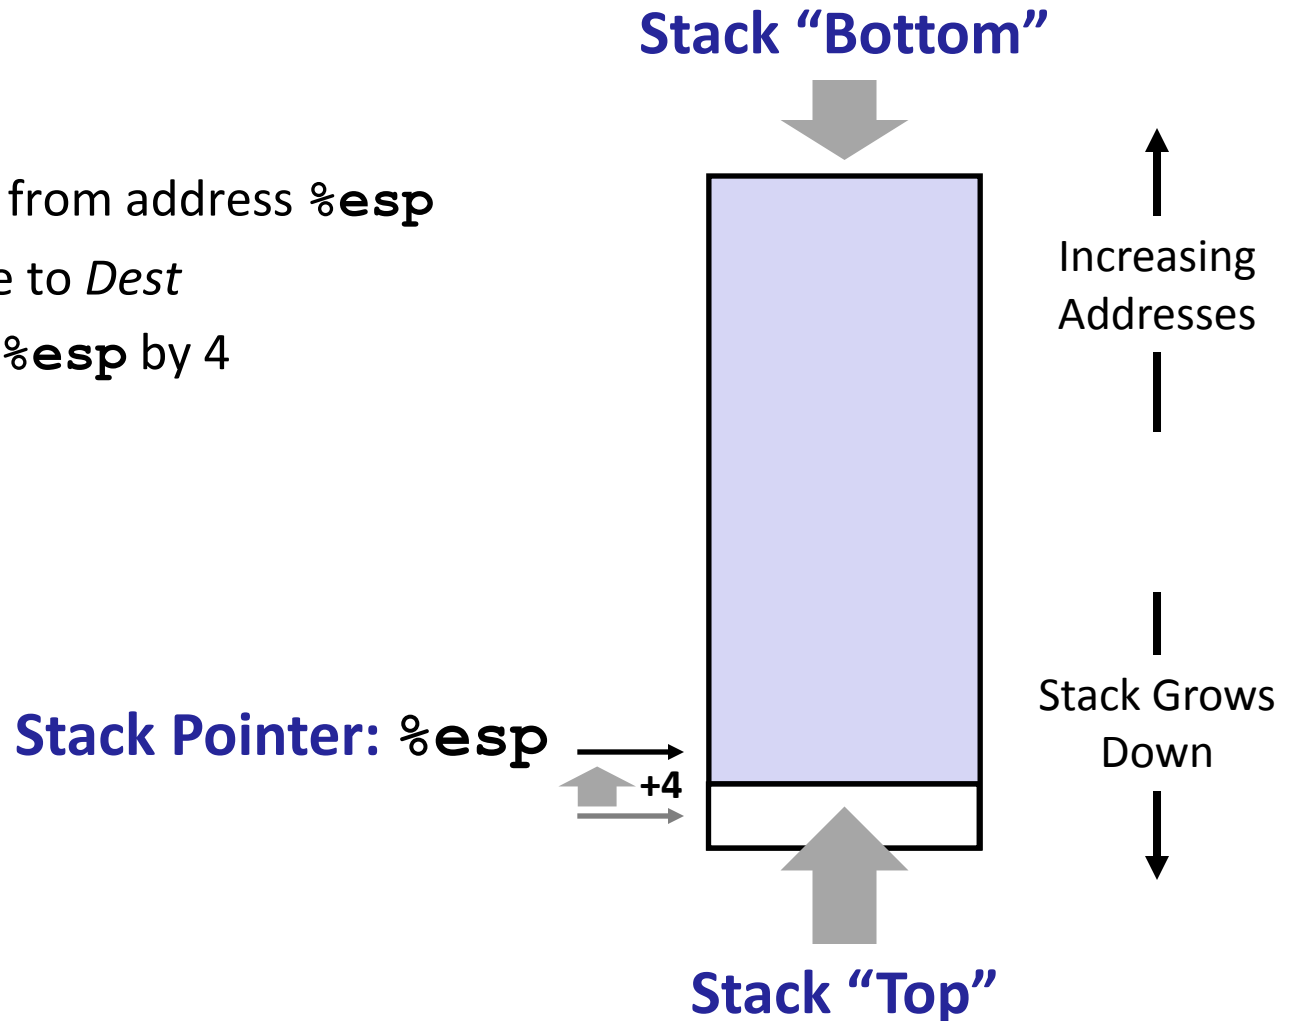
■ `popl Dest`



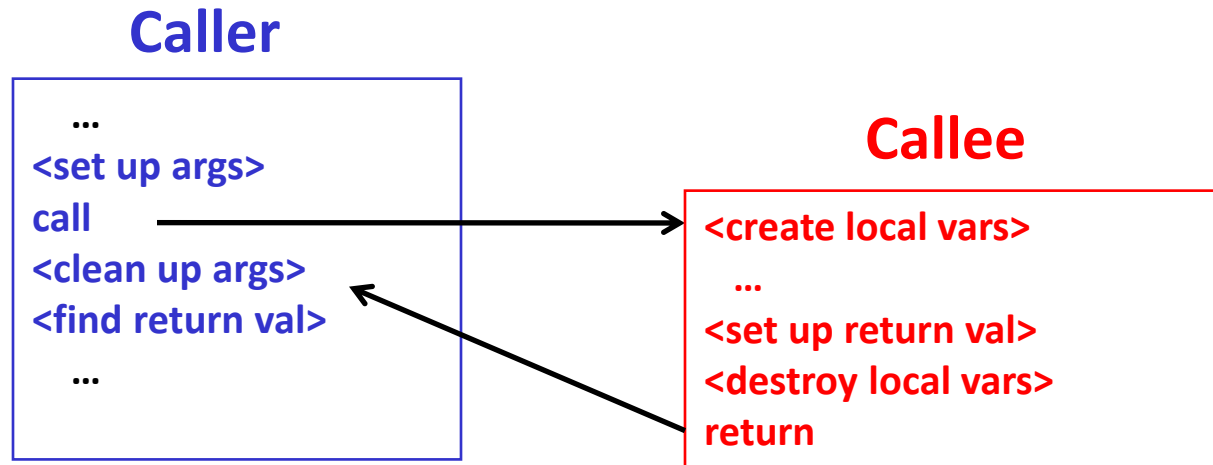
# IA32 Call Stack: Pop

## ■ `popl Dest`

- Load value from address `%esp`
- Write value to `Dest`
- Increment `%esp` by 4

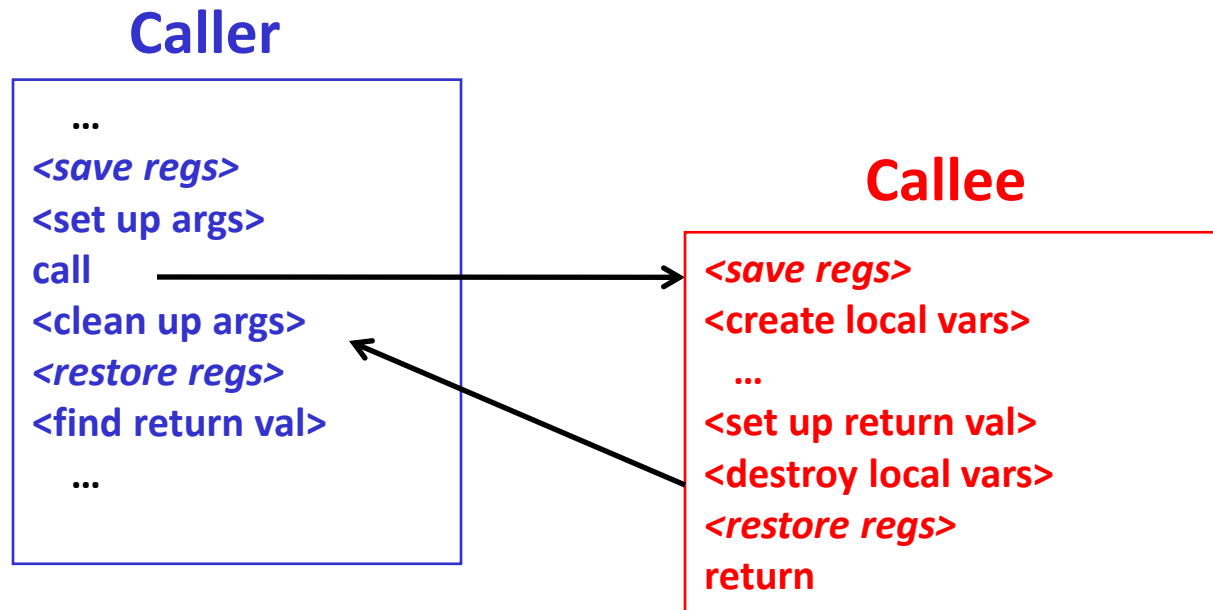


# Procedure Call Overview



- **Callee** must know where to find args
- **Callee** must know where to find “return address”
- **Caller** must know where to find return val
- **Caller** and **Callee** run on same CPU → use the same registers
  - **Caller** might need to save registers that **Callee** might use
  - **Callee** might need to save registers that **Caller** has used

# Procedure Call Overview



- The convention of where to leave/find things is called the procedure call linkage
  - Details vary between systems
  - We will see the convention for IA32/Linux in detail
  - What could happen if our program didn't follow these conventions?

# Procedure Control Flow

- Use stack to support procedure call and return
- **Procedure call:** `call label`
  - Push return address on stack
  - Jump to *label*

# Procedure Control Flow

- Use stack to support procedure call and return

- **Procedure call:** `call label`

- Push return address on stack
  - Jump to *label*

- **Return address:**

- Address of instruction after `call`
  - Example from disassembly:

804854e:	e8 3d 06 00 00	call	8048b90 <main>
8048553:	50	pushl	%eax

- Return address = `0x8048553`

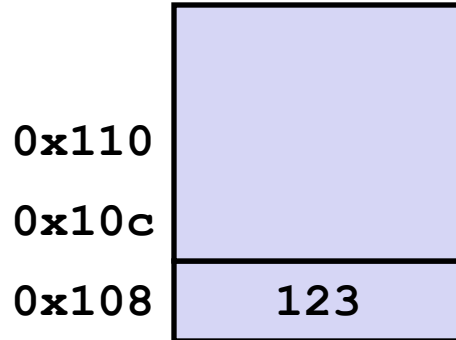
- **Procedure return:** `ret`

- Pop return address from stack
  - Jump to address

# Procedure Call Example

804854e:	e8 3d 06 00 00	call	8048b90 <main>
8048553:	50	pushl	%eax

**call 8048b90**



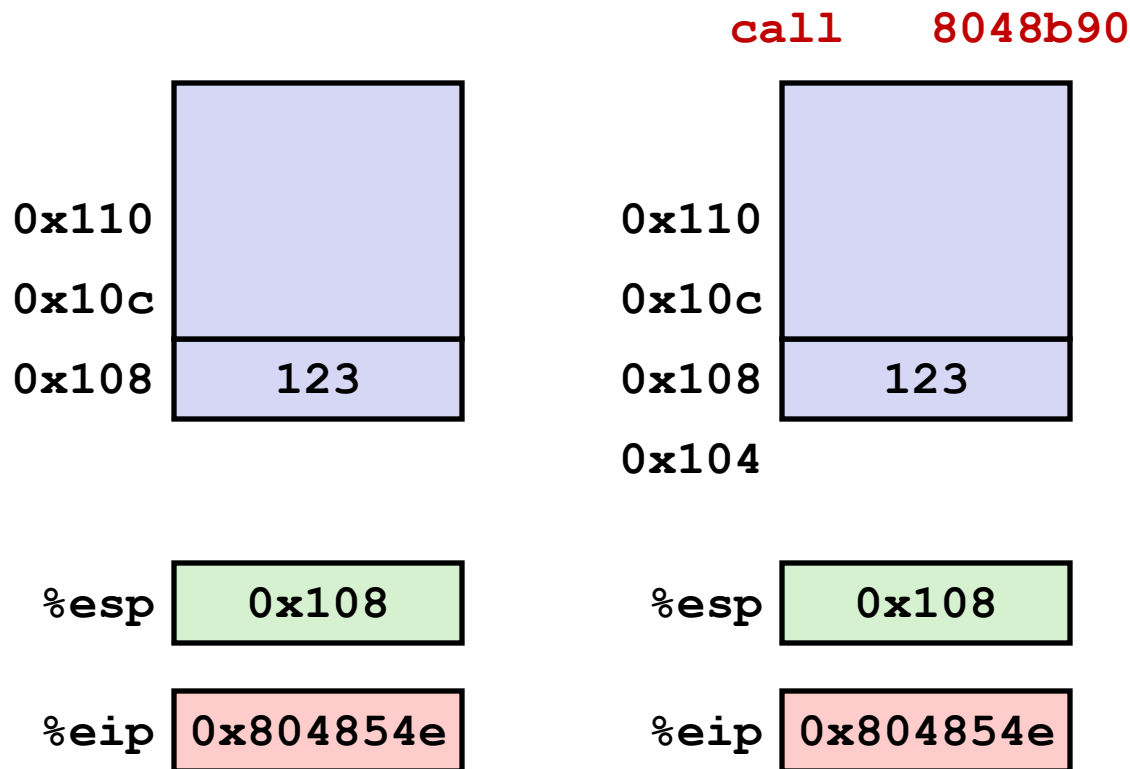
%esp 0x108

%eip 0x804854e

*%eip: program counter*

# Procedure Call Example

804854e:	e8 3d 06 00 00	call	8048b90 <main>
8048553:	50	pushl	%eax

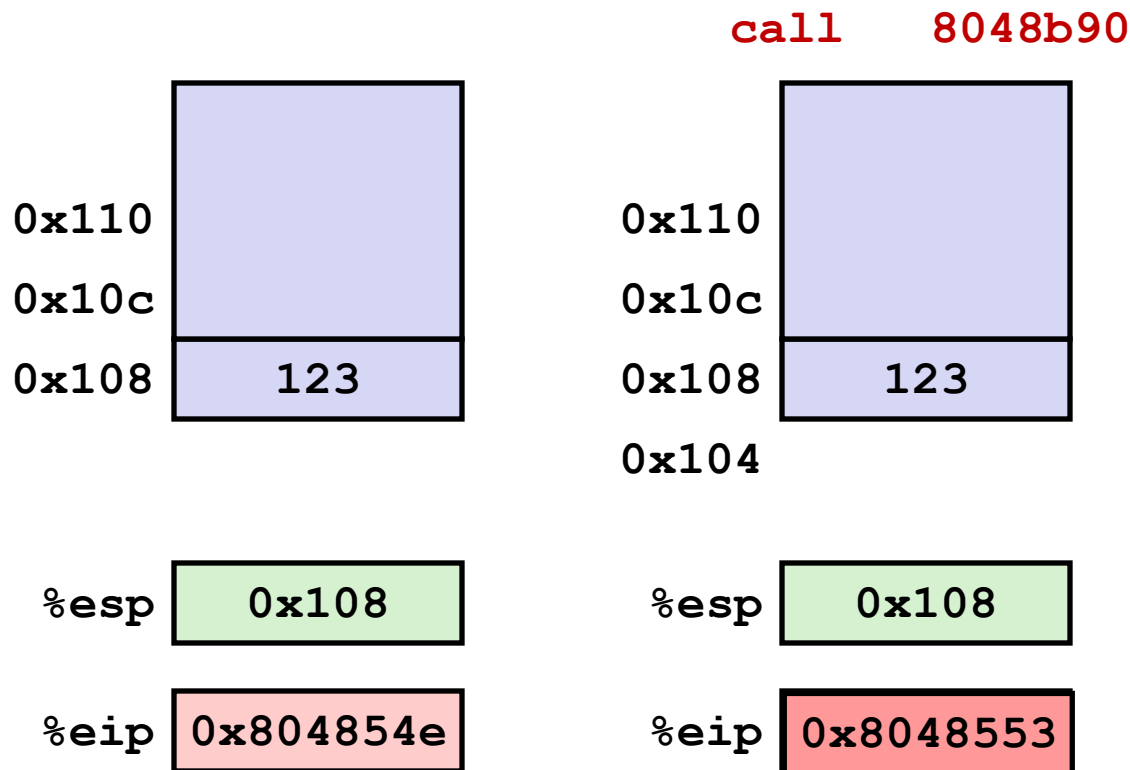


*%eip: program counter*



# Procedure Call Example

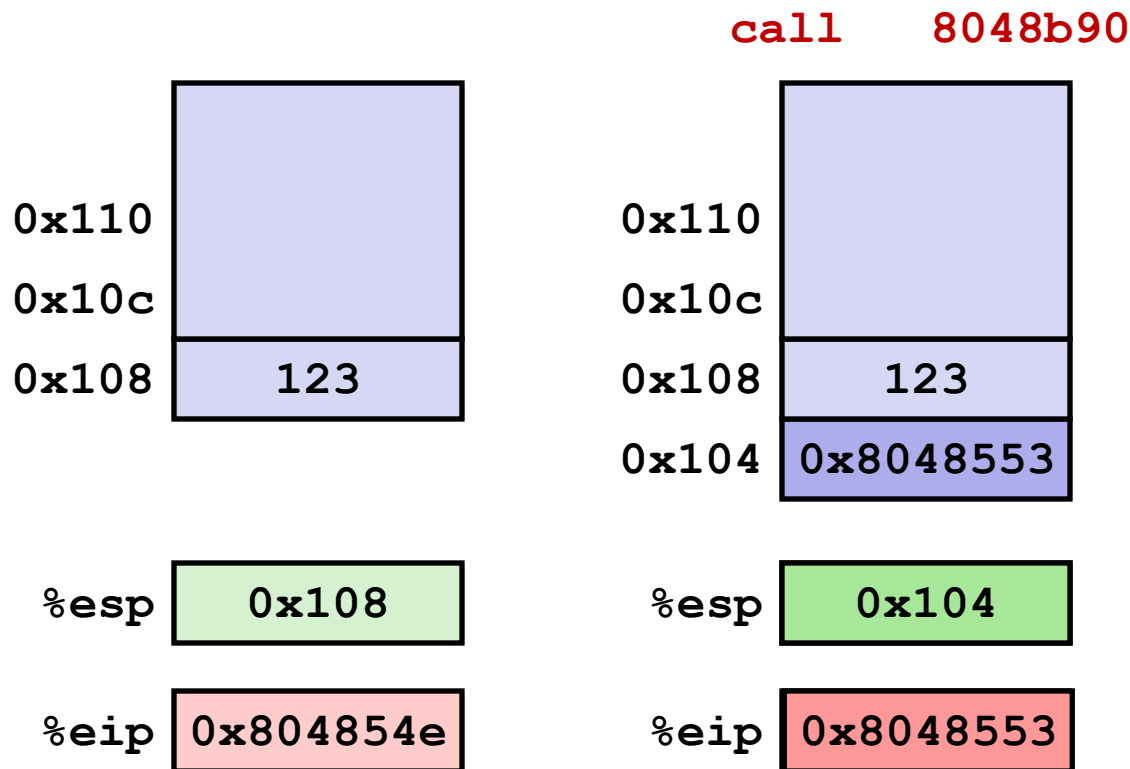
804854e:	e8 3d 06 00 00	call	8048b90 <main>
8048553:	50	pushl	%eax



*%eip: program counter*

# Procedure Call Example

804854e:	e8 3d 06 00 00	call	8048b90 <main>
8048553:	50	pushl	%eax

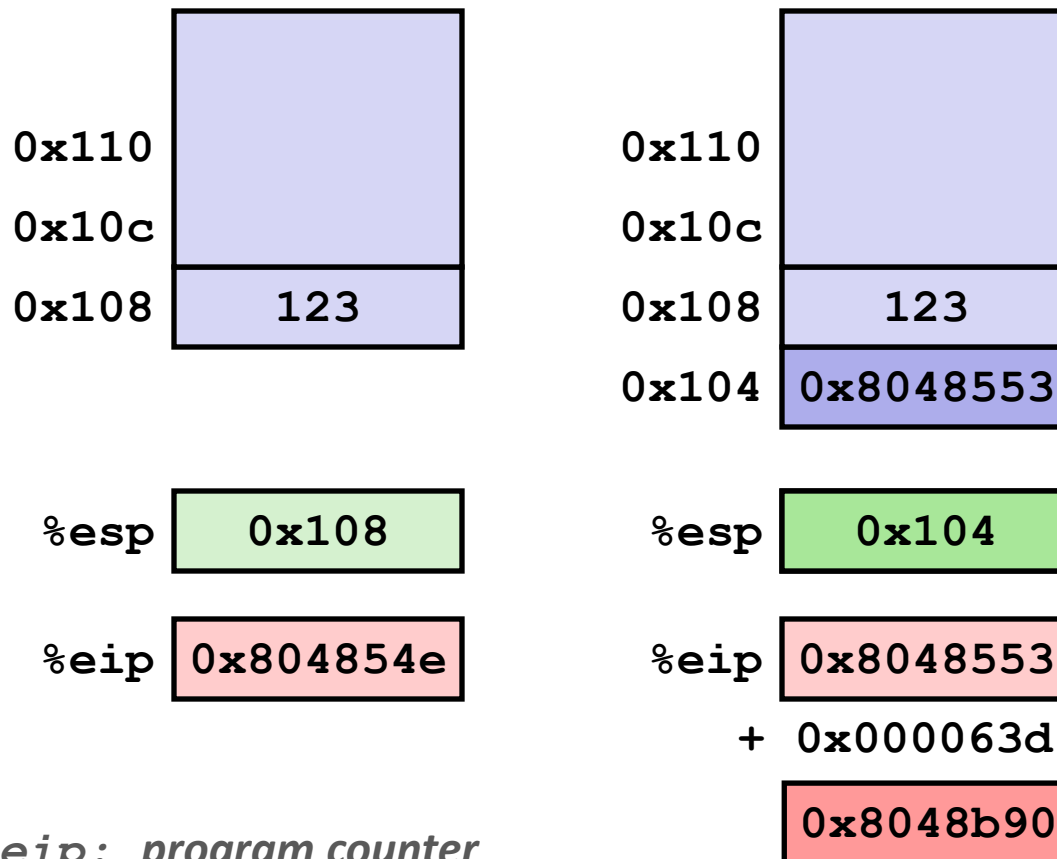


*%eip: program counter*

# Procedure Call Example

804854e:	e8 3d 06 00 00	call	8048b90 <main>
8048553:	50	pushl	%eax

**call 8048b90**

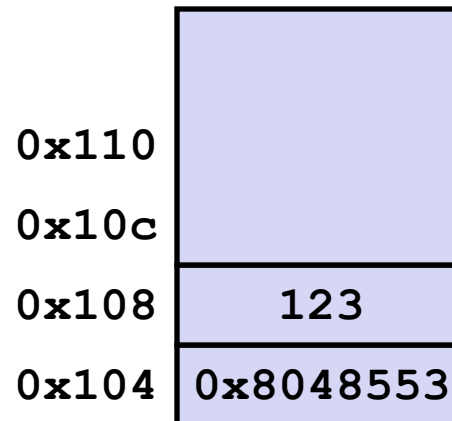


*%eip: program counter*

# Procedure Return Example

8048591:	c3	ret
----------	----	-----

ret



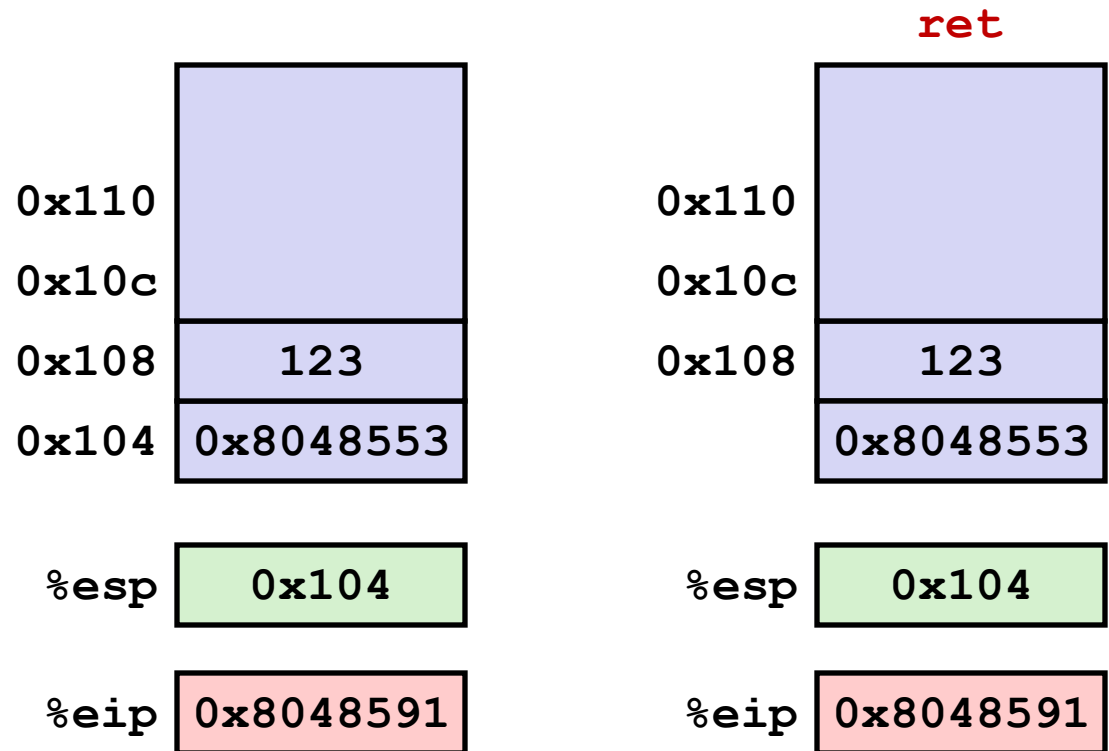
%esp	0x104
------	-------

%eip	0x8048591
------	-----------

*%eip: program counter*

# Procedure Return Example

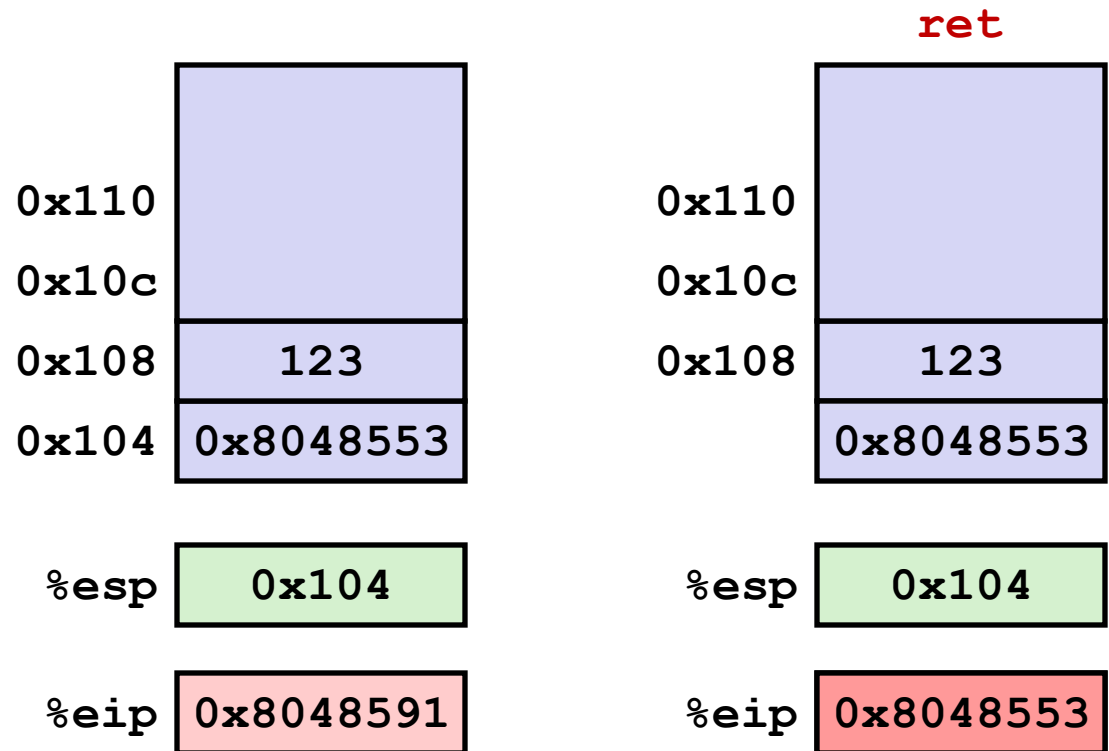
8048591:	c3	ret
----------	----	-----



*%eip: program counter*

# Procedure Return Example

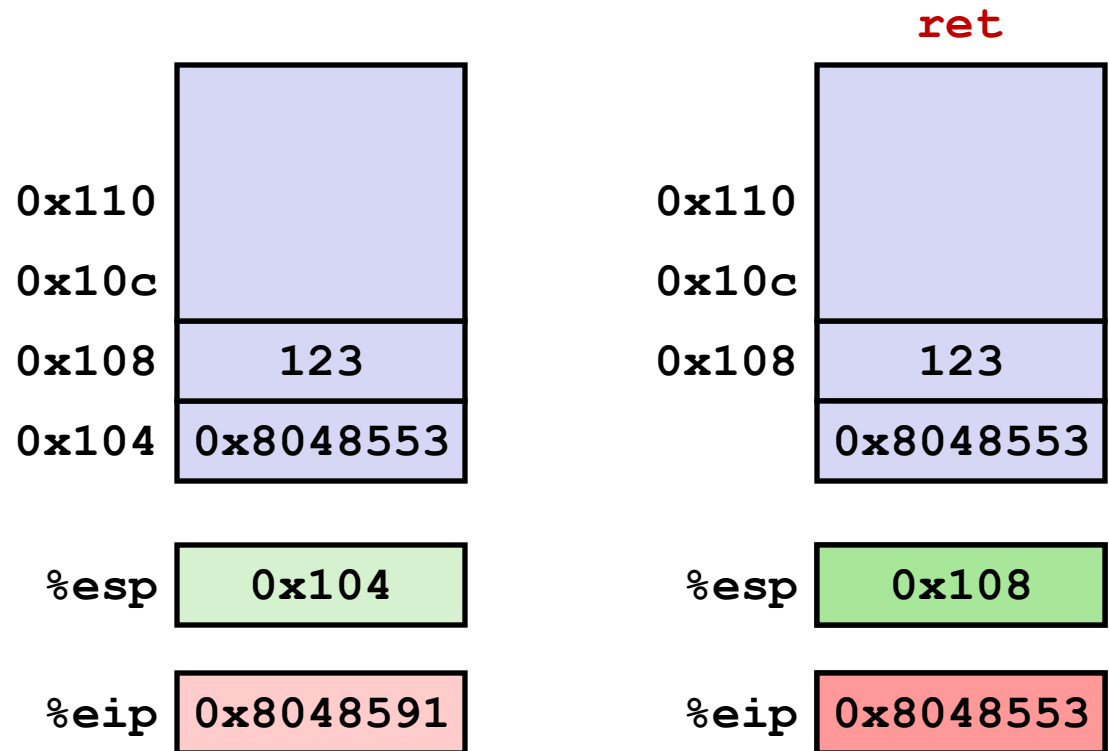
8048591:	c3	ret
----------	----	-----



*%eip: program counter*

# Procedure Return Example

8048591:	c3	ret
----------	----	-----



*%eip: program counter*

# Stack-Based Languages

## ■ Languages that support recursion

- e.g., C, Pascal, Java
- Code must be re-entrant
  - Multiple simultaneous instantiations of single procedure
- Need some place to store state of each instantiation
  - Arguments
  - Local variables
  - Return pointer

## ■ Stack discipline

- State for a given procedure needed for a limited time
  - Starting from when it is called to when it returns
- Callee always returns before caller does

## ■ Stack allocated in frames

- State for a single procedure instantiation



# Call Chain Example

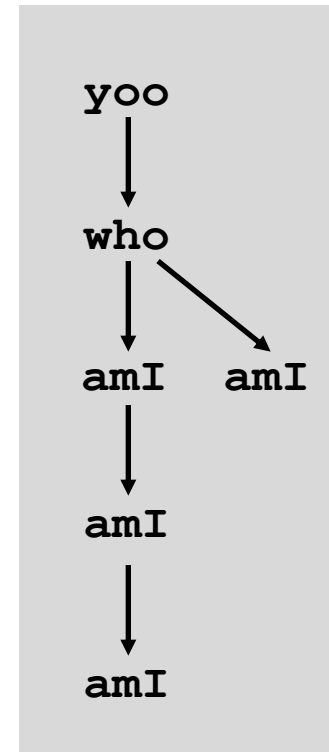
```
yoo (...)  
{  
  .  
  .  
  who () ;  
  .  
  .  
}
```

```
who (...)  
{  
  . . .  
  amI () ;  
  . . .  
  amI () ;  
  . . .  
}
```

```
amI (...)  
{  
  .  
  .  
  amI () ;  
  .  
  .  
}
```

Procedure `amI` is recursive  
(calls itself)

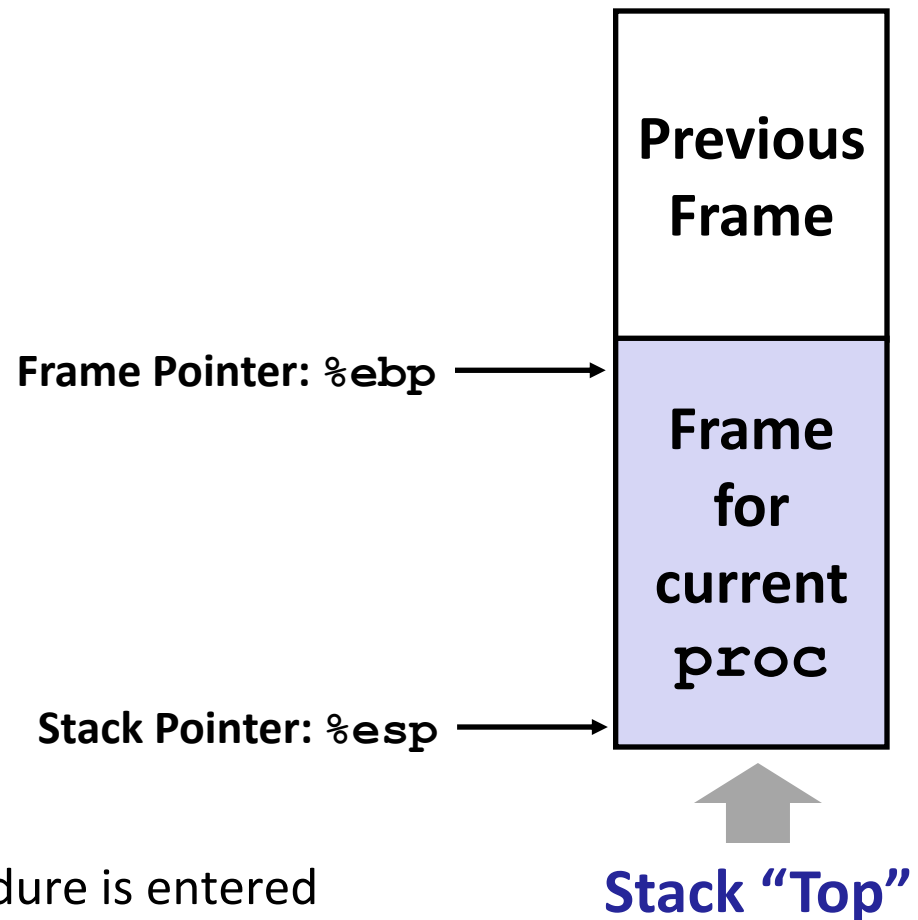
## Example Call Chain



# Stack Frames

## ■ Contents

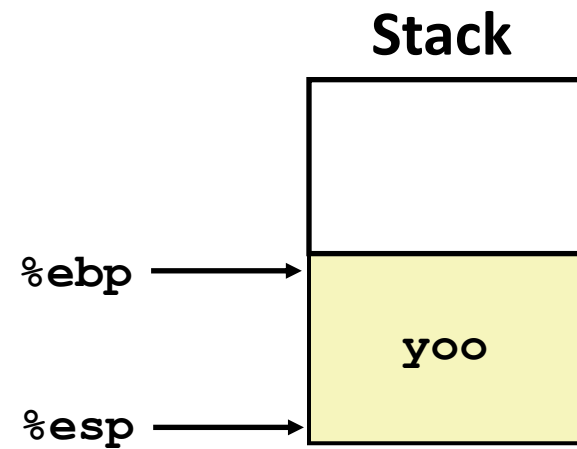
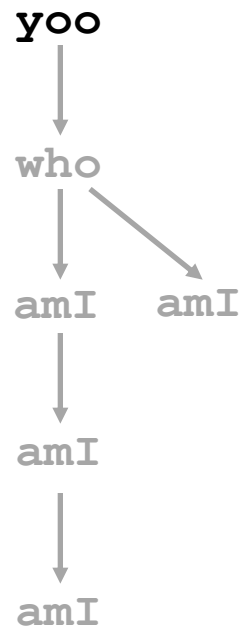
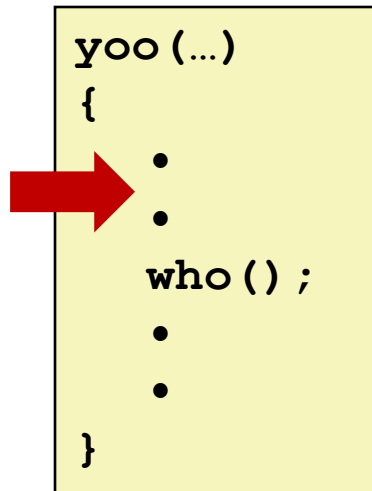
- Local variables
- Function arguments
- Return information
- Temporary space



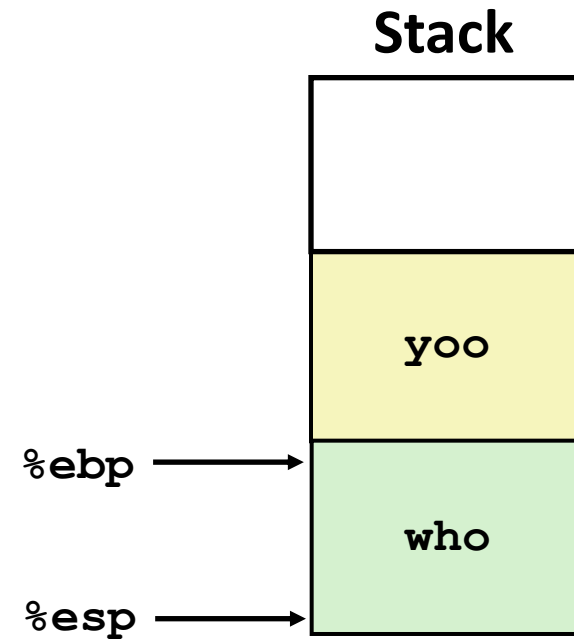
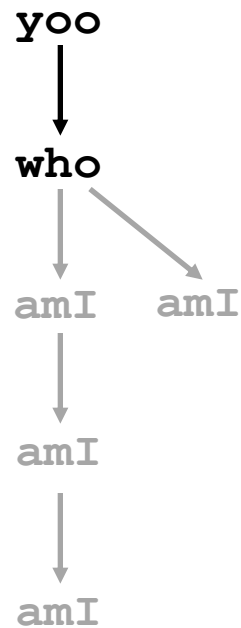
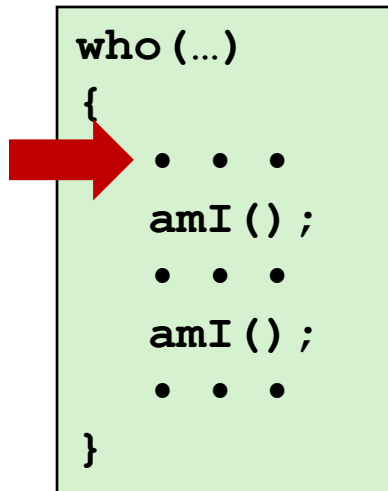
## ■ Management

- Space allocated when procedure is entered
  - "Set-up" code
- Space deallocated upon return
  - "Finish" code

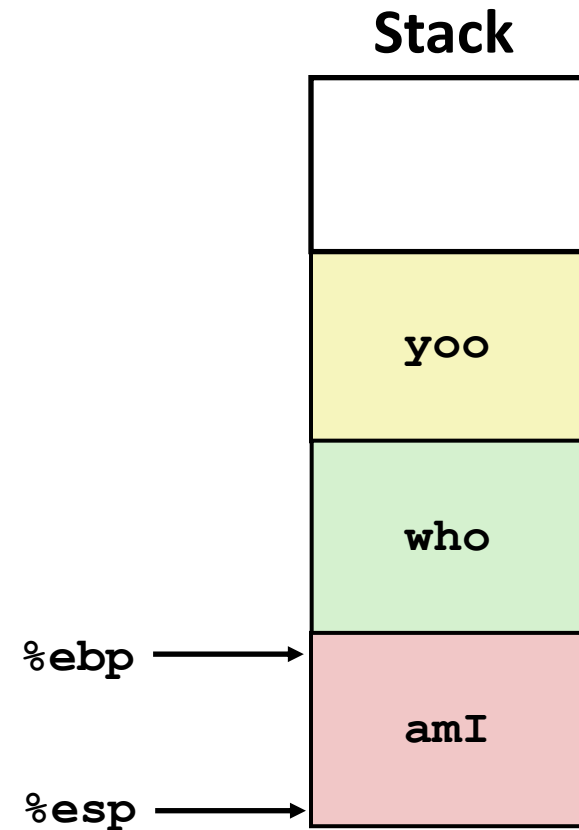
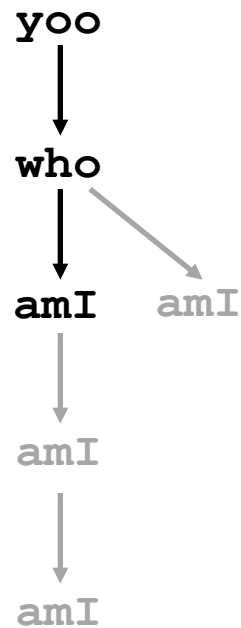
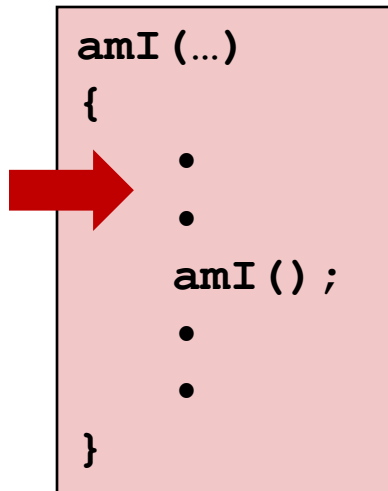
# Example



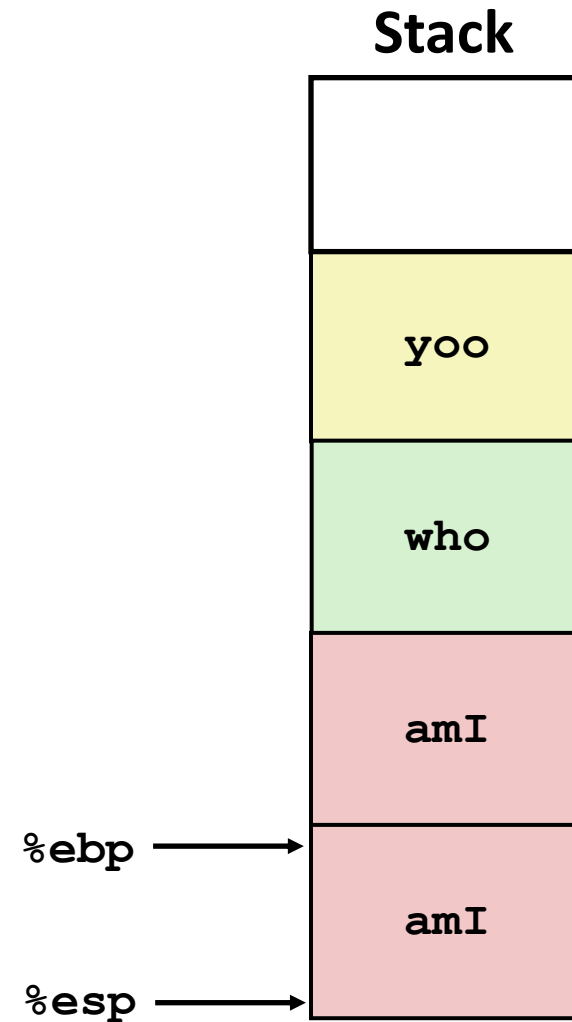
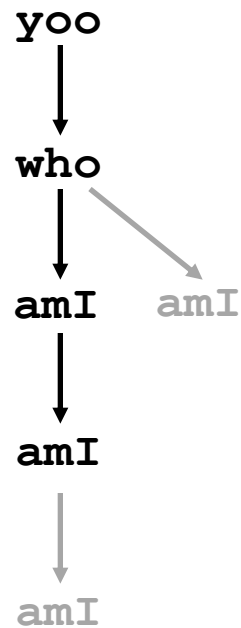
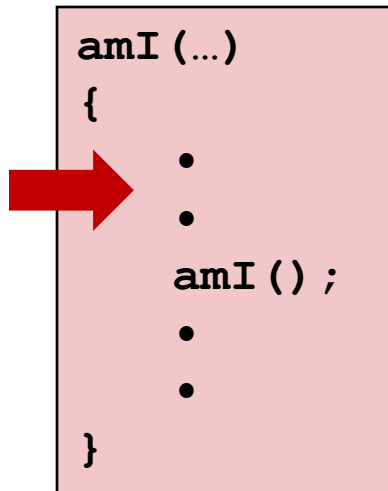
# Example



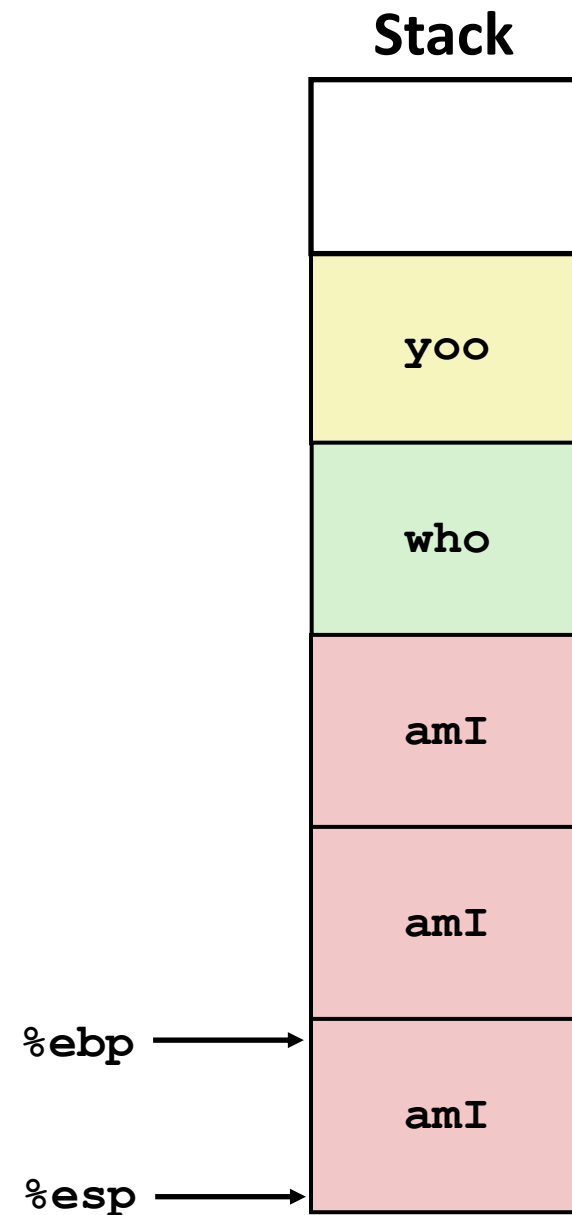
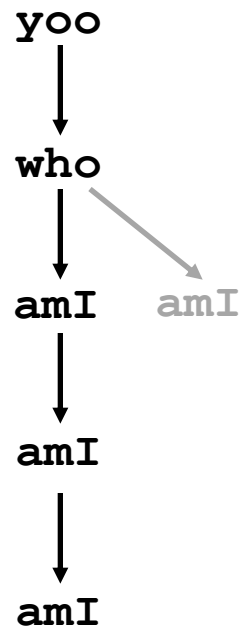
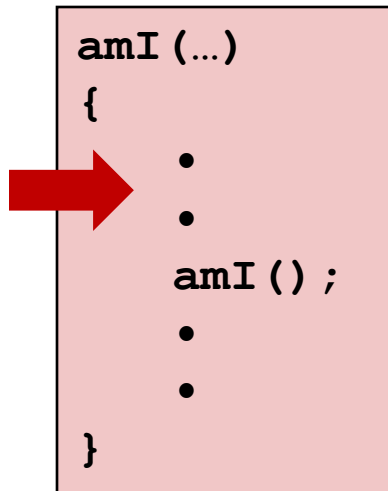
# Example



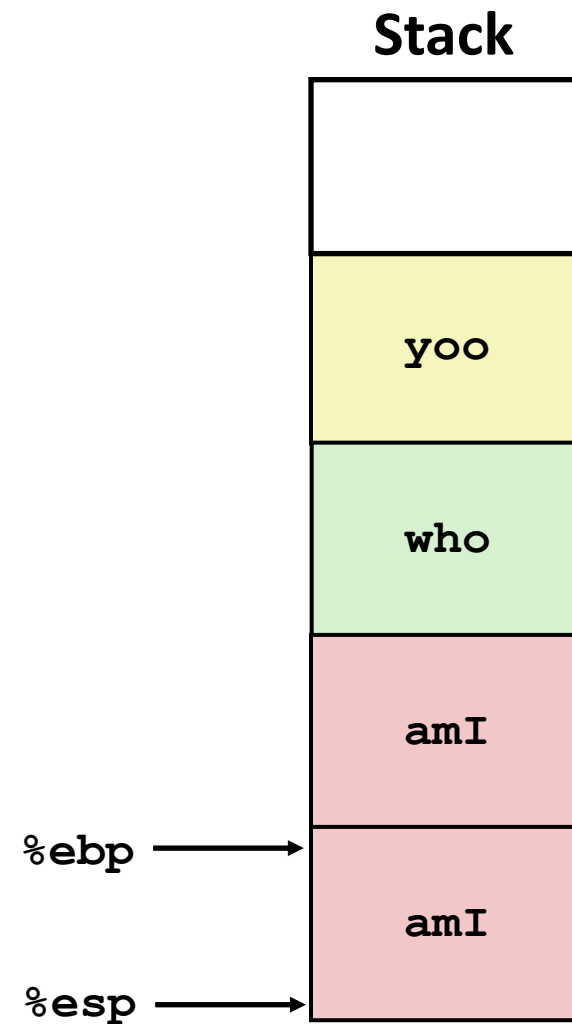
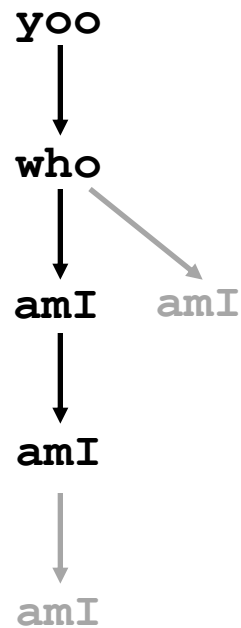
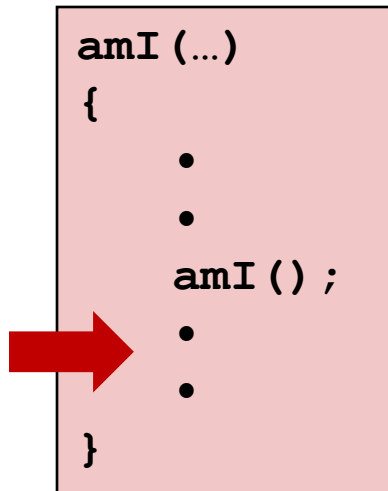
# Example



# Example

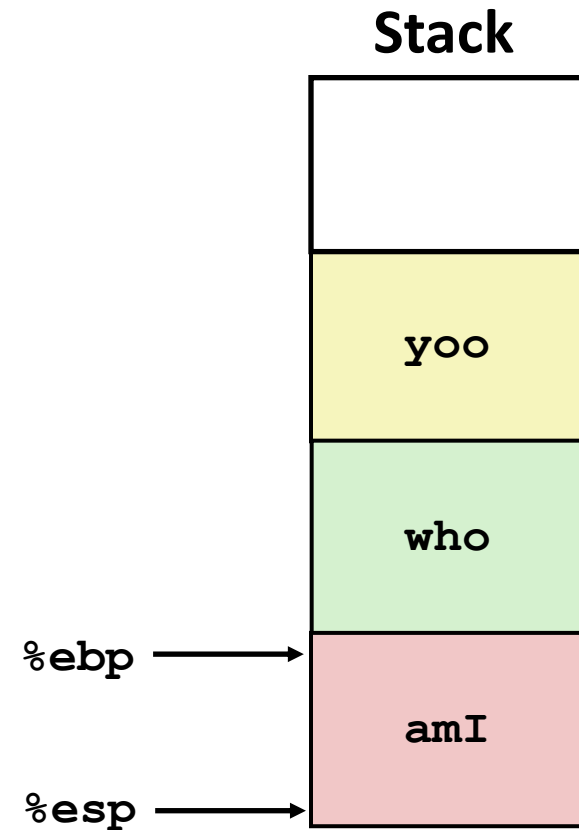
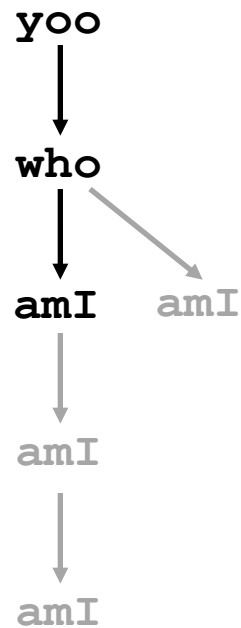
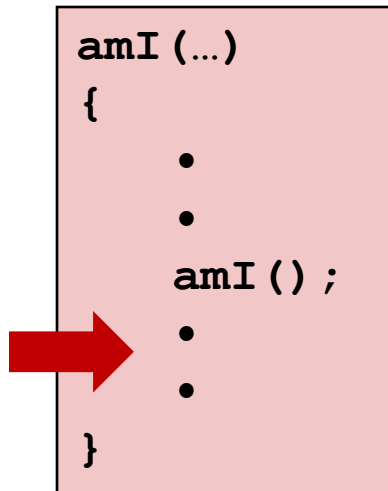


# Example

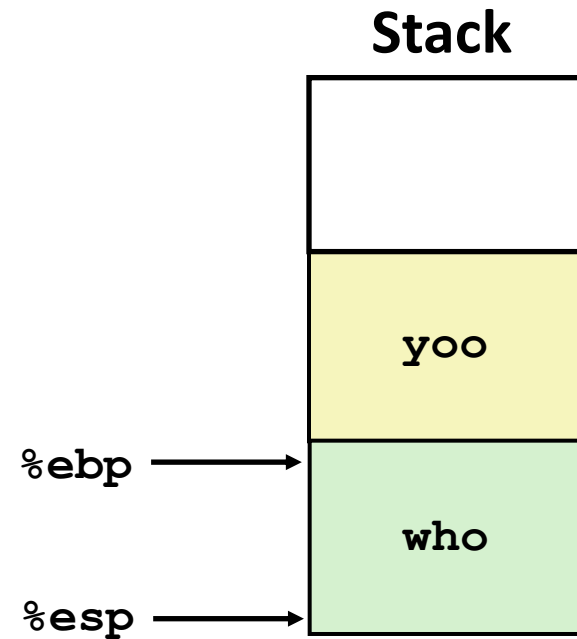
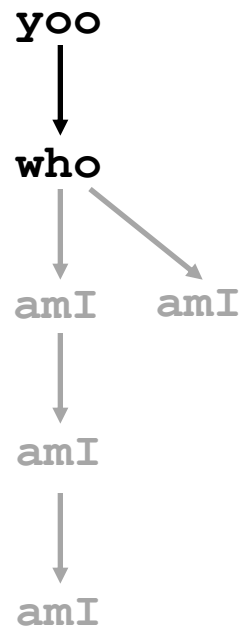
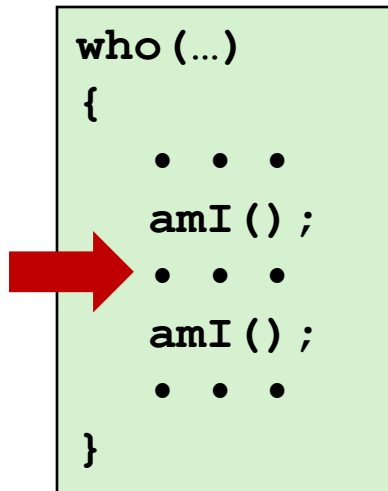




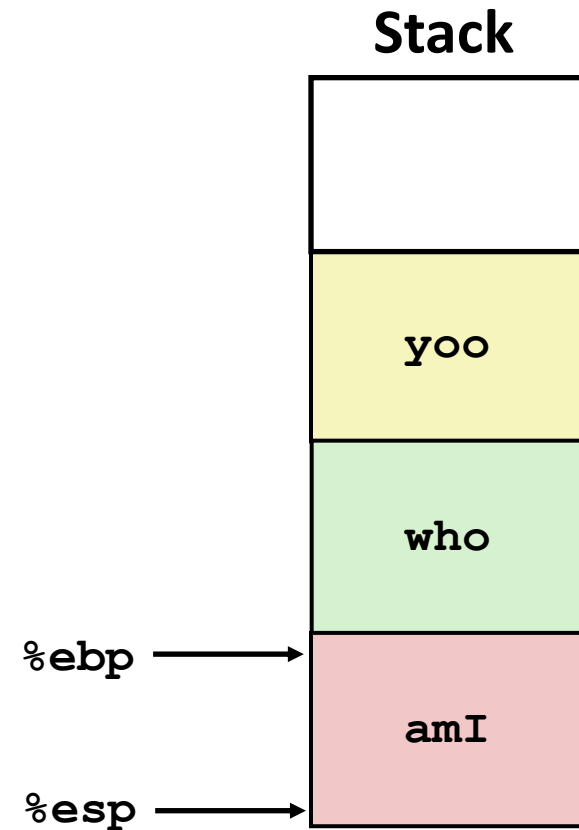
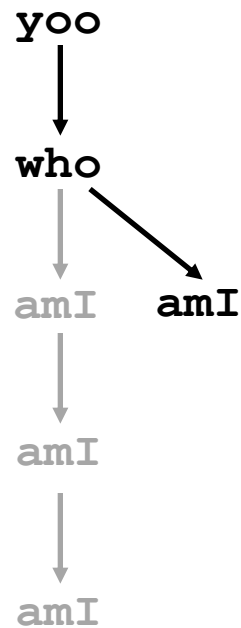
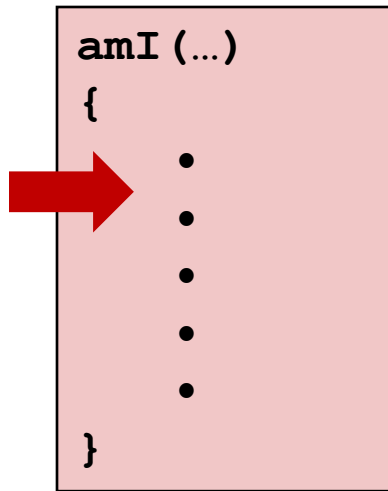
# Example



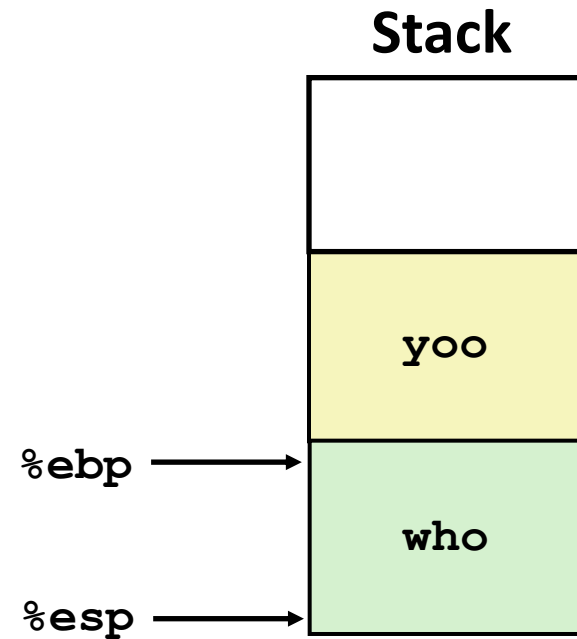
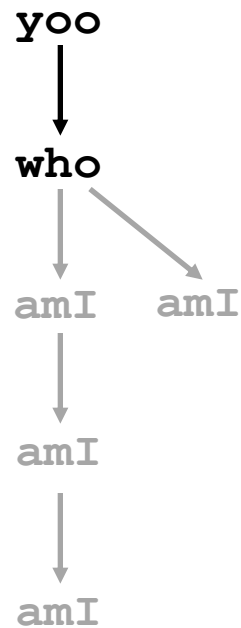
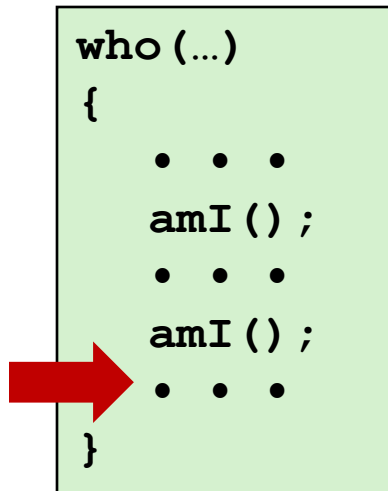
# Example



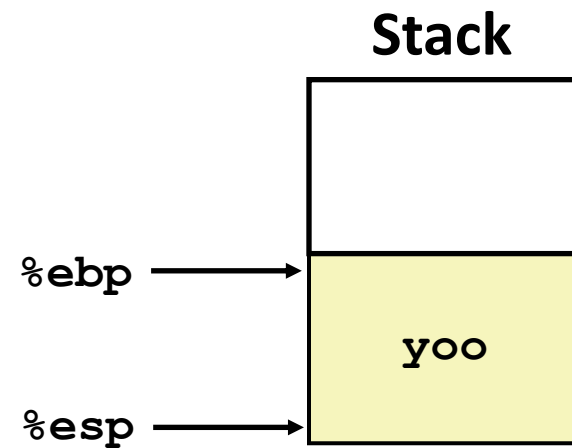
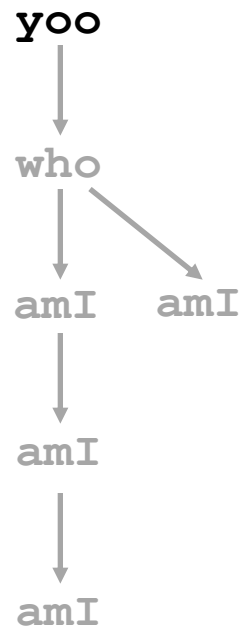
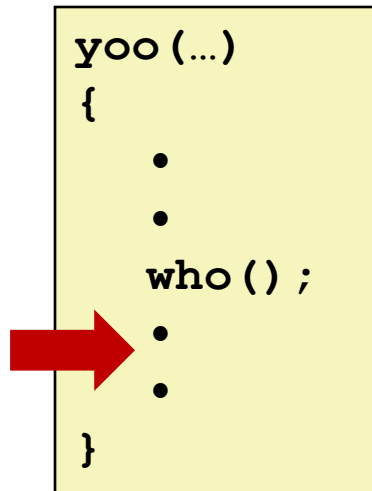
# Example



# Example



# Example



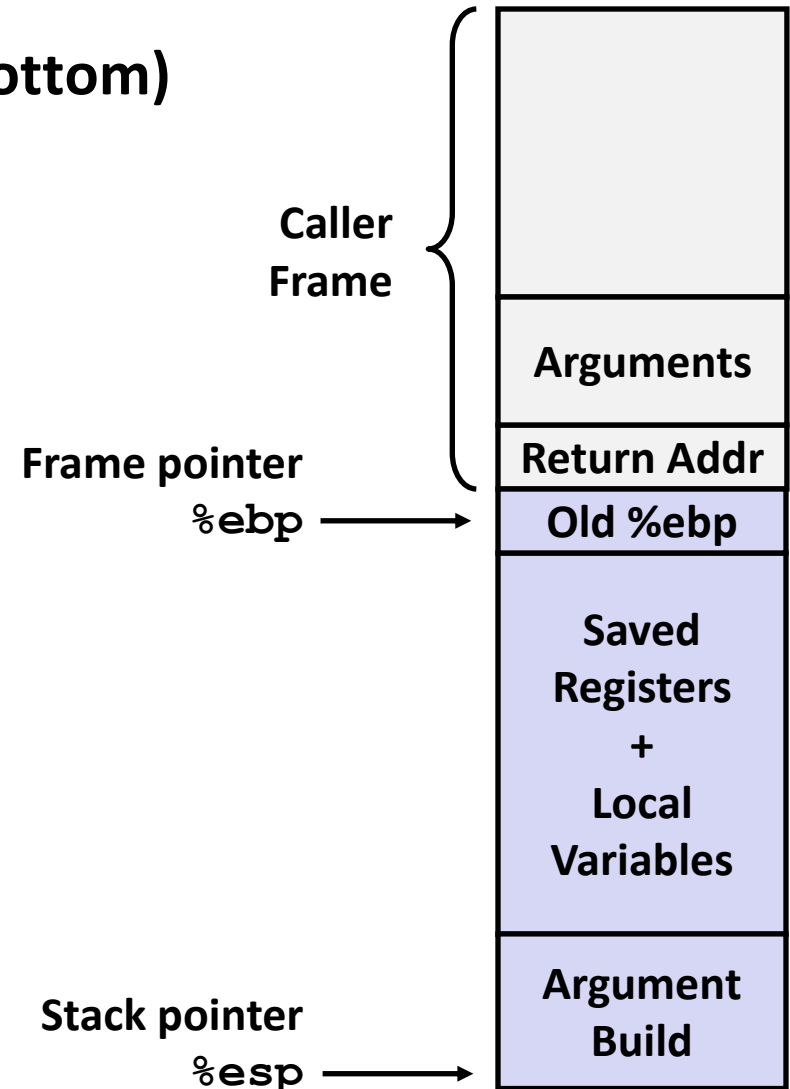
# IA32/Linux Stack Frame

## ■ Current Stack Frame (“Top” to Bottom)

- “Argument build” area (parameters for function about to be called)
- Local variables (if can’t be kept in registers)
- Saved register context (when reusing registers)
- Old frame pointer (for caller)

## ■ Caller’s Stack Frame

- Return address
  - Pushed by **call** instruction
- Arguments for this call



# Revisiting swap

```
int zip1 = 15213;
int zip2 = 98195;

void call_swap()
{
    swap(&zip1, &zip2);
}
```

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

# Revisiting swap

```
int zip1 = 15213;
int zip2 = 98195;

void call_swap()
{
    swap(&zip1, &zip2);
}
```

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

## Calling swap from call\_swap

```
call_swap:
    . . .
    pushl $zip2    # Global Var
    pushl $zip1    # Global Var
    call swap
    . . .
```



# Revisiting swap

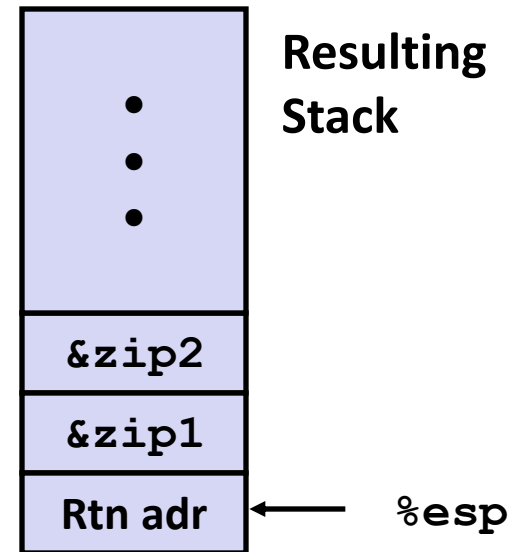
```
int zip1 = 15213;
int zip2 = 98195;

void call_swap()
{
    swap(&zip1, &zip2);
}
```

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

## Calling swap from call\_swap

```
call_swap:
    . . .
    pushl $zip2    # Global Var
    pushl $zip1    # Global Var
    call swap
    . . .
```



# Revisiting swap

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

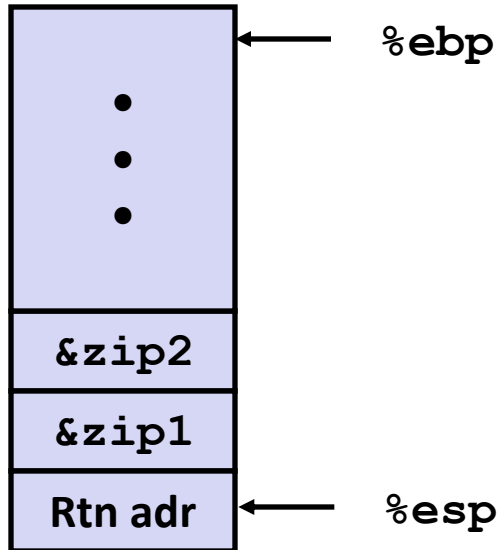
```
    pushl %ebp
    movl %esp,%ebp
    pushl %ebx
    } Set Up

    movl 12(%ebp),%ecx
    movl 8(%ebp),%edx
    movl (%ecx),%eax
    movl (%edx),%ebx
    movl %eax, (%edx)
    movl %ebx, (%ecx)
    } Body

    movl -4(%ebp),%ebx
    movl %ebp,%esp
    popl %ebp
    ret
    } Finish
```

# swap Setup #1

## Entering Stack

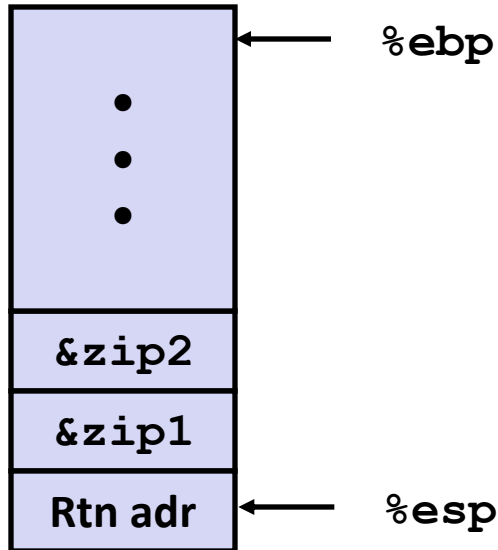


## Resulting Stack?

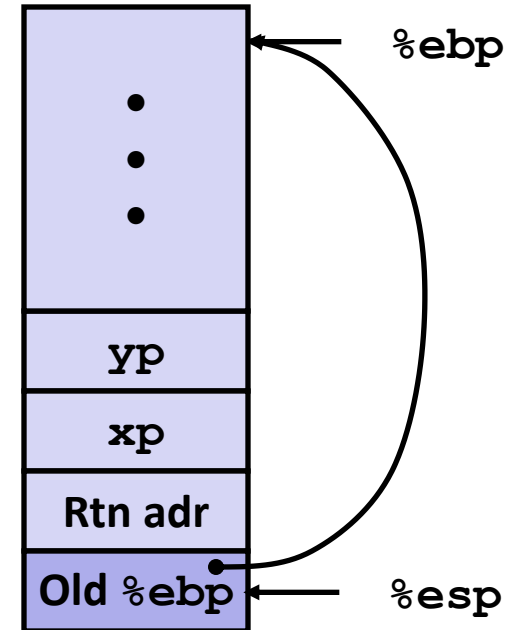
```
swap:  
    pushl %ebp  
    movl %esp,%ebp  
    pushl %ebx
```

# swap Setup #1

## Entering Stack



## Resulting Stack

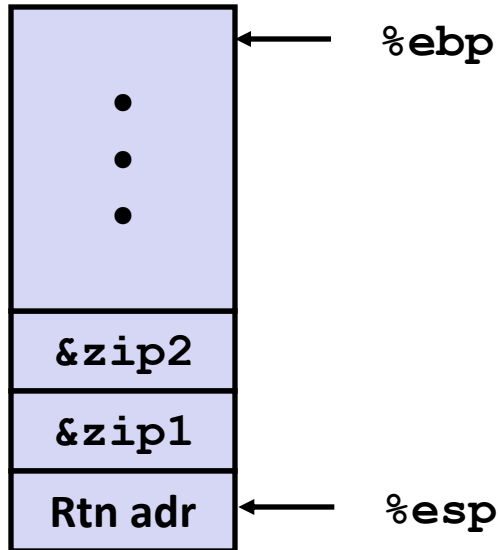


swap:

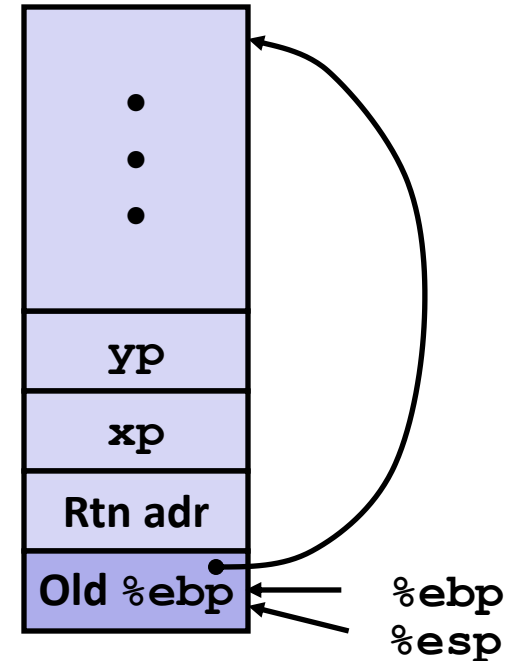
```
pushl %ebp  
movl %esp,%ebp  
pushl %ebx
```

# swap Setup #2

## Entering Stack



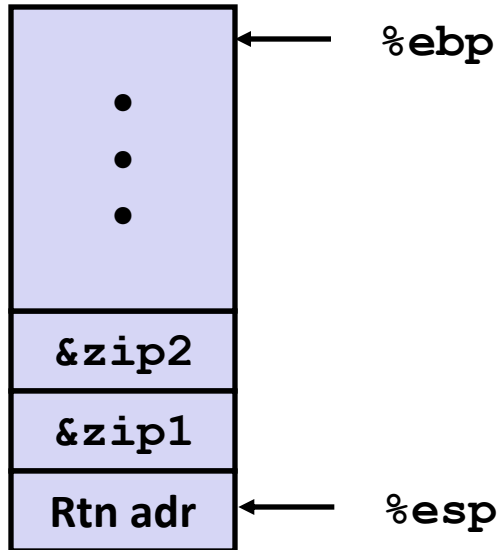
## Resulting Stack



```
swap:  
    pushl %ebp  
    movl %esp,%ebp  
    pushl %ebx
```

# swap Setup #3

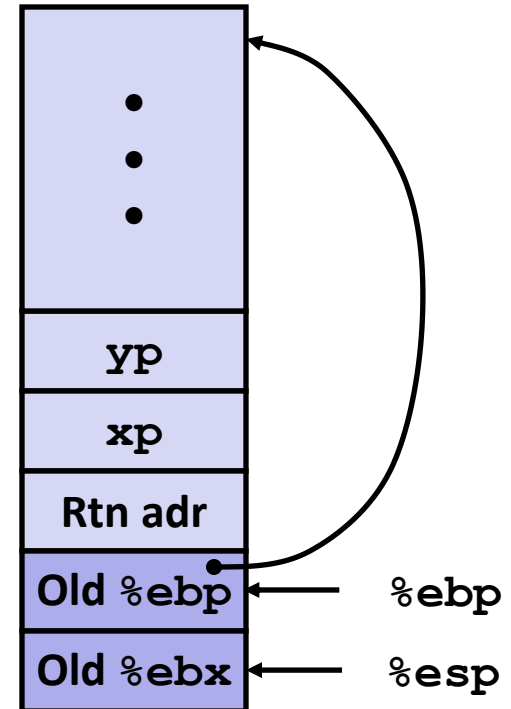
## Entering Stack



`swap:`

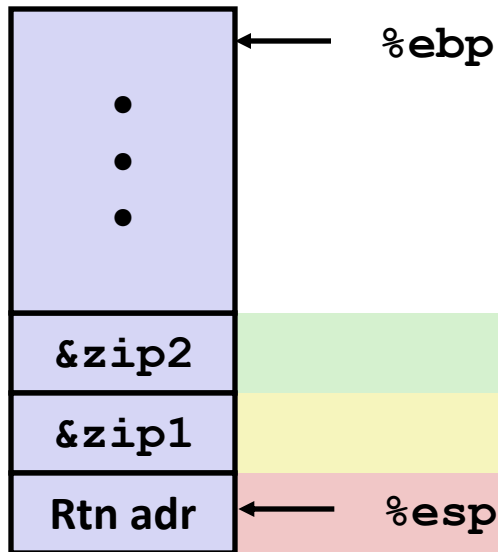
```
    pushl %ebp  
    movl %esp,%ebp  
    pushl %ebx
```

## Resulting Stack



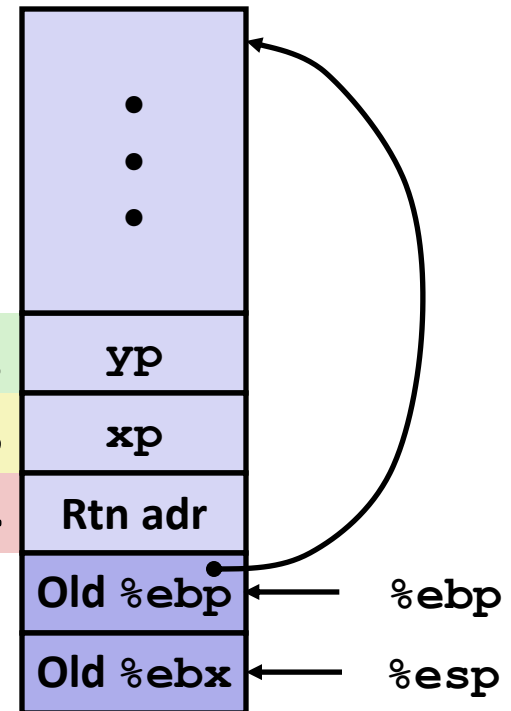
# swap Body

## Entering Stack



## Resulting Stack

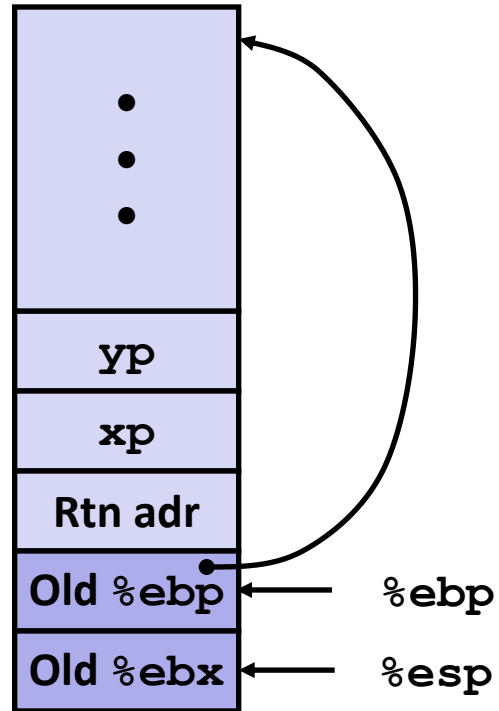
*Offset relative  
to new %ebp*



```
movl 12(%ebp), %ecx # get yp
movl 8(%ebp), %edx  # get xp
. . .
```

# swap Finish #1

swap' s Stack



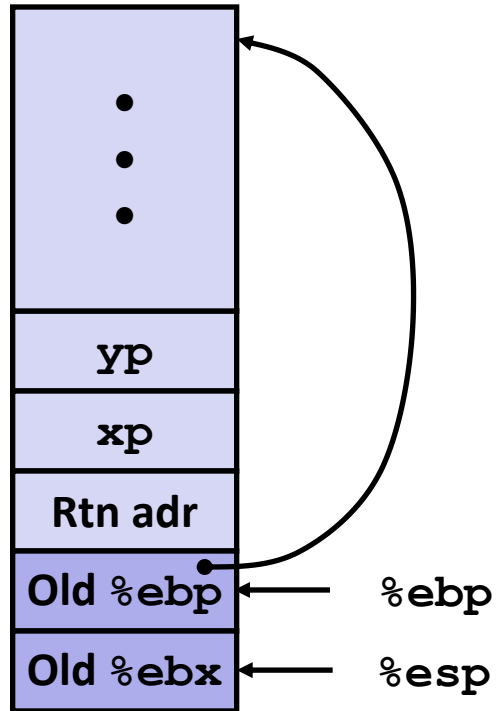
Resulting Stack?

```
movl -4(%ebp), %ebx
movl %ebp, %esp
popl %ebp
ret
```



# swap Finish #1

swap' s Stack

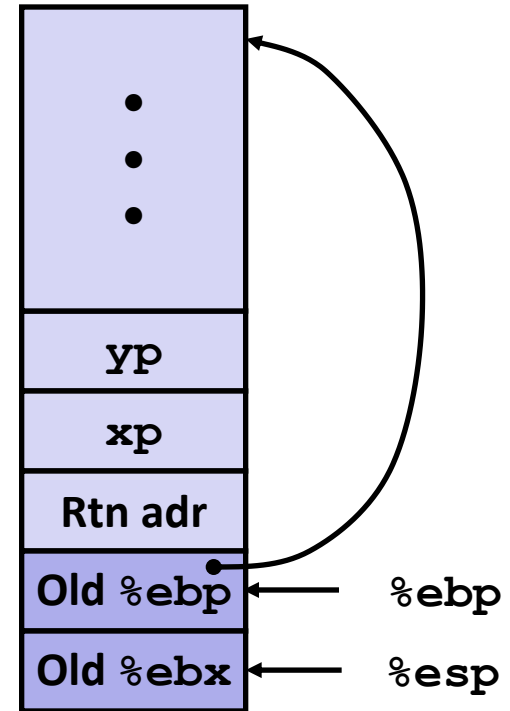


```

movl  -4(%ebp), %ebx
movl  %ebp, %esp
popl  %ebp
ret

```

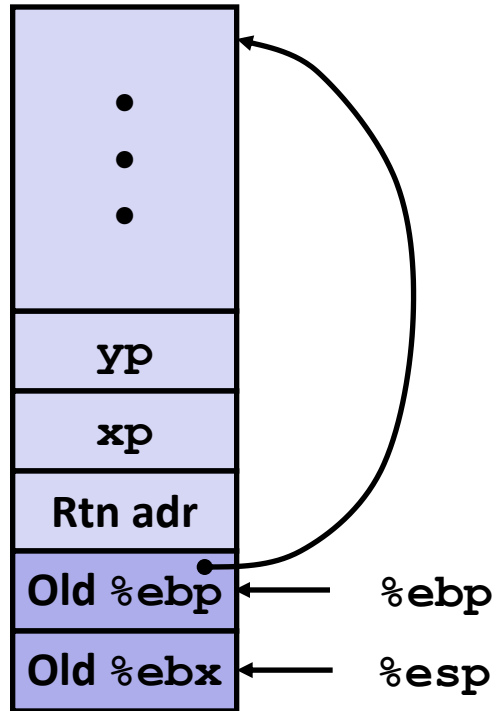
Resulting Stack



**Observation: Saved and restored register `%ebx`**

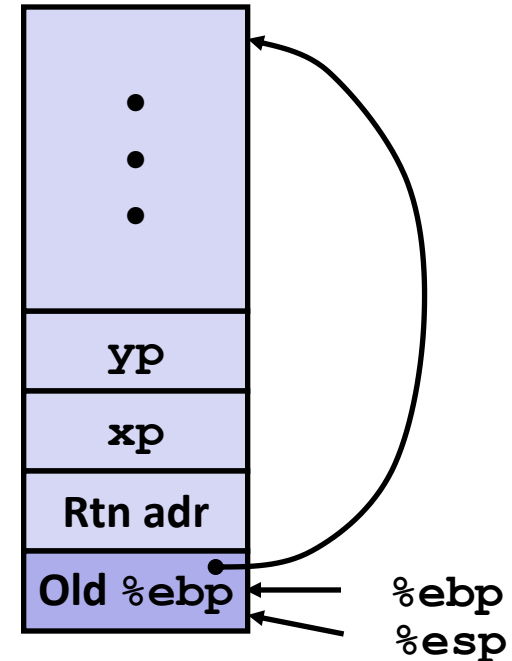
# swap Finish #2

swap' s Stack



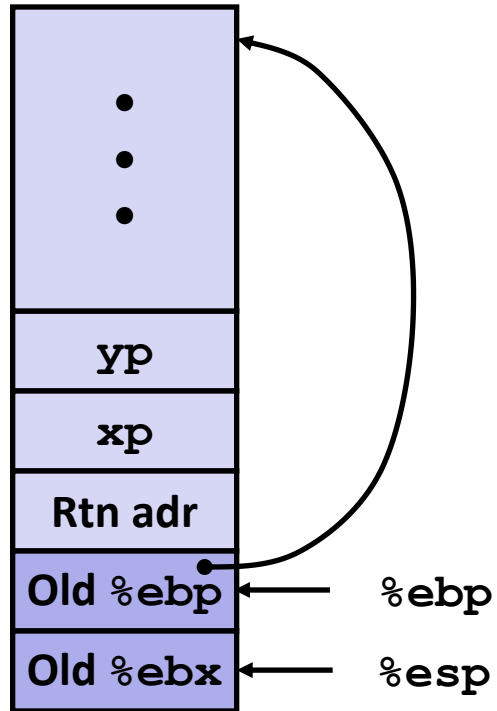
```
movl -4(%ebp), %ebx
movl %ebp, %esp
popl %ebp
ret
```

Resulting Stack



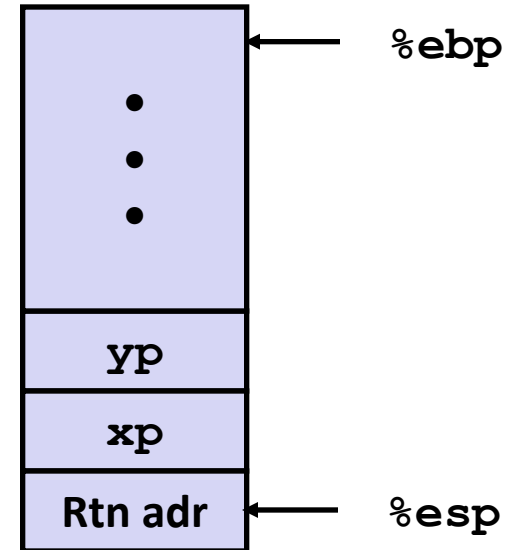
# swap Finish #3

swap' s Stack



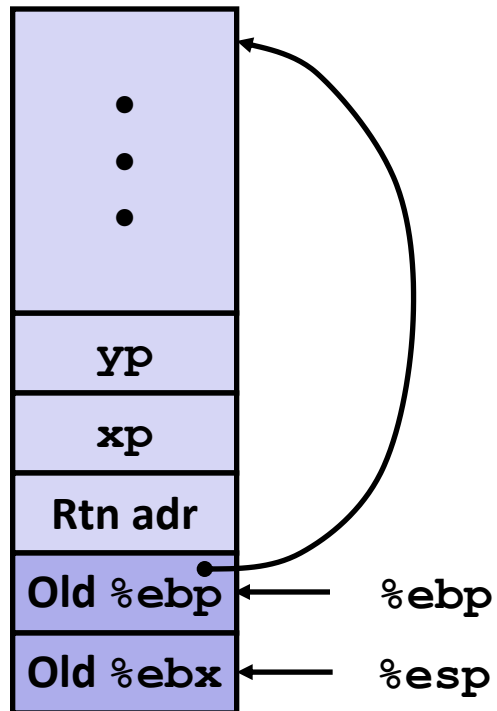
```
movl -4(%ebp), %ebx
movl %ebp, %esp
popl %ebp
ret
```

Resulting Stack



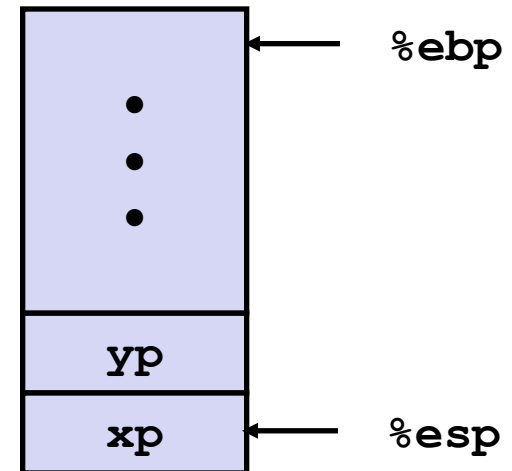
# swap Finish #4

swap' s Stack



```
movl -4(%ebp), %ebx
movl %ebp, %esp
popl %ebp
ret
```

Resulting Stack



# Disassembled swap

080483a4 <swap>:

```

80483a4:  55          push    %ebp
80483a5:  89 e5       mov     %esp, %ebp
80483a7:  53          push    %ebx
80483a8:  8b 55 08    mov     0x8(%ebp), %edx
80483ab:  8b 4d 0c    mov     0xc(%ebp), %ecx
80483ae:  8b 1a       mov     (%edx), %ebx
80483b0:  8b 01       mov     (%ecx), %eax
80483b2:  89 02       mov     %eax, (%edx)
80483b4:  89 19       mov     %ebx, (%ecx)
80483b6:  5b          pop     %ebx
80483b7:  c9          leave  %ecx
80483b8:  c3          ret

```



```

mov    %ebp, %esp
pop    %ebp

```

## Calling Code

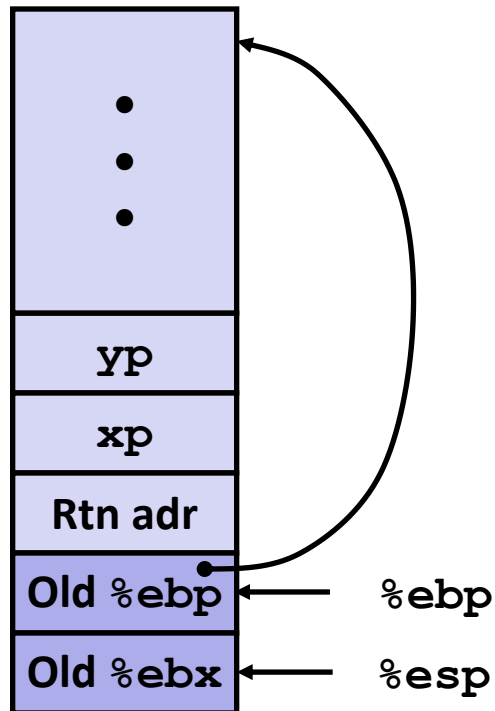
```

8048409:  e8 96 ff ff ff  call  80483a4 <swap>
804840e:  8b 45 f8       mov     0xfffffffff8(%ebp), %eax

```

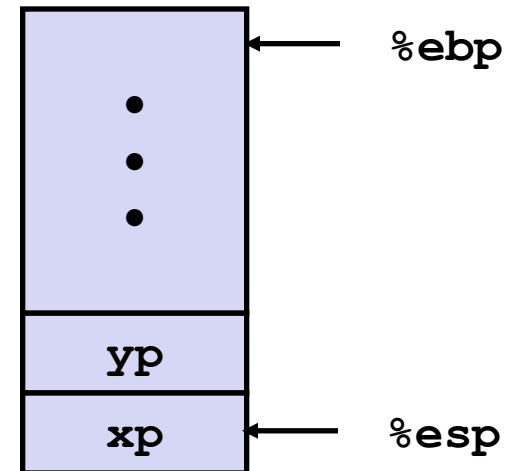
# swap Finish #4

swap' s Stack



```
movl -4(%ebp), %ebx
movl %ebp, %esp
popl %ebp
ret
```

Resulting Stack



## ■ Observation

- Saved & restored register %ebx
- Didn't do so for %eax, %ecx, or %edx

# Register Saving Conventions

## ■ When procedure `yoo` calls `who`:

- `yoo` is the *caller*
- `who` is the *callee*

## ■ Can a register be used for temporary storage?

`yoo:`

```
    . . .  
    movl $12345, %edx  
    call who  
    addl %edx, %eax  
    . . .  
    ret
```

`who:`

```
    . . .  
    movl 8(%ebp), %edx  
    addl $98195, %edx  
    . . .  
    ret
```

- Contents of register `%edx` overwritten by `who`

# Register Saving Conventions

- When procedure `yoo` calls `who`:
  - `yoo` is the *caller*
  - `who` is the *callee*
- Can a register be used for temporary storage?
- Conventions
  - “*Caller Save*”
    - Caller saves temporary values in its frame before calling
  - “*Callee Save*”
    - Callee saves temporary values in its frame before using



# IA32/Linux Register Usage

## ■ **%eax, %edx, %ecx**

- Caller saves prior to call if values are used later

## ■ **%eax**

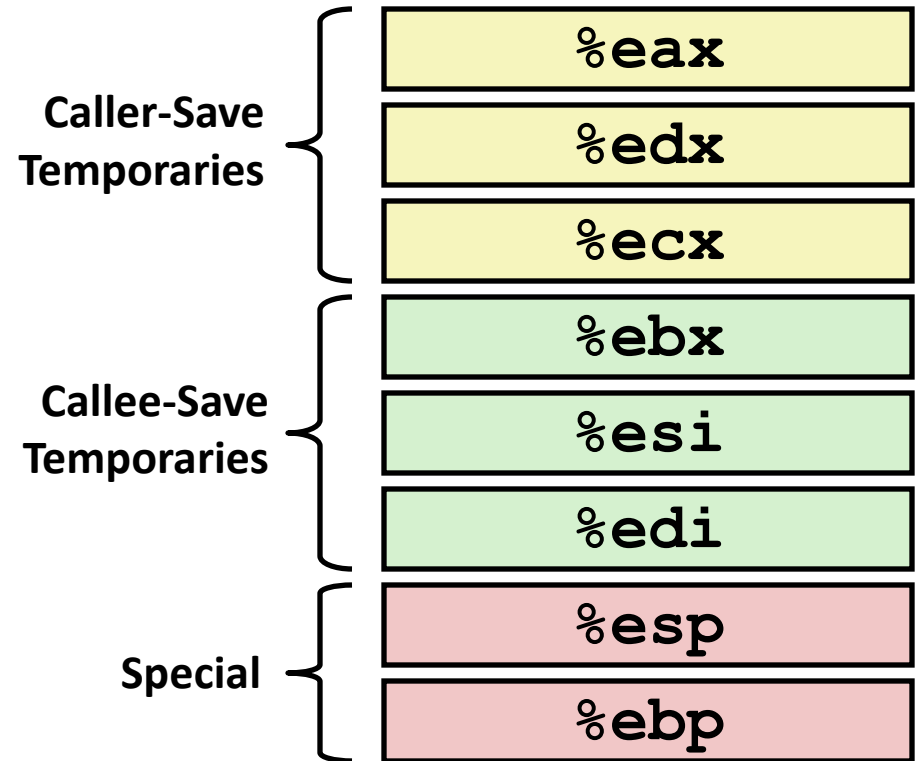
- also used to return integer value

## ■ **%ebx, %esi, %edi**

- Callee saves if wants to use them

## ■ **%esp, %ebp**

- special form of callee save – restored to original values upon exit from procedure



# Example: Pointers to Local Variables

## Recursive Procedure

```
void s_helper
(int x, int *accum)
{
    if (x <= 1)
        return;
    else {
        int z = *accum * x;
        *accum = z;
        s_helper (x-1, accum);
    }
}
```

## Top-Level Call

```
int sfact(int x)
{
    int val = 1;
    s_helper(x, &val);
    return val;
}
```

- Pass pointer to update location

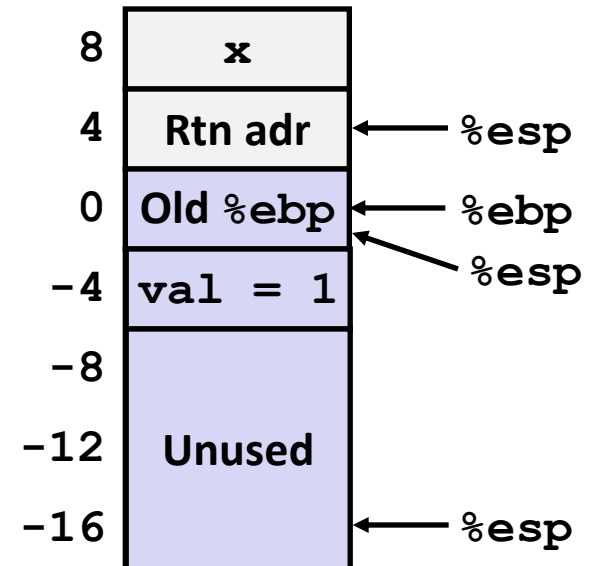
# Creating & Initializing Pointer

```
int sfact(int x)
{
    int val = 1;
    s_helper(x, &val);
    return val;
}
```

- Variable `val` must be stored on stack
  - Because: Need to create pointer to it
- Compute pointer as `-4 (%ebp)`
- Push on stack as second argument

## Initial part of `sfact`

```
_sfact:
    pushl %ebp           # Save %ebp
    movl %esp,%ebp       # Set %ebp
    subl $16,%esp        # Add 16 bytes
    movl 8(%ebp),%edx     # edx = x
    movl $1,-4(%ebp)      # val = 1
```



# Passing Pointer

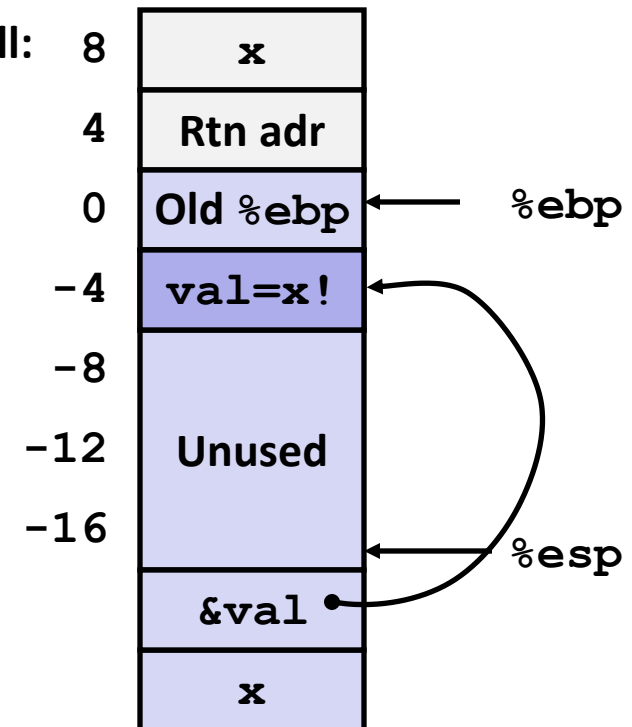
```
int sfact(int x)
{
    int val = 1;
    s_helper(x, &val);
    return val;
}
```

- Variable `val` must be stored on stack
  - Because: Need to create pointer to it
- Compute pointer as `-4 (%ebp)`
- Push on stack as second argument

Stack at time of call:

Calling `s_helper` from `sfact`

```
leal -4(%ebp), %eax # Compute &val
pushl %eax          # Push on stack
pushl %edx          # Push x
call s_helper       # call
movl -4(%ebp), %eax # Return val
. . .              # Finish
```



# IA 32 Procedure Summary

## ■ Important points:

- IA32 procedures are a *combination of instructions and conventions*
  - Conventions prevent functions from disrupting each other
- Stack is the right data structure for procedure call / return
  - If P calls Q, then Q returns before P

## ■ Recursion handled by normal calling conventions

- Can safely store values in local stack frame and in callee-saved registers
- Put function arguments at top of stack
- Result returned in **%eax**

