

# **The Hardware/Software Interface**

CSE351 Spring 2013

## **Data Structures II: Structs and Unions**

# Data Structures in Assembly

## ■ Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

## ■ Structs

- Alignment

## ■ Unions

# Structures

```
struct rec {  
    int i;  
    int a[3];  
    int *p;  
};
```

# Structures

```
struct rec {  
    int i;  
    int a[3];  
    int *p;  
};
```

## Memory Layout



## ■ Characteristics

- Contiguously-allocated region of memory
- Refer to members within structure by names
- Members may be of different types

# Structures

## ■ Accessing Structure Member

- Given an instance of the struct, we can use the `.` operator, just like Java:

- `struct rec r1; r1.i = val;`

- What if we have a *pointer* to a struct: `struct rec *r = &r1;`

```
struct rec {  
    int i;  
    int a[3];  
    int *p;  
};
```

# Structures

## ■ Accessing Structure Member

- Given an instance of the struct, we can use the `.` operator, just like Java:
  - `struct rec r1; r1.i = val;`
- What if we have a *pointer* to a struct: `struct rec *r = &r1;`
  - Using `*` and `.` operators: `(*r).i = val;`
  - Or, use `->` operator for short: `r->i = val;`
- Pointer indicates first byte of structure; access members with offsets

```
struct rec {
    int i;
    int a[3];
    int *p;
};
```


```
void
set_i(struct rec *r,
      int val)
{
    r->i = val;
}
```

## IA32 Assembly

```
# %eax = val
# %edx = r
movl %eax, (%edx)    # Mem[r] = val
```

# Generating Pointer to Structure Member

```
struct rec {
    int i;
    int a[3];
    int *p;
};
```



## ■ Generating Pointer to Array Element


- Offset of each structure member determined at compile time

```
int *find_a
(struct rec *r, int idx)
{
    return &r->a[idx];
}
```

```
# %ecx = idx
# %edx = r
leal 0(,%ecx,4),%eax    # 4*idx
leal 4(%eax,%edx),%eax  # r+4*idx+4
```

# Structures & Alignment

## ■ Unaligned Data




```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

## ■ How would it look like if data were *aligned*?



# Structures & Alignment

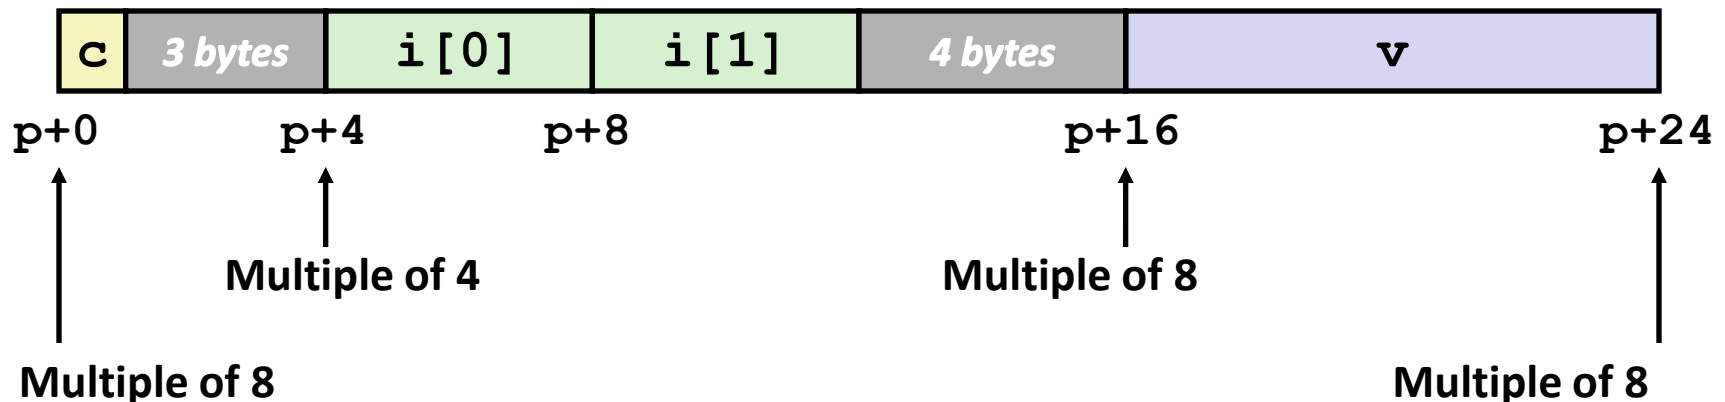
## ■ Unaligned Data



```
struct S1 {
    char c;
    int i[2];
    double v;
} *p;
```

## ■ Aligned Data

- Primitive data type requires  $K$  bytes
- Address must be multiple of  $K$



# Alignment Principles

- **Aligned Data**
  - Primitive data type requires K bytes
  - Address must be multiple of K
- **Aligned data is required on some machines; it is *advised* on IA32**
  - Treated differently by IA32 Linux, x86-64 Linux, and Windows!
- **What is the motivation for alignment?**

# Alignment Principles

## ■ Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K

## ■ Aligned data is required on some machines; it is *advised* on IA32

- Treated differently by IA32 Linux, x86-64 Linux, and Windows!

## ■ Motivation for Aligning Data

- Physical memory is accessed by aligned chunks of 4 or 8 bytes (system-dependent)
  - Inefficient to load or store datum that spans quad word boundaries
- Also, virtual memory is very tricky when datum spans two pages (later...)

## ■ Compiler

- Inserts padding in structure to ensure correct alignment of fields
- `sizeof()` should be used to get true size of structs

# Specific Cases of Alignment (IA32)

- **1 byte: char, ...**
  - no restrictions on address
- **2 bytes: short, ...**
  - lowest 1 bit of address must be  $0_2$
- **4 bytes: int, float, char \*, ...**
  - lowest 2 bits of address must be  $00_2$
- **8 bytes: double, ...**
  - Windows (and most other OSs & instruction sets): lowest 3 bits  $000_2$
  - Linux: lowest 2 bits of address must be  $00_2$ 
    - i.e., treated the same as a 4-byte primitive data type
- **12 bytes: long double**
  - Windows, Linux: lowest 2 bits of address must be  $00_2$

# Satisfying Alignment with Structures

- **Within structure:**
  - Must satisfy every member's alignment requirement
- **Overall structure placement**
  - Each structure has alignment requirement  $K$ 
    - $K$  = Largest alignment of any element
  - Initial address & structure length must be multiples of  $K$
- **Example (under Windows or x86-64):**  $K = ?$

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p1;
```

# Satisfying Alignment with Structures

## ■ Within structure:

- Must satisfy every member's alignment requirement

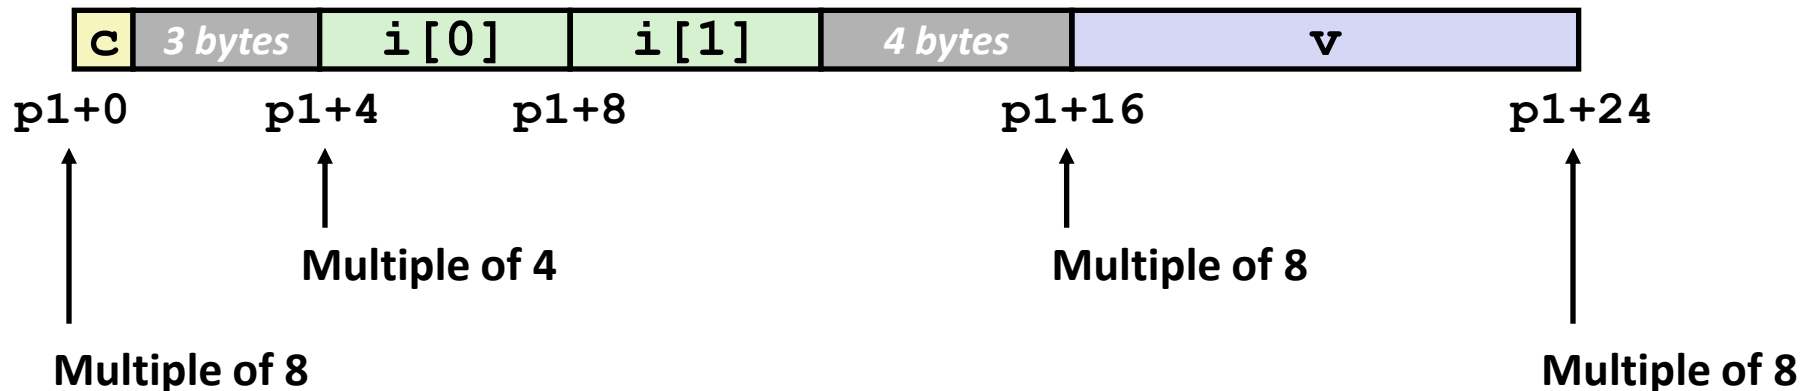
## ■ Overall structure placement

- Each structure has alignment requirement  $K$ 
  - $K$  = Largest alignment of any element
- Initial address & structure length must be multiples of  $K$

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p1;
```

## ■ Example (under Windows or x86-64): $K = ?$


- $K = 8$ , due to **double** member



# Different Alignment Conventions

- **IA32 Windows or x86-64:**
  - $K = 8$ , due to **double** member

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p1;
```




- **IA32 Linux:**  $K = ?$

# Different Alignment Conventions

## ■ IA32 Windows or x86-64:


- $K = 8$ , due to **double** member

```
struct S1 {
    char c;
    int i[2];
    double v;
} *p1;
```



## ■ IA32 Linux: $K = ?$


- $K = 4$ ; **double** aligned like a 4-byte data type





# Saving Space


```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p1;
```



# Saving Space

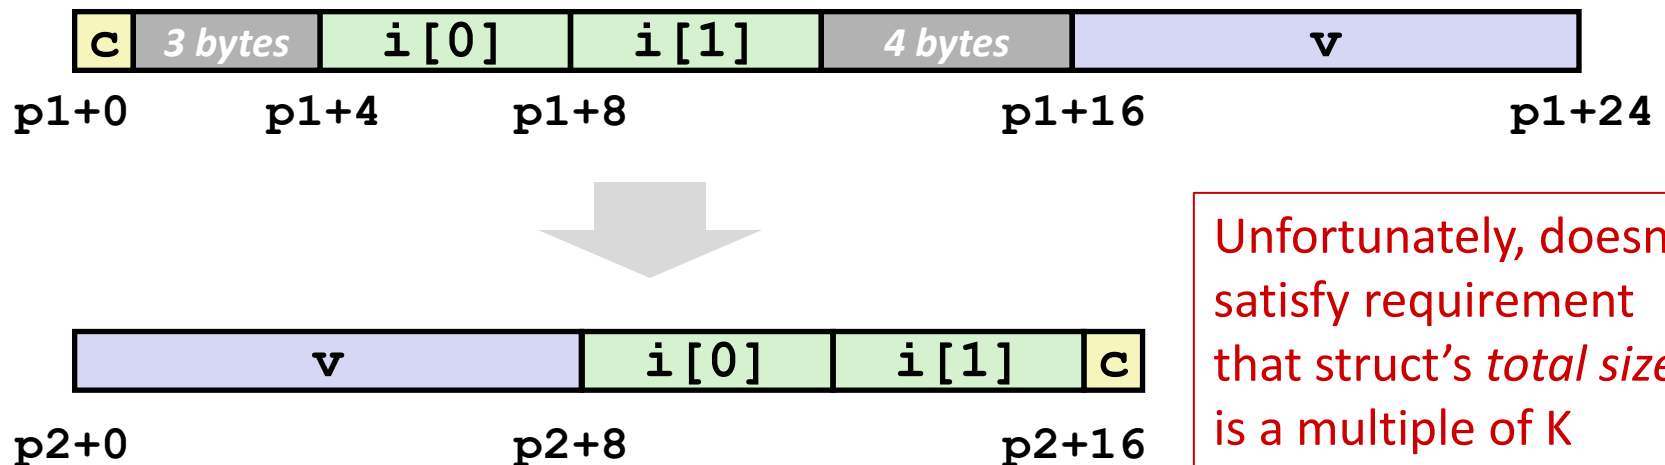
- Put large data types first:

```
struct S1 {
    char c;
    int i[2];
    double v;
} *p1;
```



```
struct S2 {
    double v;
    int i[2];
    char c;
} *p2;
```

- Effect (example x86-64, both have  $K=8$ )

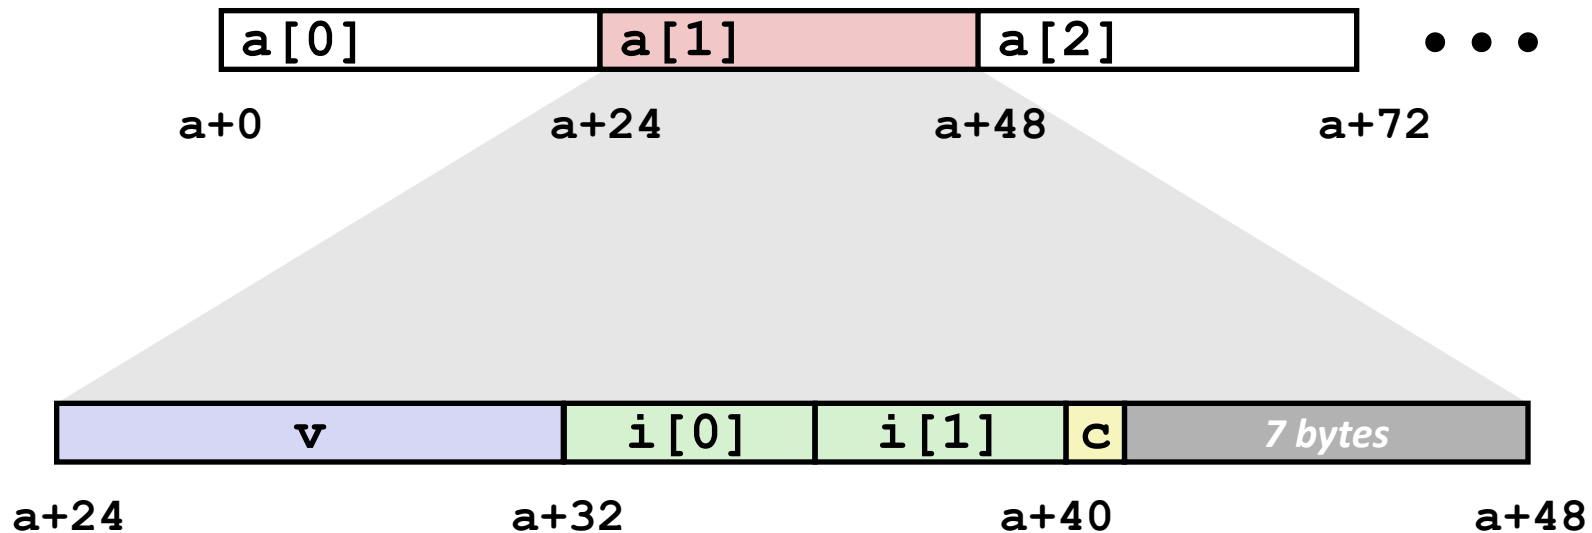


Unfortunately, doesn't satisfy requirement that struct's *total size* is a multiple of  $K$

# Arrays of Structures

- Satisfy alignment requirement for every element
- How would accessing an element work?


```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



# Saving Space



- Put large data types first:

```
struct S3 {  
    char c;  
    int i;  
    char d;  
} *p3;
```



```
struct S4 {  
    int i;  
    char c;  
    char d;  
} *p4;
```

- Effect (K=4)




- This strategy *can* save some space for certain structs.

# Unions

- Allocated according to largest element
- Can only use one member at a time

```
union U1 {  
    char c;  
    int i[2];  
    double v;  
} *up;
```

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *sp;
```





# Unions

- Allocated according to largest element
- Can only use one member at a time

```
union U1 {
    char c;
    int i[2];
    double v;
} *up;
```

```
struct S1 {
    char c;
    int i[2];
    double v;
} *sp;
```



# What Are Unions Good For?

- **Unions allow the same region of memory to be referenced as different types**
  - Different “views” of the same memory location
  - Can be used to circumvent C’s type system (bad idea)
- **Better idea: use a struct inside a union to access some memory location either as a whole or by its parts**

# Unions For Embedded Programming

```
typedef union
{
    unsigned char byte;
    struct {
        unsigned char b0:1;
        unsigned char b1:1;
        unsigned char b2:1;
        unsigned char b3:1;
        unsigned char reserved:4;
    } bits;
} hw_register;

hw_register reg;
reg.byte = 0x3F;           // 001111112
reg.bits.b2 = 0;           // 001110112
reg.bits.b3 = 0;           // 001100112
unsigned short a = reg.byte;
printf("0x%X\n", a);      // output: 0x33
```

(Note: the placement of these fields and other parts of this example are implementation-dependent)



# Summary

## ■ Arrays in C

- Contiguous allocations of memory
- No bounds checking
- Can usually be treated like a pointer to first element

## ■ Structures

- Allocate bytes in order declared
- Pad in middle and at end to satisfy alignment

## ■ Unions

- Provide different views of the same memory location