


The Hardware/Software Interface

CSE351 Spring 2013

Buffer Overflow



Buffer Overflow

- Buffer overflows are possible because C doesn't check array boundaries
- Buffer overflows are *dangerous* because buffers for user input are often stored on the stack
 - Probably the most common type of security vulnerability
- Today we'll go over:
 - Address space layout
 - Input buffers on the stack
 - Overflowing buffers and injecting code
 - Defenses against buffer overflows

IA32 Linux Memory Layout

■ Stack

- Runtime stack (8MB limit)

■ Heap

- Dynamically allocated storage
- Allocated by `malloc()`, `calloc()`, `new()`

■ Data



- Statically allocated data
 - Read-only: string literals
 - Read/write: global arrays and variables

■ Text

- Executable machine instructions
- Read-only

not drawn to scale

Upper 2 hex digits
= 8 bits of address



Memory Allocation Example

```
char big_array[1<<24]; /* 16 MB */
char huge_array[1<<28]; /* 256 MB */


int beyond;
char *p1, *p2, *p3, *p4;

int useless() { return 0; }

int main()
{
    p1 = malloc(1 << 28); /* 256 MB */
    p2 = malloc(1 << 8); /* 256 B */
    p3 = malloc(1 << 28); /* 256 MB */
    p4 = malloc(1 << 8); /* 256 B */
    /* Some print statements ... */
}
```

Where does everything go?

not drawn to scale




IA32 Example Addresses

address range $\sim 2^{32}$

<code>\$esp</code>	<code>0xffffbcd0</code>
<code>p3</code>	<code>0x65586008</code>
<code>p1</code>	<code>0x55585008</code>
<code>p4</code>	<code>0x1904a110</code>
<code>p2</code>	<code>0x1904a008</code>
<code>&p2</code>	<code>0x18049760</code>
<code>beyond</code>	<code>0x08049744</code>
<code>big_array</code>	<code>0x18049780</code>
<code>huge_array</code>	<code>0x08049760</code>
<code>main()</code>	<code>0x080483c6</code>
<code>useless()</code>	<code>0x08049744</code>
<code>final malloc()</code>	<code>0x006be166</code>

`malloc()` is dynamically linked
address determined at runtime

not drawn to scale



Internet Worm

- **These characteristics of the traditional IA32 Linux memory layout provide opportunities for malicious programs**
 - Stack grows “backwards” in memory
 - Data and instructions both stored in the same memory

- **November, 1988**
 - Internet Worm attacks thousands of Internet hosts.
 - How did it happen?

- **The Internet Worm was based on *stack buffer overflow* exploits!**
 - Many Unix functions do not check argument sizes
 - Allows target buffers to overflow

String Library Code

■ Implementation of Unix function gets ()

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getchar();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getchar();
    }
    *p = '\0';
    return dest;
}
```

- What could go wrong in this code?

String Library Code

■ Implementation of Unix function `gets()`

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getchar();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getchar();
    }
    *p = '\0';
    return dest;
}
```

- No way to specify limit on number of characters to read

■ Similar problems with other Unix functions

- **strcpy**: Copies string of arbitrary length
- **scanf**, **fscanf**, **sscanf**, when given **%s** conversion specification

Vulnerable Buffer Code

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
int main()  
{  
    printf("Type a string:");  
    echo();  
    return 0;  
}
```

```
unix>./bufdemo  
Type a string:1234567  
1234567
```

```
unix>./bufdemo  
Type a string:12345678  
Segmentation Fault
```

```
unix>./bufdemo  
Type a string:123456789ABC  
Segmentation Fault
```

Buffer Overflow Disassembly

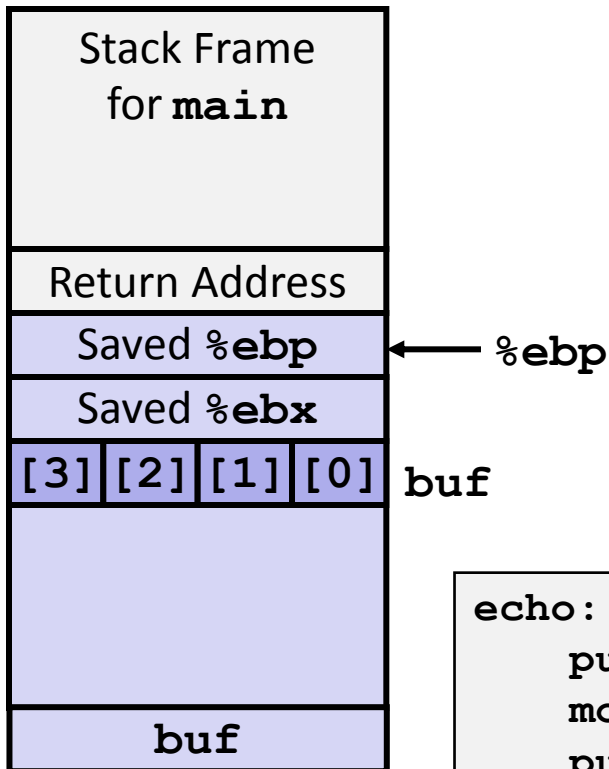
080484f0 <echo>:

80484f0:	55	push	%ebp
80484f1:	89 e5	mov	%esp, %ebp
80484f3:	53	push	%ebx
80484f4:	8d 5d f8	lea	0xfffffffff8(%ebp), %ebx
80484f7:	83 ec 14	sub	\$0x14, %esp
80484fa:	89 1c 24	mov	%ebx, (%esp)
80484fd:	e8 ae ff ff ff	call	80484b0 <gets>
8048502:	89 1c 24	mov	%ebx, (%esp)
8048505:	e8 8a fe ff ff	call	8048394 <puts@plt>
804850a:	83 c4 14	add	\$0x14, %esp
804850d:	5b	pop	%ebx
804850e:	c9	leave	
804850f:	c3	ret	

80485f2:	e8 f9 fe ff ff	call	80484f0 <echo>
80485f7:	8b 5d fc	mov	0xfffffffffc(%ebp), %ebx
80485fa:	c9	leave	
80485fb:	31 c0	xor	%eax, %eax
80485fd:	c3	ret	

Buffer Overflow Stack

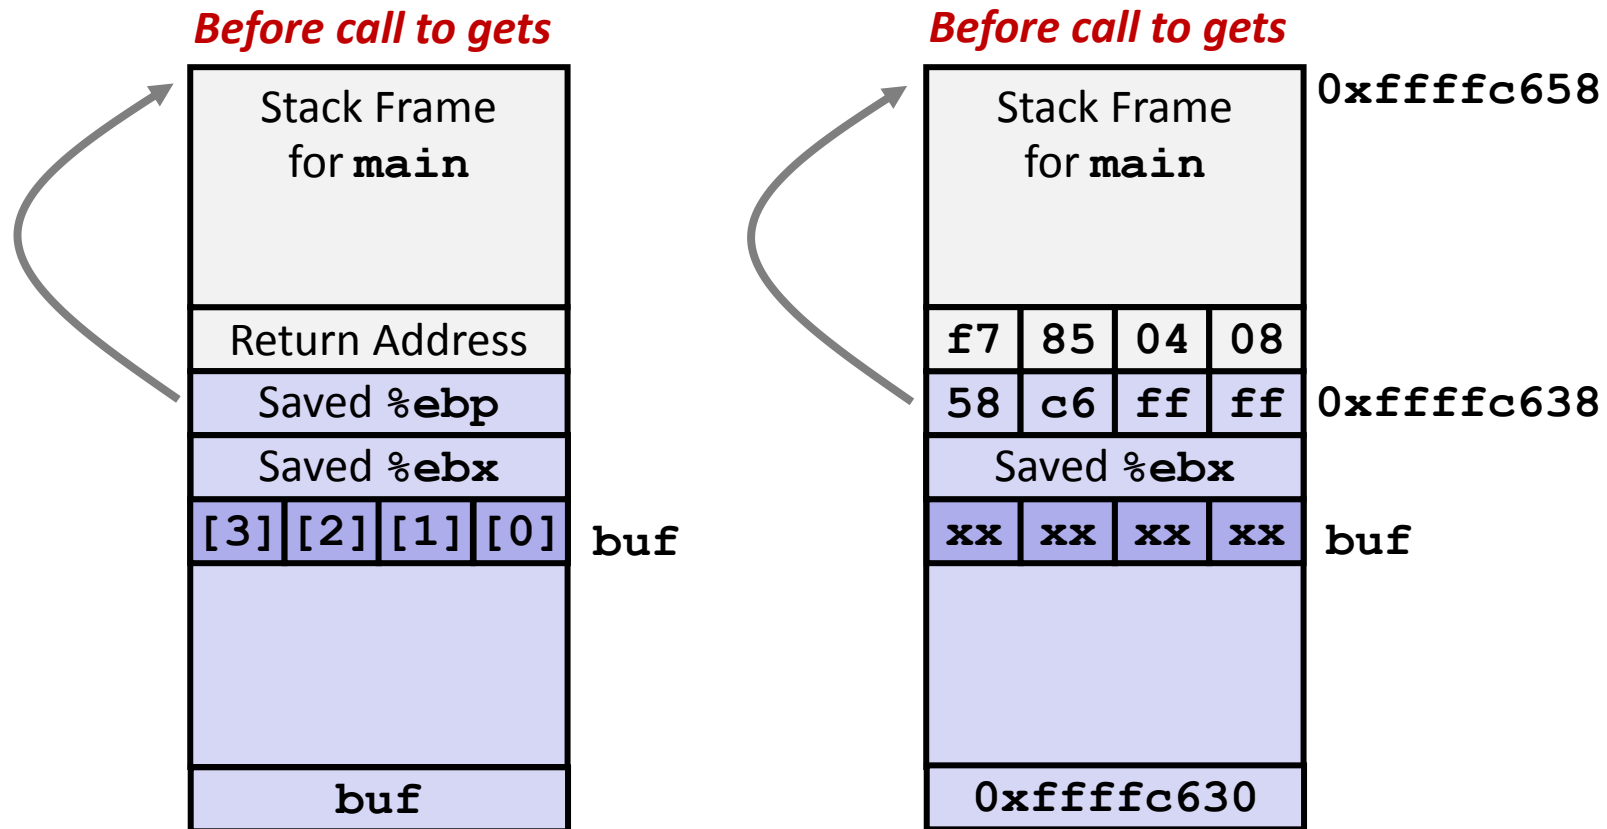
Before call to gets



```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

```
echo:
    pushl %ebp                # Save %ebp on stack
    movl %esp, %ebp
    pushl %ebx                # Save %ebx
    leal -8(%ebp), %ebx       # Compute buf as %ebp-8
    subl $20, %esp            # Allocate stack space
    movl %ebx, (%esp)         # Push buf addr on stack
    call gets                  # Call gets
    . . .
```

Buffer Overflow Stack Example




```
80485f2: call 80484f0 <echo>
```


```
80485f7: mov 0xffffffffc(%ebp), %ebx # Return Point
```

Buffer Overflow Example #1

Before call to gets




Input 1234567



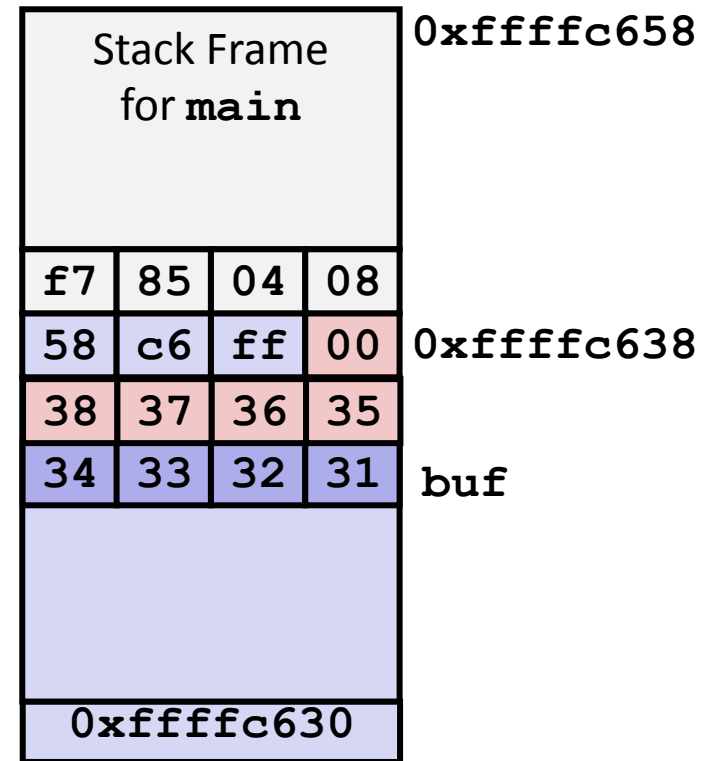
Overflow buf, and corrupt saved `%ebx`, but no problem

Buffer Overflow Example #2

Before call to gets



Input 12345678




Frame pointer corrupted

```

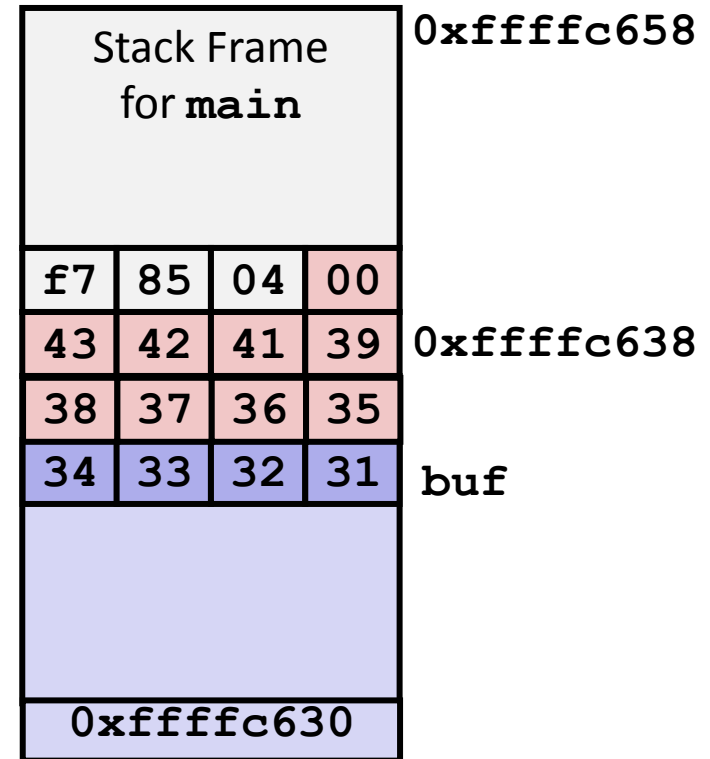
. . . .
804850a: 83 c4 14  add    $0x14,%esp  # deallocate space
804850d: 5b         pop     %ebx      # restore %ebx
804850e: c9         leave   # movl %ebp, %esp; popl %ebp
804850f: c3         ret      # Return
  
```

Buffer Overflow Example #3

Before call to gets



Input 123456789ABC

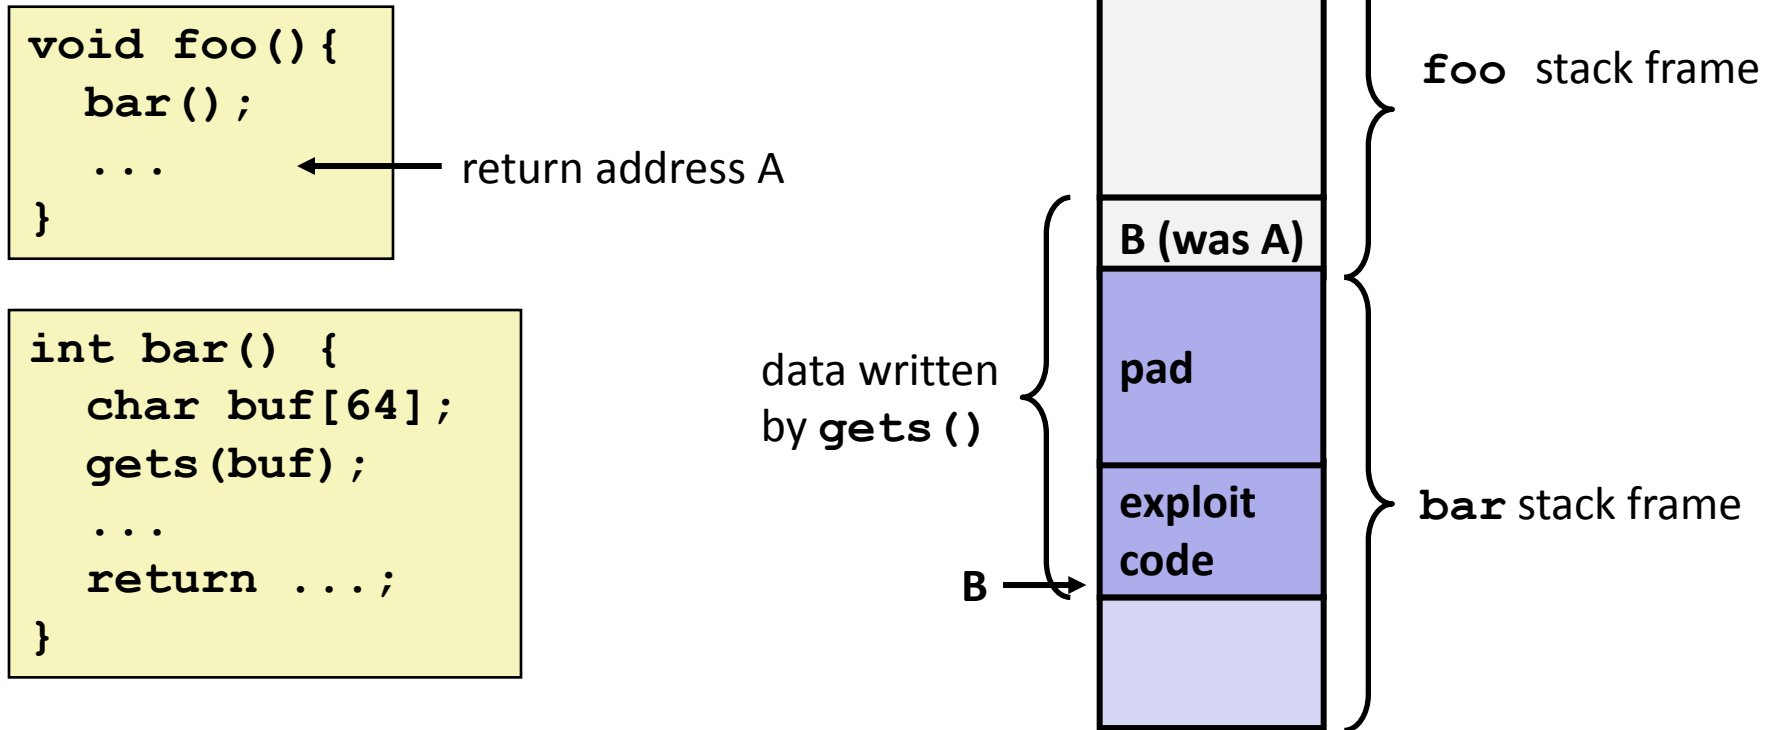


Return address corrupted

```
080485f2: call 80484f0 <echo>
```

```
080485f7: mov 0xffffffffc(%ebp),%ebx # Return Point
```

Malicious Use of Buffer Overflow



- Input string contains byte representation of executable code
- Overwrite return address A with address of buffer (need to know B)
- When `bar()` executes `ret`, will jump to exploit code (instead of A)

Exploits Based on Buffer Overflows

- *Buffer overflow bugs allow remote machines to execute arbitrary code on victim machines*
- **Internet worm**
 - Early versions of the finger server (fingerd) used `gets ()` to read the argument sent by the client:
 - `finger droh@cs.cmu.edu`
 - Worm attacked fingerd server by sending phony argument:
 - `finger "exploit-code padding new-return-address"`
 - exploit code: executed a root shell on the victim machine with a direct TCP connection to the attacker

Avoiding Overflow Vulnerability

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    fgets(buf, 4, stdin);  
    puts(buf);  
}
```

- **Use library routines that limit string lengths**
 - **fgets** instead of **gets** (second argument to fgets sets limit)
 - **strncpy** instead of **strcpy**
 - Don't use **scanf** with **%s** conversion specification
 - Use **fgets** to read the string
 - Or use **%ns** where **n** is a suitable integer

System-Level Protections

■ Randomized stack offsets

- At start of program, allocate random amount of space on stack
- Makes it difficult for exploit to predict beginning of inserted code

■ Use techniques to *detect* stack corruption

■ Nonexecutable code segments

- Only allow code to execute from “text” sections of memory
- Do NOT execute code in stack, data, or heap regions
- Hardware support needed

