

# Roadmap

C:

```
car *c = malloc(sizeof(car)) ;
c->miles = 100;
c->gals = 17;
float mpg = get_mpg(c);
free(c) ;
```

Java:

```
Car c = new Car();
c.setMiles(100);
c.setGals(17);
float mpg =
    c.getMPG();
```


Assembly  
language:

```
get_mpg:
    pushq    %rbp
    movq     %rsp, %rbp
    ...
    popq     %rbp
    ret
```

Machine  
code:

```
0111010000011000
100011010000010000000010
1000100111000010
110000011111101000011111
```

Computer  
system:



Data & addressing  
Integers & floats  
Machine code & C  
x86 assembly  
programming  
Procedures &  
stacks  
Arrays & structs  
Memory & caches  
Processes  
Virtual memory  
**Memory allocation**  
Java vs. C

OS:



# Memory Allocation Topics

## ■ Dynamic memory allocation

- Size/number of data structures may only be known at run time
- Need to allocate space on the heap
- Need to de-allocate (free) unused memory so it can be re-allocated

## ■ Implementation

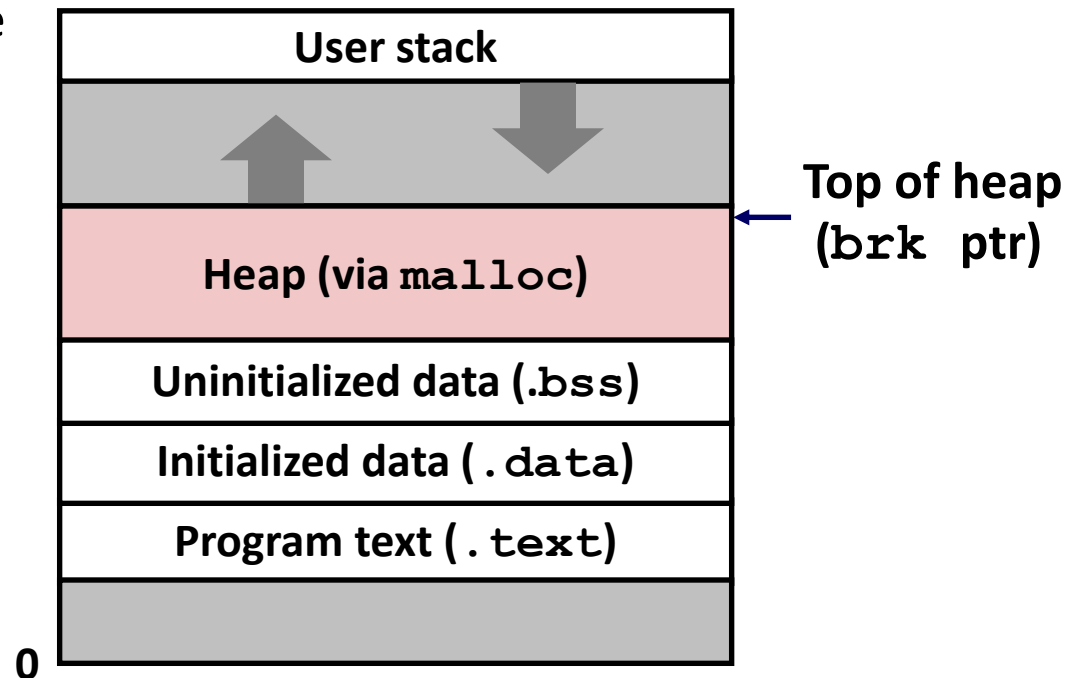

- Implicit free lists
- Explicit free lists – subject of next programming assignment
- Segregated free lists

## ■ Garbage collection

## ■ Common memory-related bugs in C programs

# Dynamic Memory Allocation

- Programmers use *dynamic memory allocators* (such as `malloc`) to acquire VM at run time.
  - For data structures whose size is only known at runtime.
- Dynamic memory allocators manage an area of process virtual memory known as the *heap*.



# Dynamic Memory Allocation

- **Allocator maintains heap as collection of variable sized *blocks*, which are either *allocated* or *free***
  - Allocator requests space in heap region; VM hardware and kernel allocate these pages to the process
  - Application objects are typically smaller than pages, so the allocator manages blocks *within* pages
- **Types of allocators**
  - ***Explicit allocator:*** application allocates and frees space
    - E.g. `malloc` and `free` in C
  - ***Implicit allocator:*** application allocates, but does not free space
    - E.g. garbage collection in Java, ML, and Lisp

# The malloc Package

```
#include <stdlib.h>
```

```
void *malloc(size_t size)
```

- Successful:
  - Returns a pointer to a memory block of at least **size** bytes (typically) aligned to 8-byte boundary
  - If **size == 0**, returns NULL
- Unsuccessful: returns NULL and sets **errno**

```
void free(void *p)
```

- Returns the block pointed at by **p** to pool of available memory
- **p** must come from a previous call to **malloc** or **realloc**

## Other functions

- **calloc**: Version of **malloc** that initializes allocated block to zero.
- **realloc**: Changes the size of a previously allocated block.
- **sbrk**: Used internally by allocators to grow or shrink the heap.

# Malloc Example

```
void foo(int n, int m) {
    int i, *p;

    /* allocate a block of n ints */
    p = (int *)malloc(n * sizeof(int));
    if (p == NULL) {
        perror("malloc");
        exit(0);
    }
    for (i=0; i<n; i++) p[i] = i;


    /* add space for m ints to end of p block */
    if ((p = (int *)realloc(p, (n+m) * sizeof(int))) == NULL) {
        perror("realloc");
        exit(0);
    }
    for (i=n; i < n+m; i++) p[i] = i;

    /* print new array */
    for (i=0; i<n+m; i++)
        printf("%d\n", p[i]);

    free(p); /* return p to available memory pool */
}
```

# Assumptions Made in This Lecture

- **Memory is word addressed (each word can hold a pointer)**
  - block size is a multiple of words




# Allocation Example


```
p1 = malloc(4)
```




```
p2 = malloc(5)
```



```
p3 = malloc(6)
```



```
free(p2)
```



```
p4 = malloc(2)
```





# How are going to implement that?!?

- *What information does the allocator need to keep track of?*

# Constraints

## ■ Applications

- Can issue arbitrary sequence of malloc() and free() requests
- free() requests must be made only for a previously malloc()'d block

## ■ Allocators

- Can't control number or size of allocated blocks
- Must respond immediately to malloc() requests
  - *i.e.*, can't reorder or buffer requests
- Must allocate blocks from free memory
  - *i.e.*, blocks can't overlap
- Must align blocks so they satisfy all alignment requirements
  - 8 byte alignment for GNU malloc (**libc** malloc) on Linux boxes
- Can't move the allocated blocks once they are malloc()'d
  - *i.e.*, compaction is not allowed. *Why not?*

# Performance Goal: Throughput

- Given some sequence of **malloc** and **free** requests:

- $R_0, R_1, \dots, R_k, \dots, R_{n-1}$

- Goals: maximize throughput and peak memory utilization

- These goals are often conflicting

- Throughput:

- Number of completed requests per unit time

- Example:

- 5,000 **malloc**() calls and 5,000 **free**() calls in 10 seconds

- Throughput is 1,000 operations/second

# Performance Goal: Peak Memory Utilization


- Given some sequence of `malloc` and `free` requests:
  - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- **Def: Aggregate payload  $P_k$** 
  - `malloc(p)` results in a block with a **payload** of `p` bytes
  - After request  $R_k$  has completed, the **aggregate payload**  $P_k$  is the sum of currently allocated payloads
- **Def: Current heap size =  $H_k$** 
  - Assume  $H_k$  is monotonically nondecreasing
    - Allocator can increase size of heap using `sbrk()`
- **Def: Peak memory utilization after  $k$  requests**
  - $U_k = (\max_{i \leq k} P_i) / H_k$
  - Goal: maximize utilization for a sequence of requests.
  - *Why is this hard? And what happens to throughput?*

# Fragmentation

- Poor memory utilization is caused by *fragmentation*
  - *internal* fragmentation
  - *external* fragmentation

# Internal Fragmentation

- For a given block, *internal fragmentation* occurs if payload is smaller than block size



- **Caused by**
  - overhead of maintaining heap data structures (inside block, outside payload)
  - padding for alignment purposes
  - explicit policy decisions (e.g., to return a big block to satisfy a small request)  
*why would anyone do that?*
- **Depends only on the pattern of *previous* requests**
  - thus, easy to measure


# External Fragmentation

- Occurs when there is enough aggregate heap memory, but no single free block is large enough


```
p1 = malloc(4)
```



```
p2 = malloc(5)
```



```
p3 = malloc(6)
```



```
free(p2)
```



```
p4 = malloc(6)
```

*Oops! (what would happen now?)*

- Depends on the pattern of future requests
  - Thus, difficult to measure